

# **CS-200**

# **Computer Architecture**

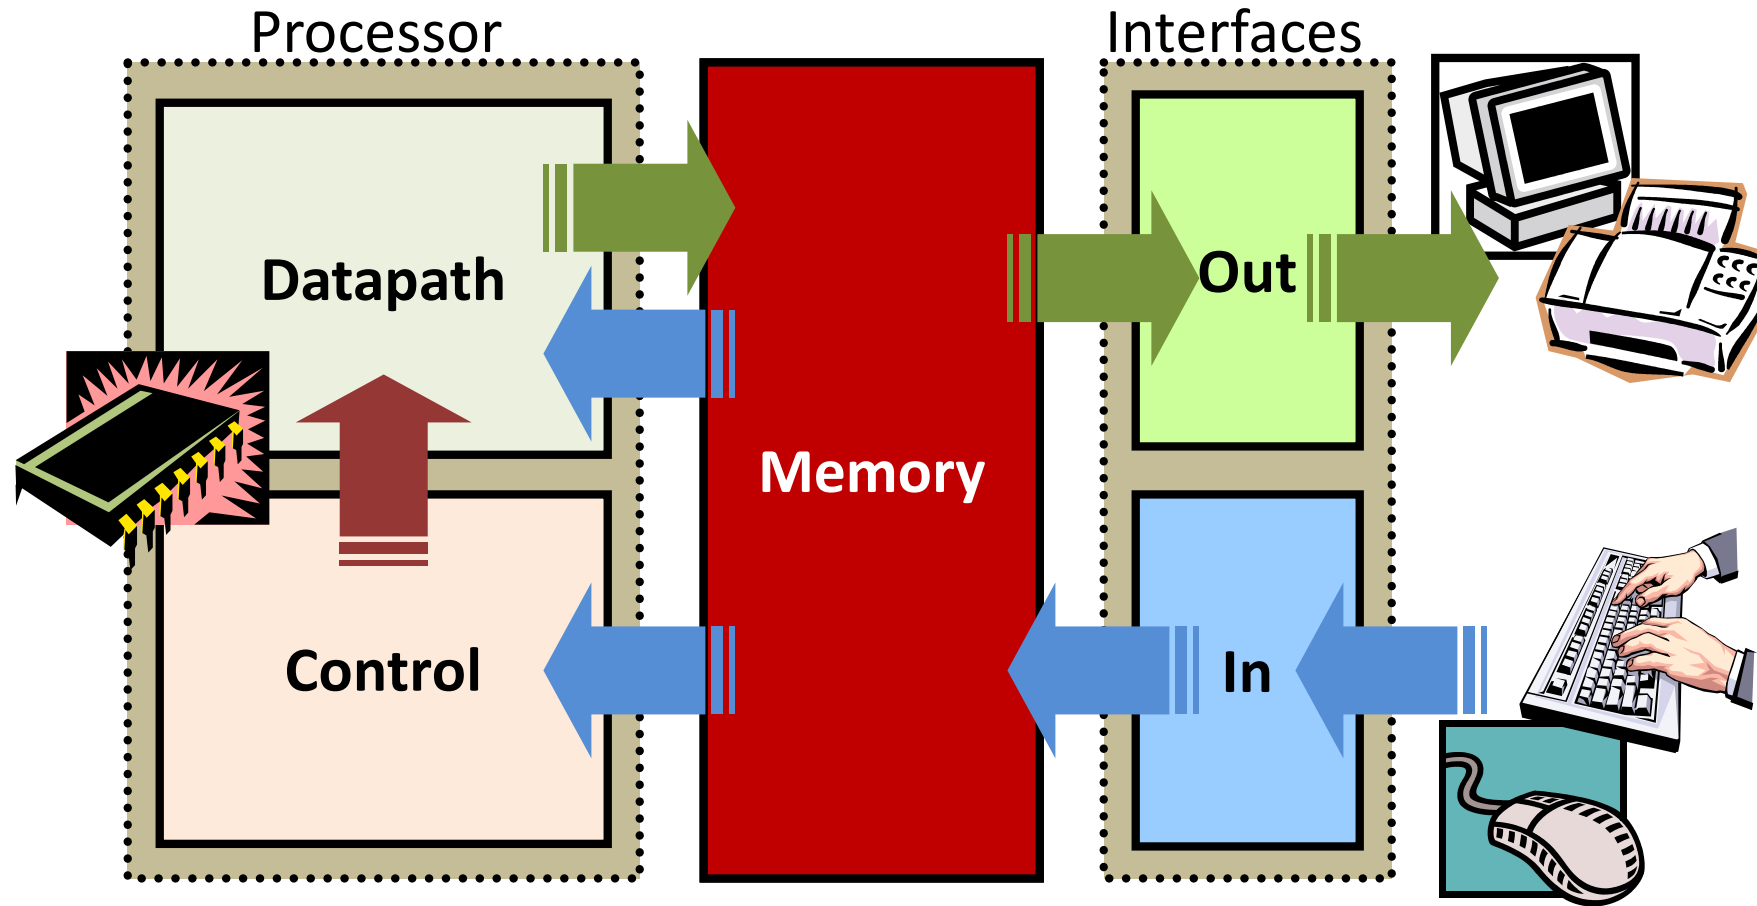
---

## **Part 3a. Memory Hierarchy**


## **Caches**

Paolo Ienne  
<paolo.ienne@epfl.ch>


# The Five Classic Components of a Computer



# Memory Performance over Time



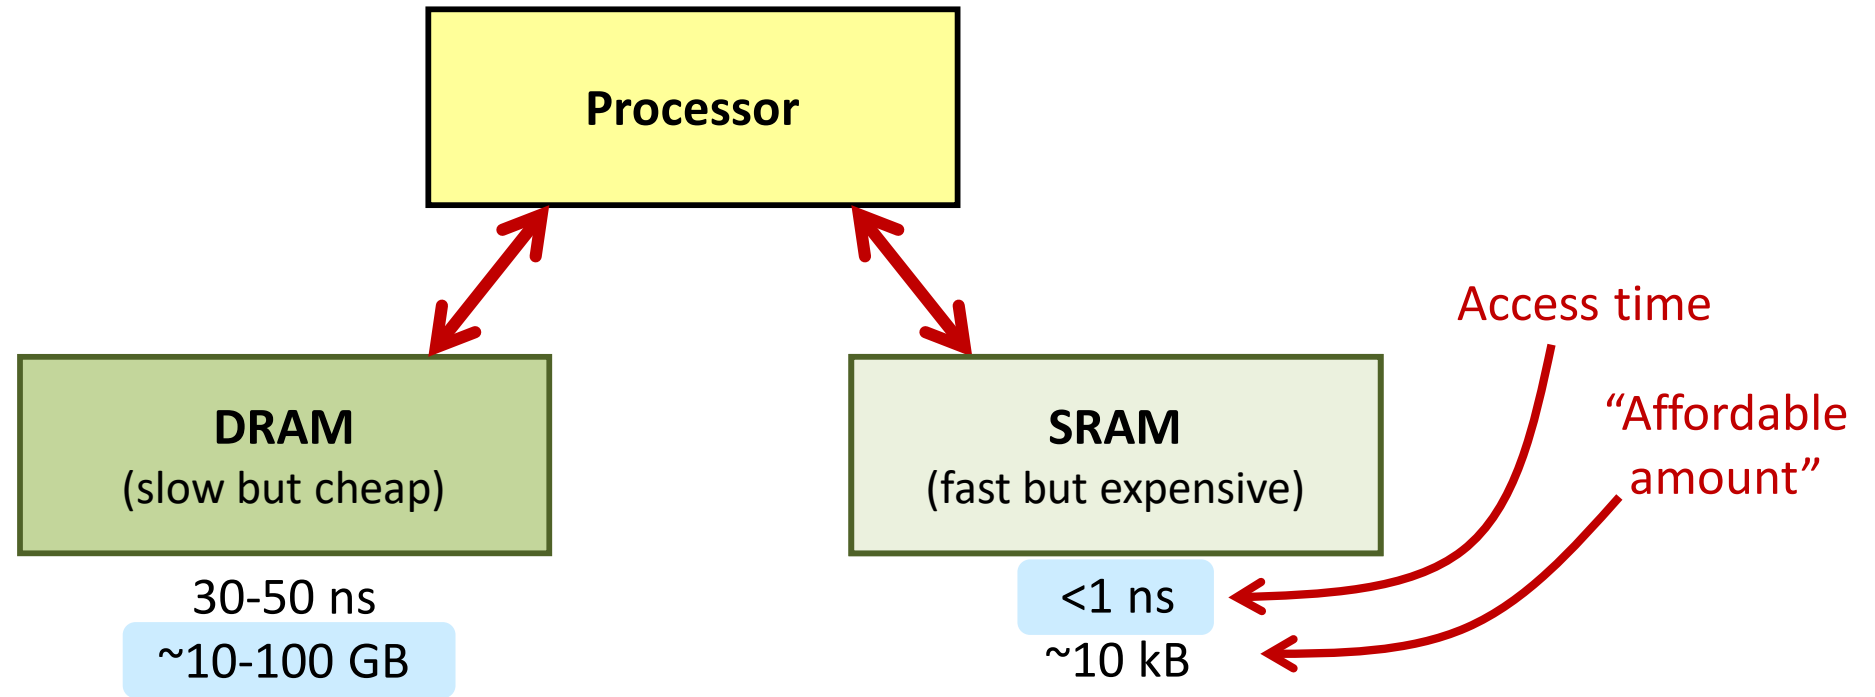
# The Tension between Big and Fast!



This is your memory subsystem  
**You want it big!**

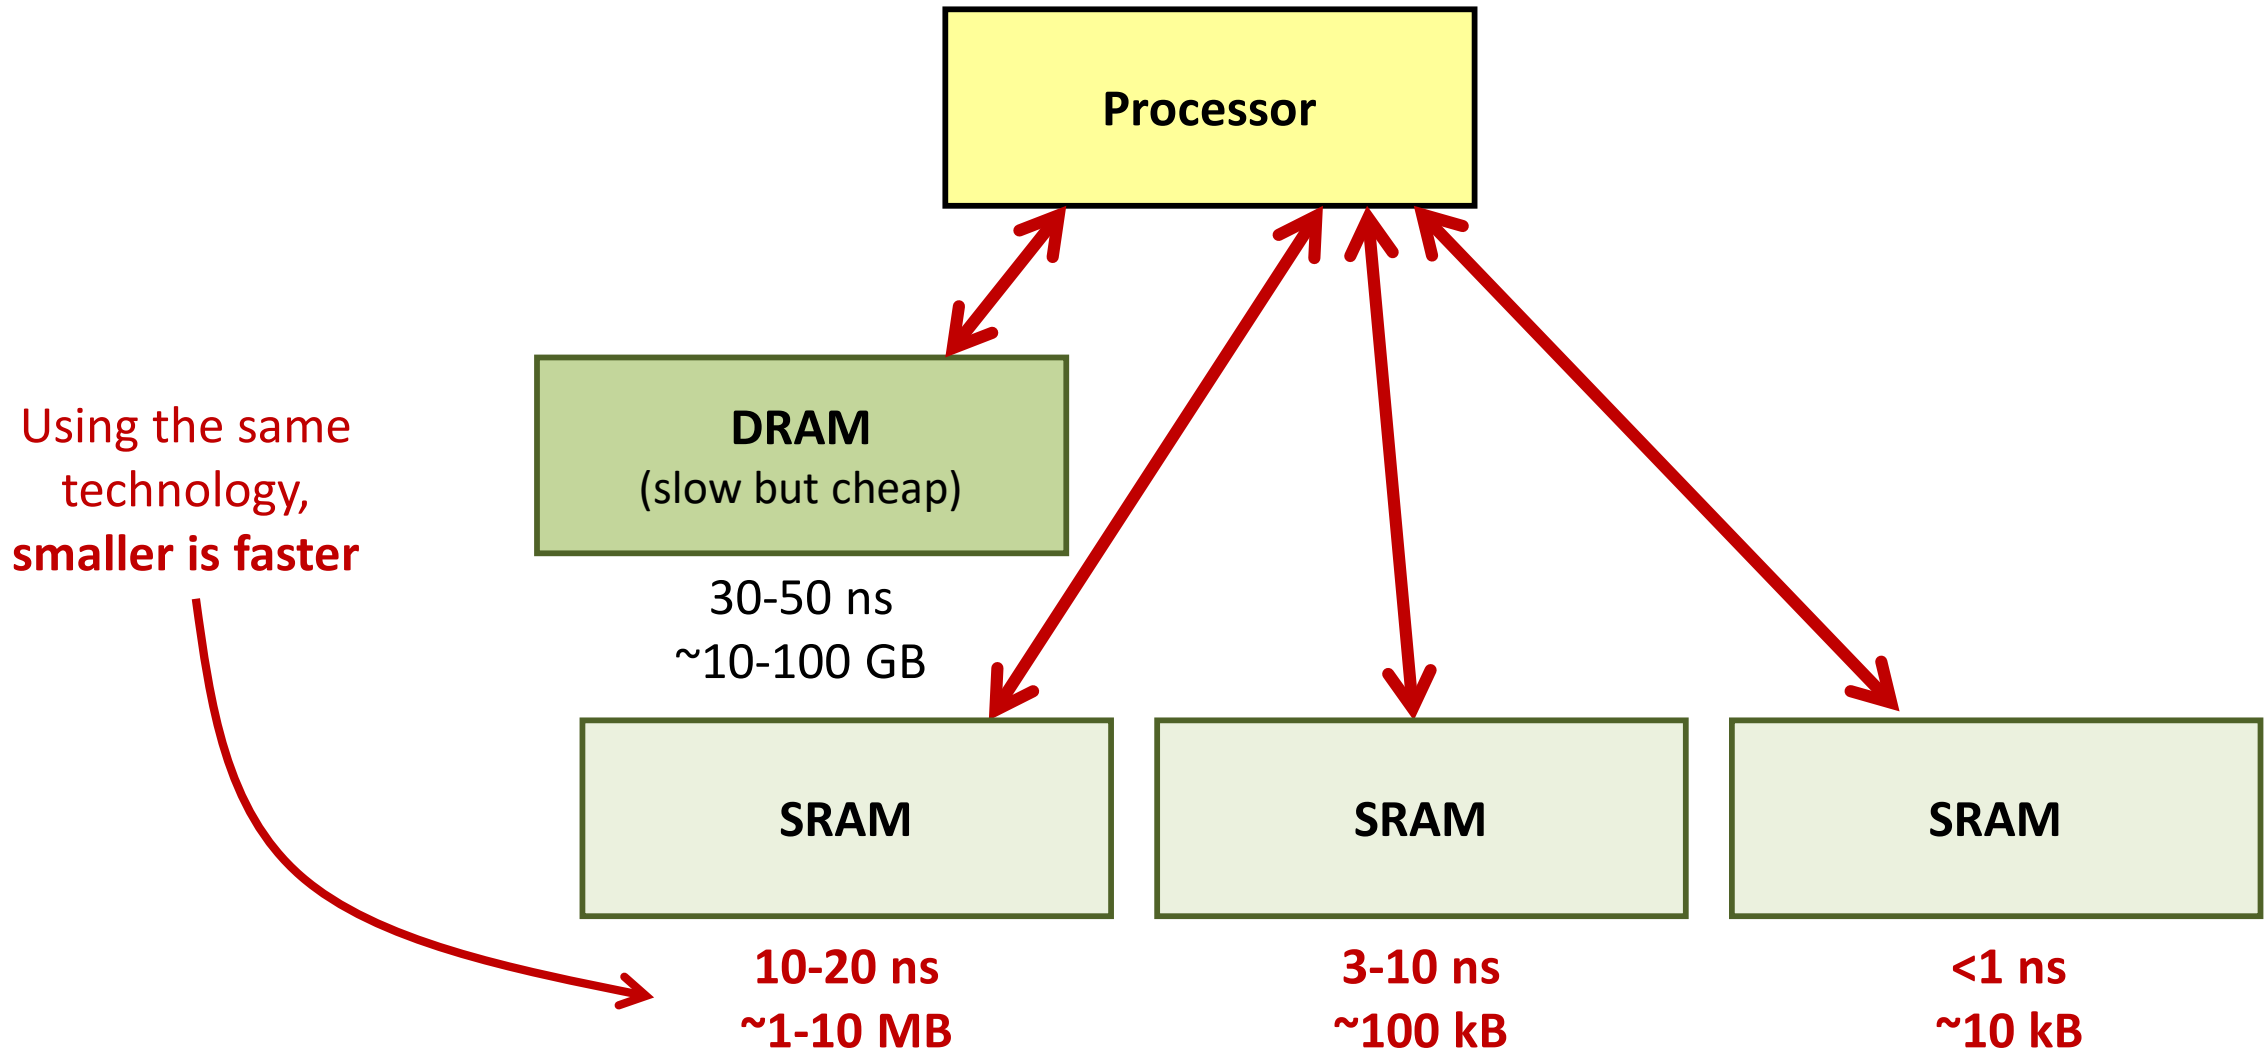
This is your CPU  
**You want it fast!**

# Our Goal Today: Use Different Memories

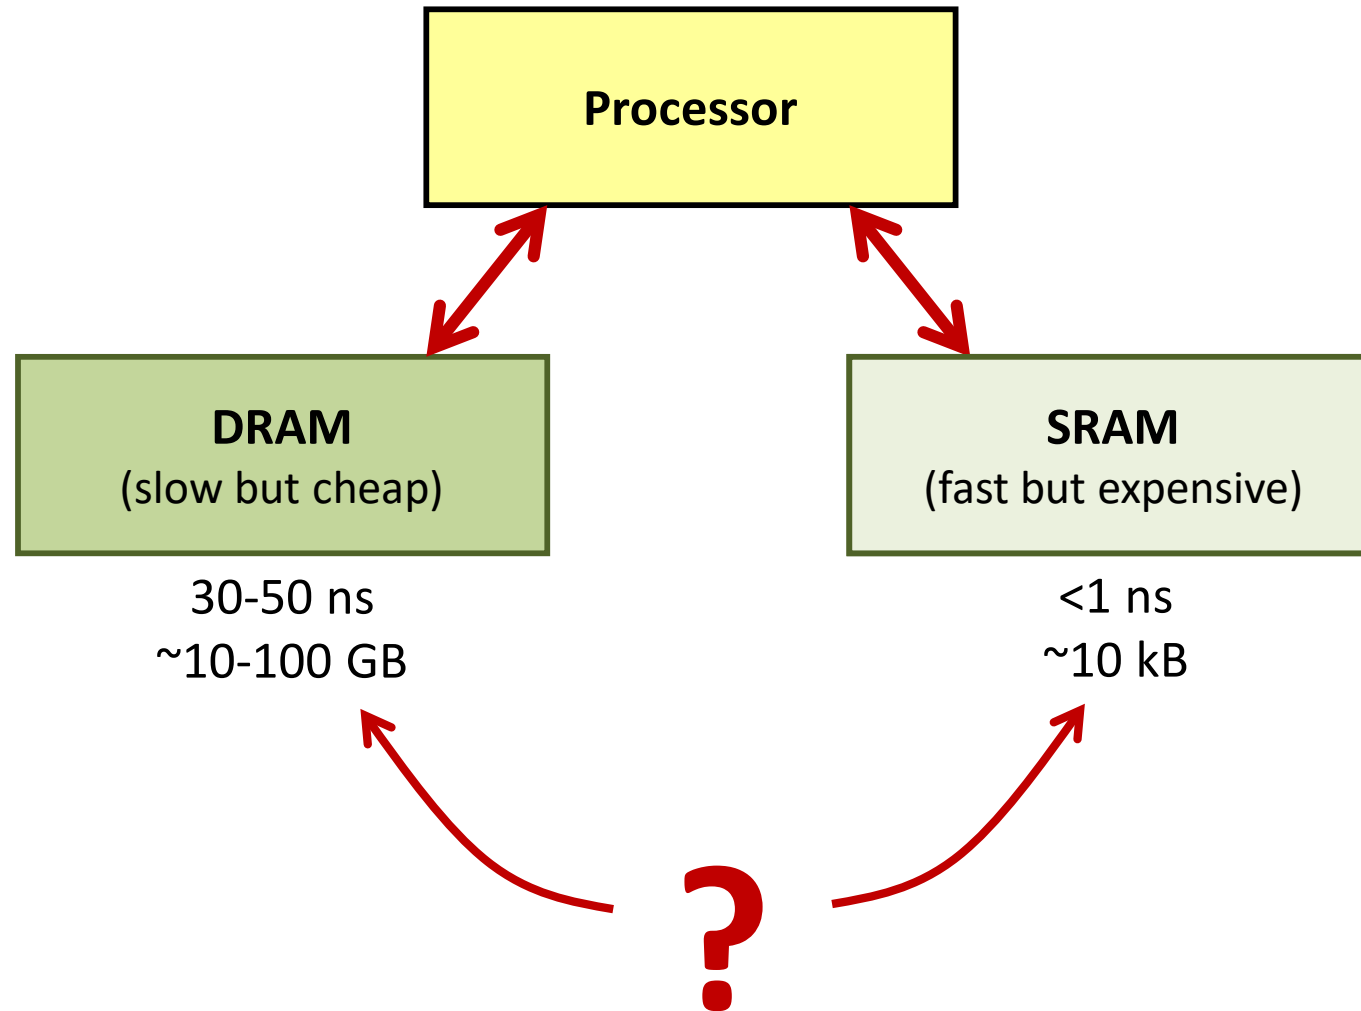


Can we get the best of both worlds?!

# Our Goal Today: Use Different Memories



# Where to Put What—and How?



# What Memory to Use?

- Instructions corresponding to lines 3-5 are **read over and over**: should be in fast memory
- If variables **i** and **sum** are stored in memory, they are also **used often** and should be stored in fast memory
- One would like to anticipate the future and load **following** instructions and vector elements

```
1:    i = 0;  
2:    sum = 0;  
3:    while (i < 1024) {  
4:        sum = sum + a[i];  
5:        i = i + 1;  
6:    }
```



# Spatial and Temporal Locality

Two important criteria to decide on placement:

- **Temporal Locality**

- Data that have been **used recently**, have high likelihood of being used again
  - Code: loops, functions,...
  - Data: local variables and data structures

- **Spatial Locality**

- Data which **follow in the memory other data** that are currently being used are likely to be used in the future
  - Code: usually read sequentially
  - Data: arrays

# Our Placement Policy Must Be...

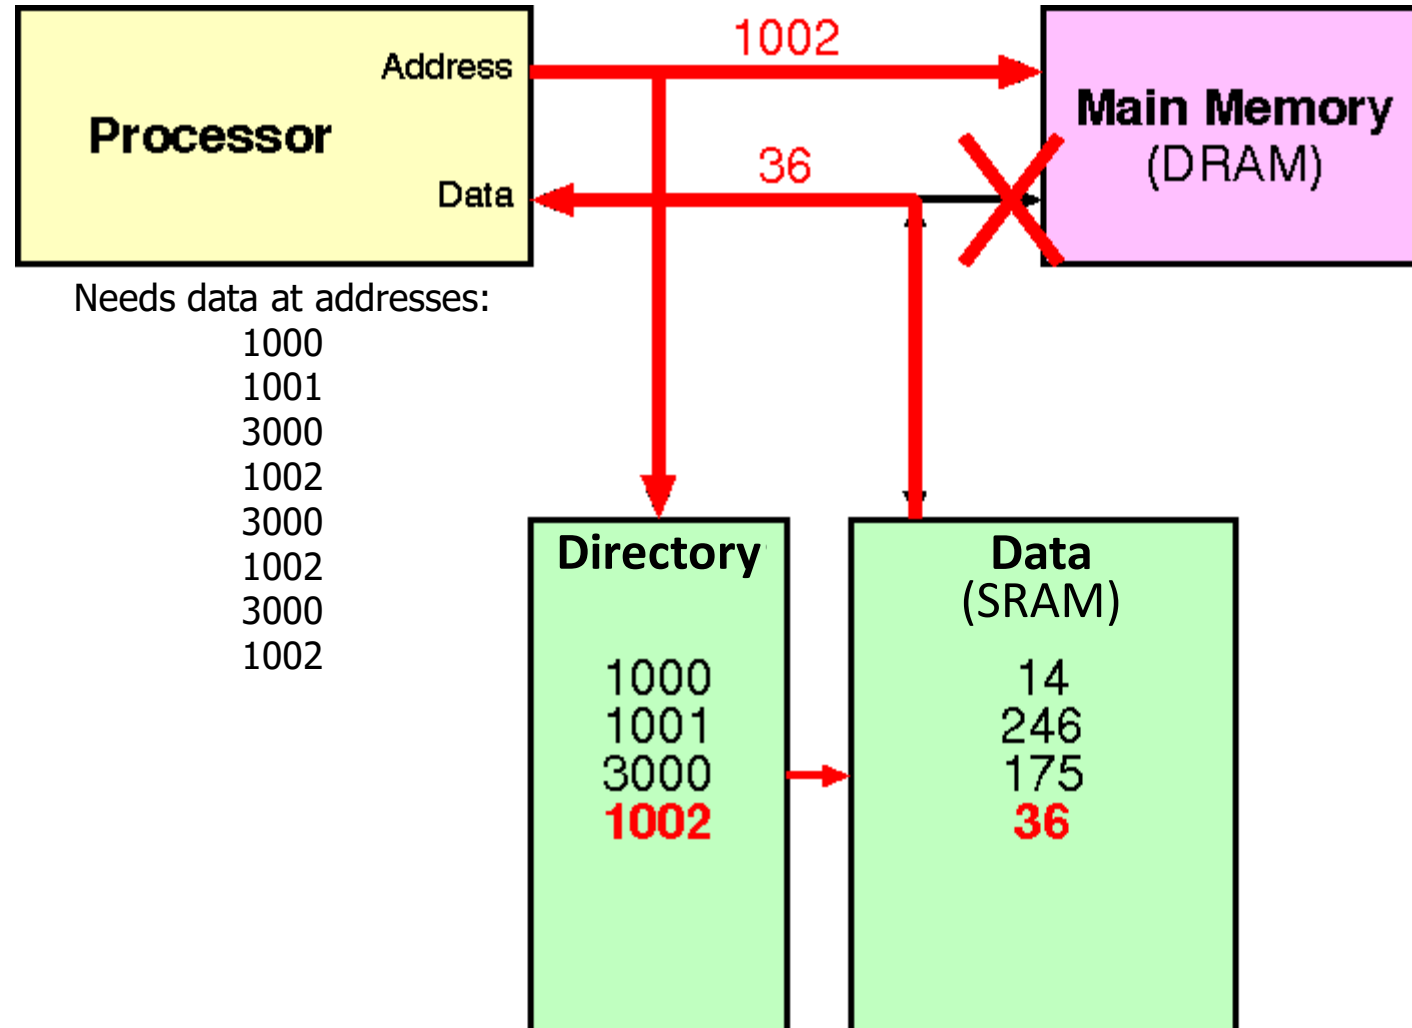
- **Invisible to the programmer**

- One could analyse data structures and program semantics to detect heavily used variables/arrays and thus decide placement → Ok in some contexts (embedded) but we want to have the programmers not go through this hassle
- We will add **hardware** to help

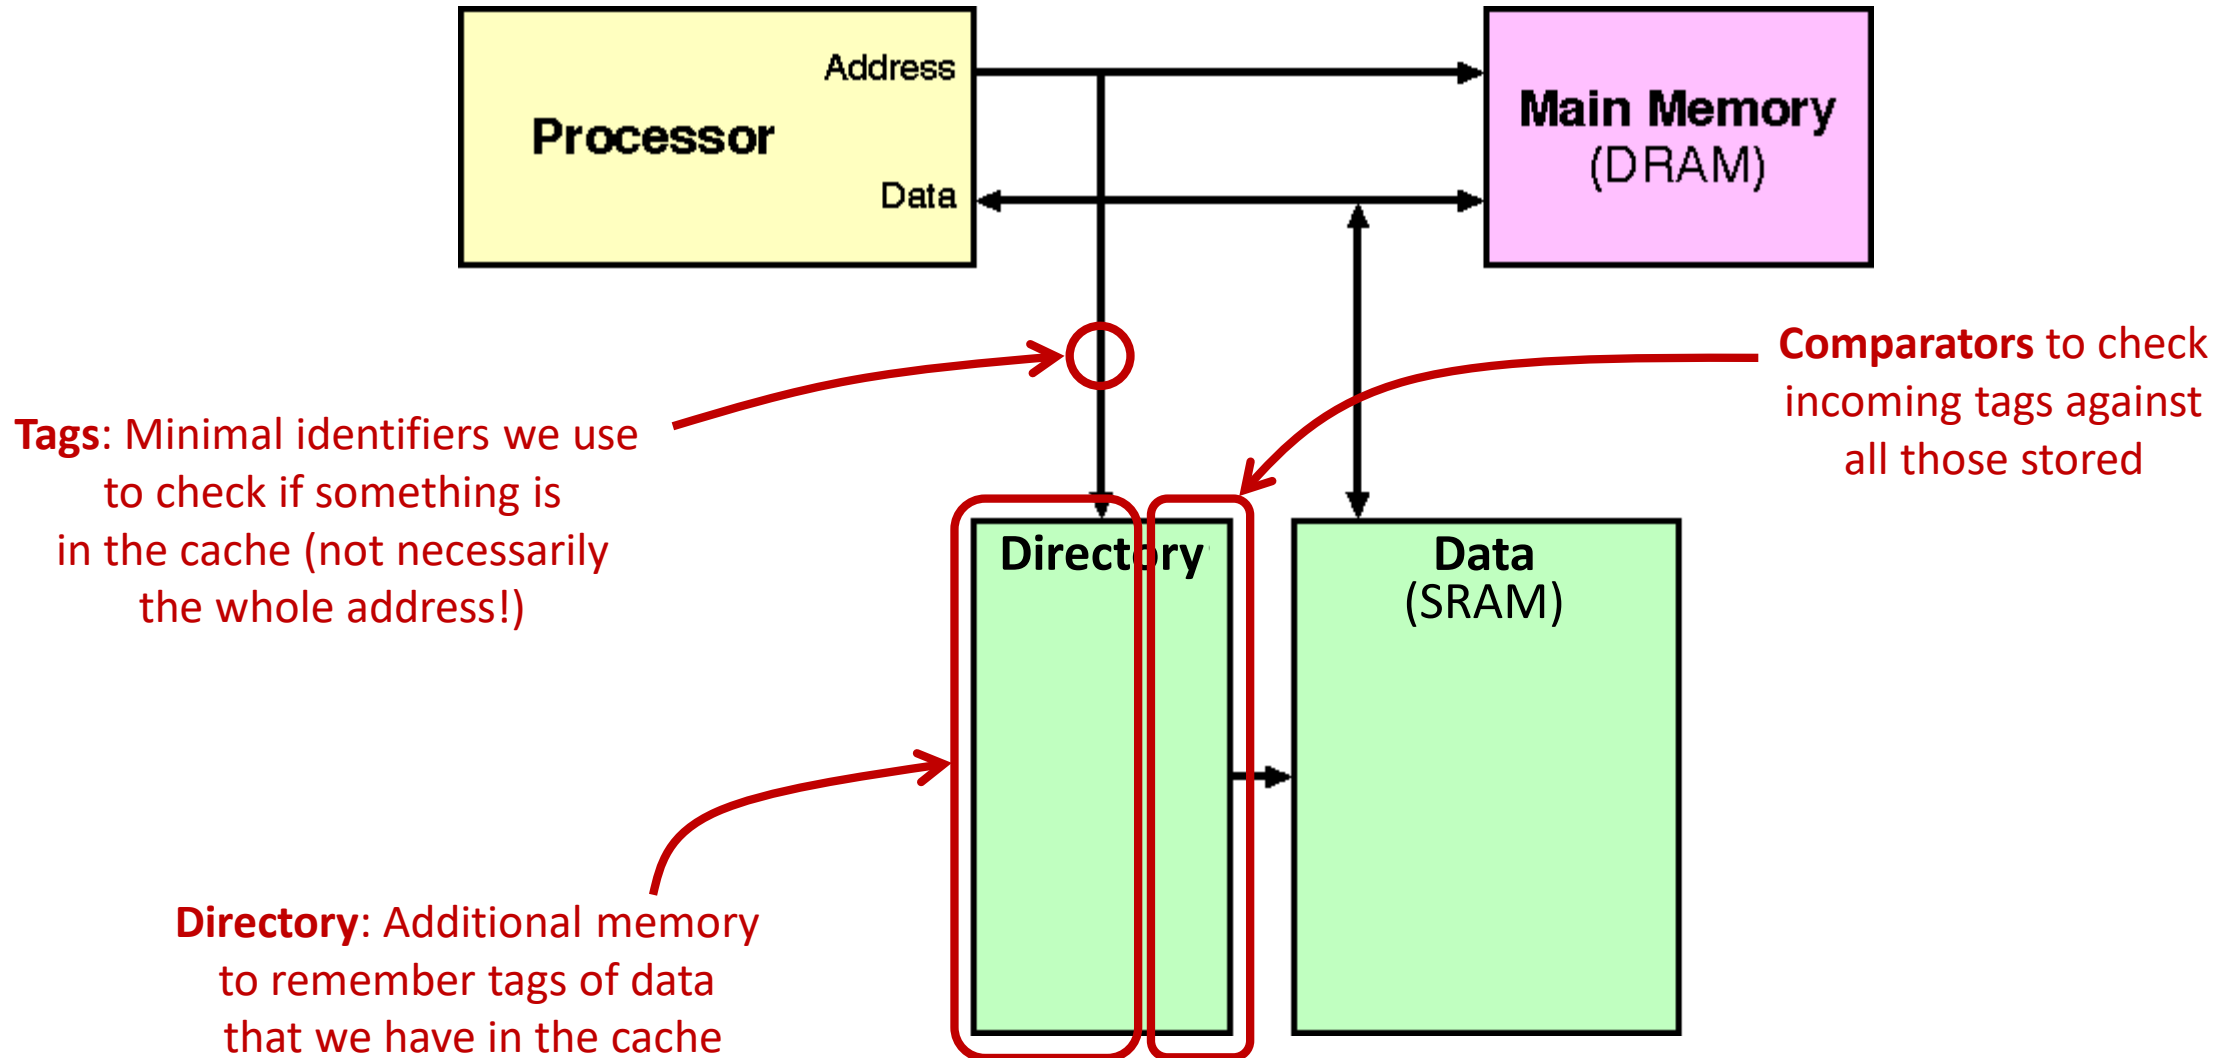
- **Extremely simple and fast**

- If decisions are to be made in hardware, they need to be simple
- Goal is to access fast memory in the order of a ns or less: **not much time** to make complex decisions...

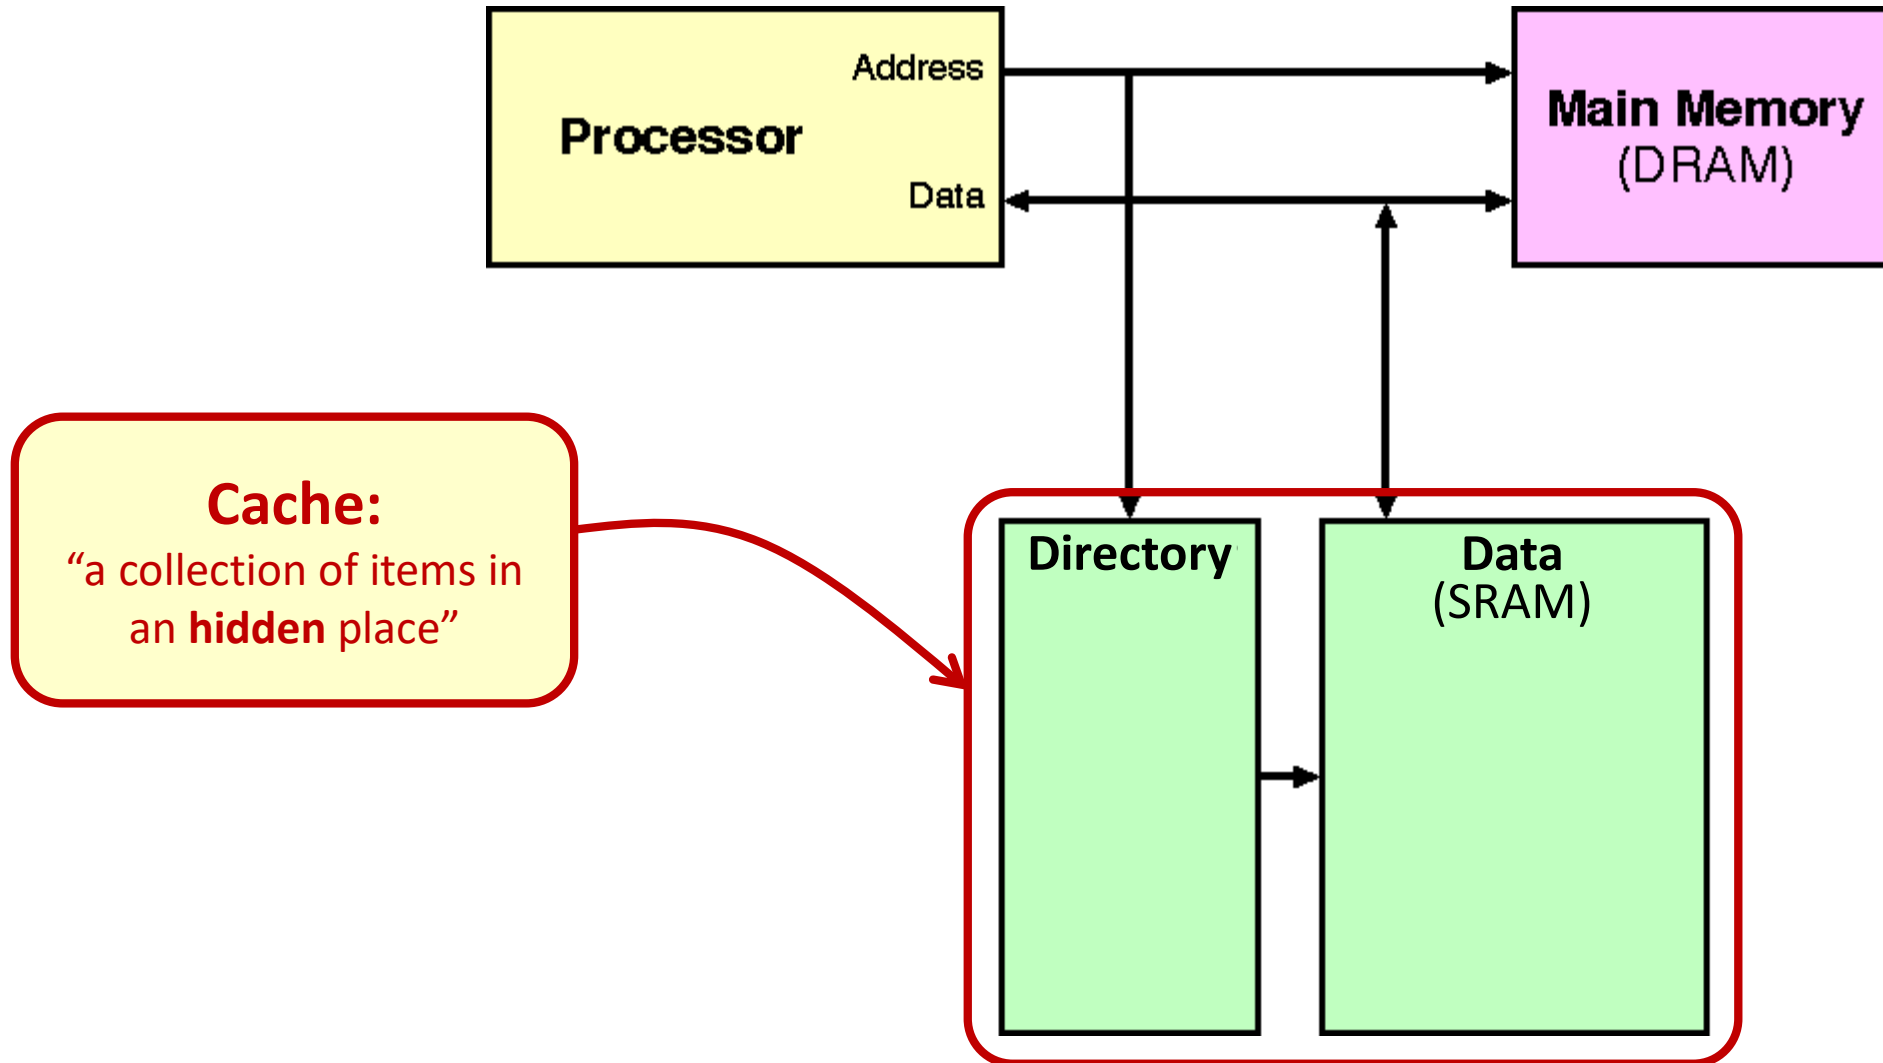
# Cache: The Idea



# Not Just Fast Memory: Directory and Tags



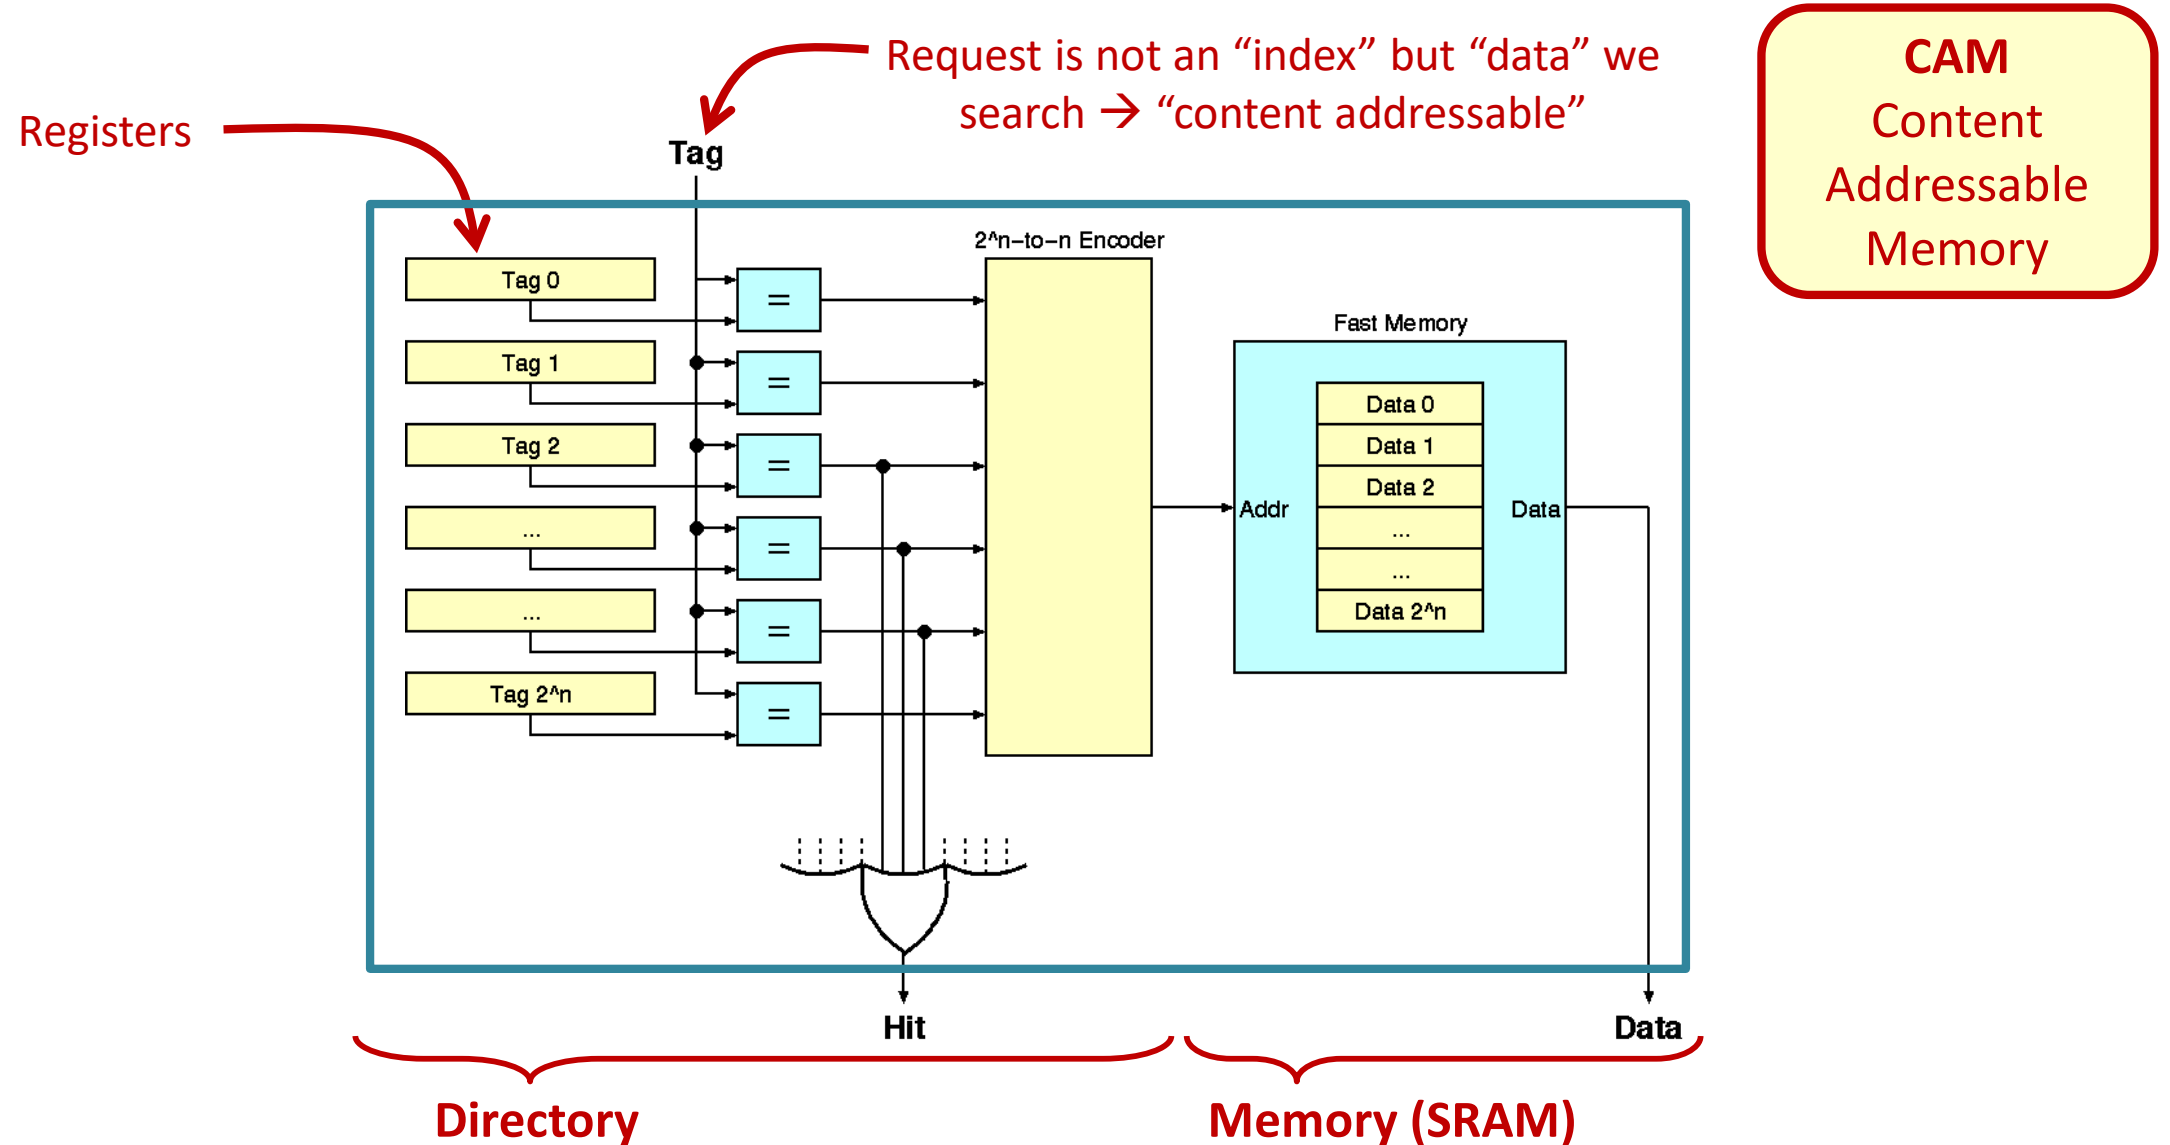
# A Cache!



# Cache Hits and Misses


- A **cache** is any form of storage which takes **automatically** advantage of locality of accesses
  - The idea works so well that now they are **not only in processors!**
  - Web browsers have caches, network routers have routing information and even data caches, DNSs cache frequent names, databases cache queries...
- When we find the data required in the cache, we call it a **Hit**; otherwise it is a **Miss**
- **Hit (or Miss) Rate** is the number of hits (or misses) over the total number of accesses

# Fully-Associative Cache



# Fully-Associative Cache

The representation  
we will use



Tag


	=	
	=	
0x100	=	M[0x100]
	=	
	=	
	=	

A **line** or a **block** of the cache


Hit    Data



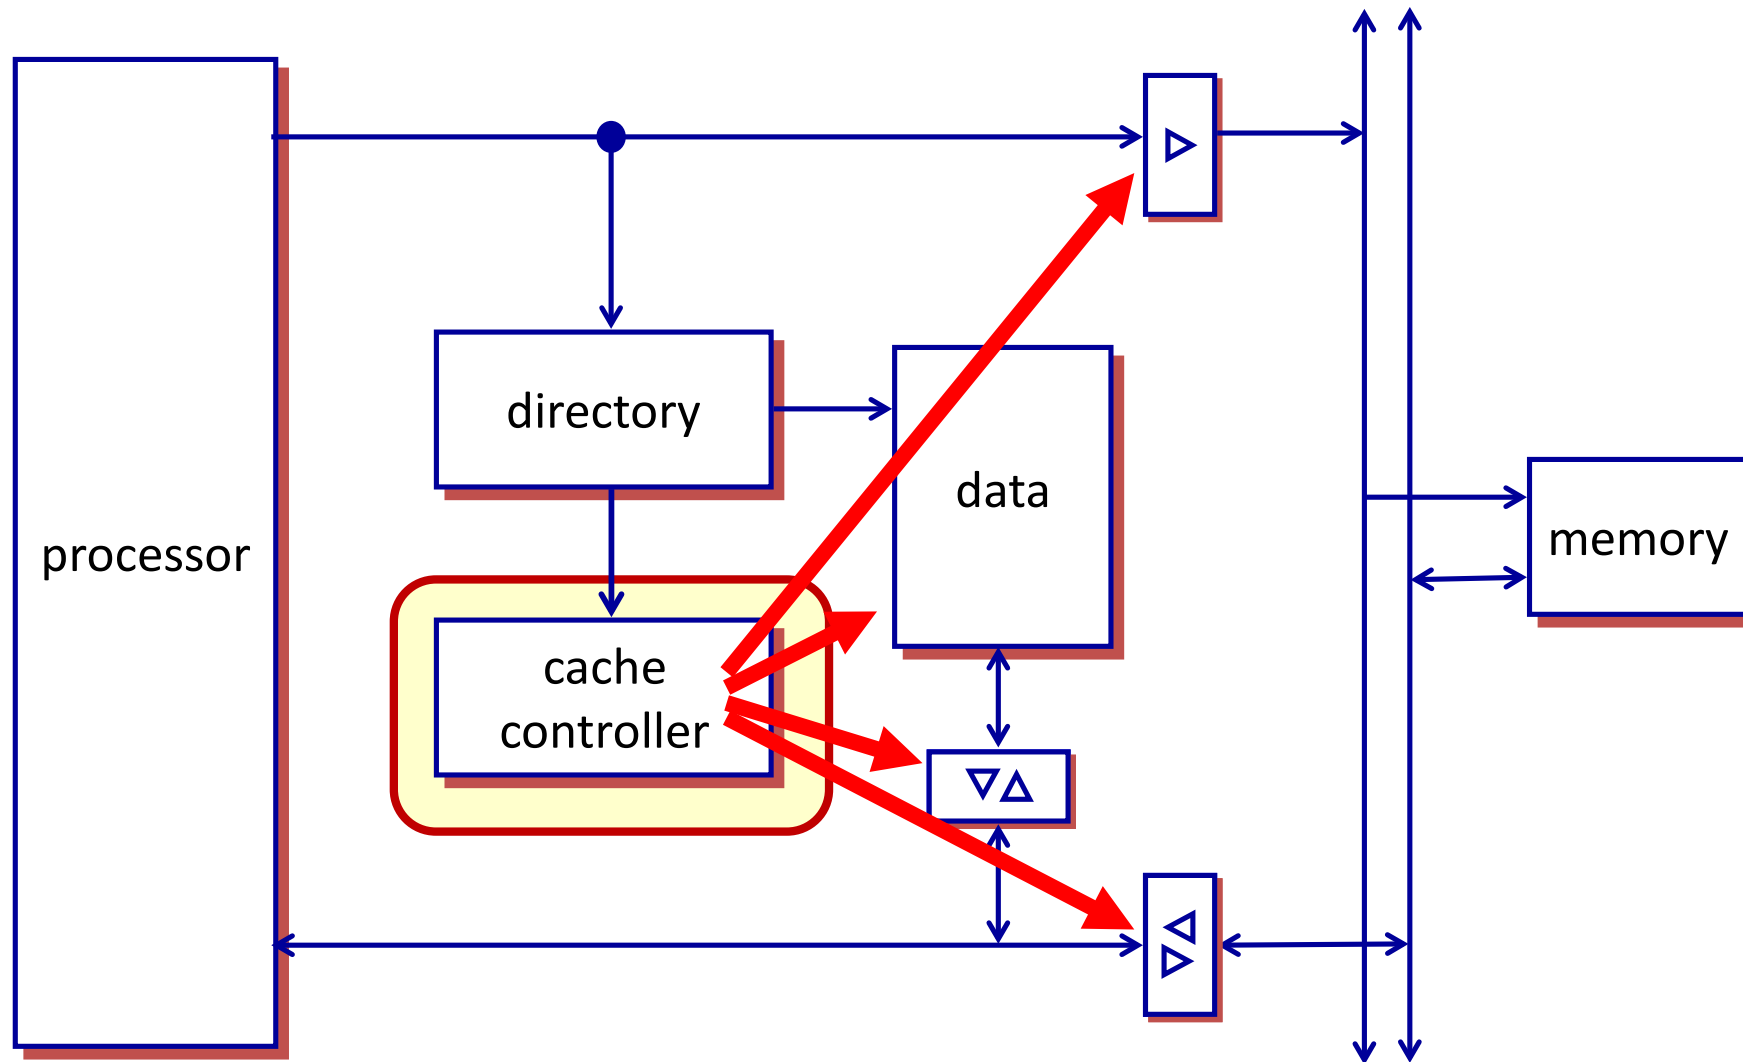
# Cache and Cache Controller



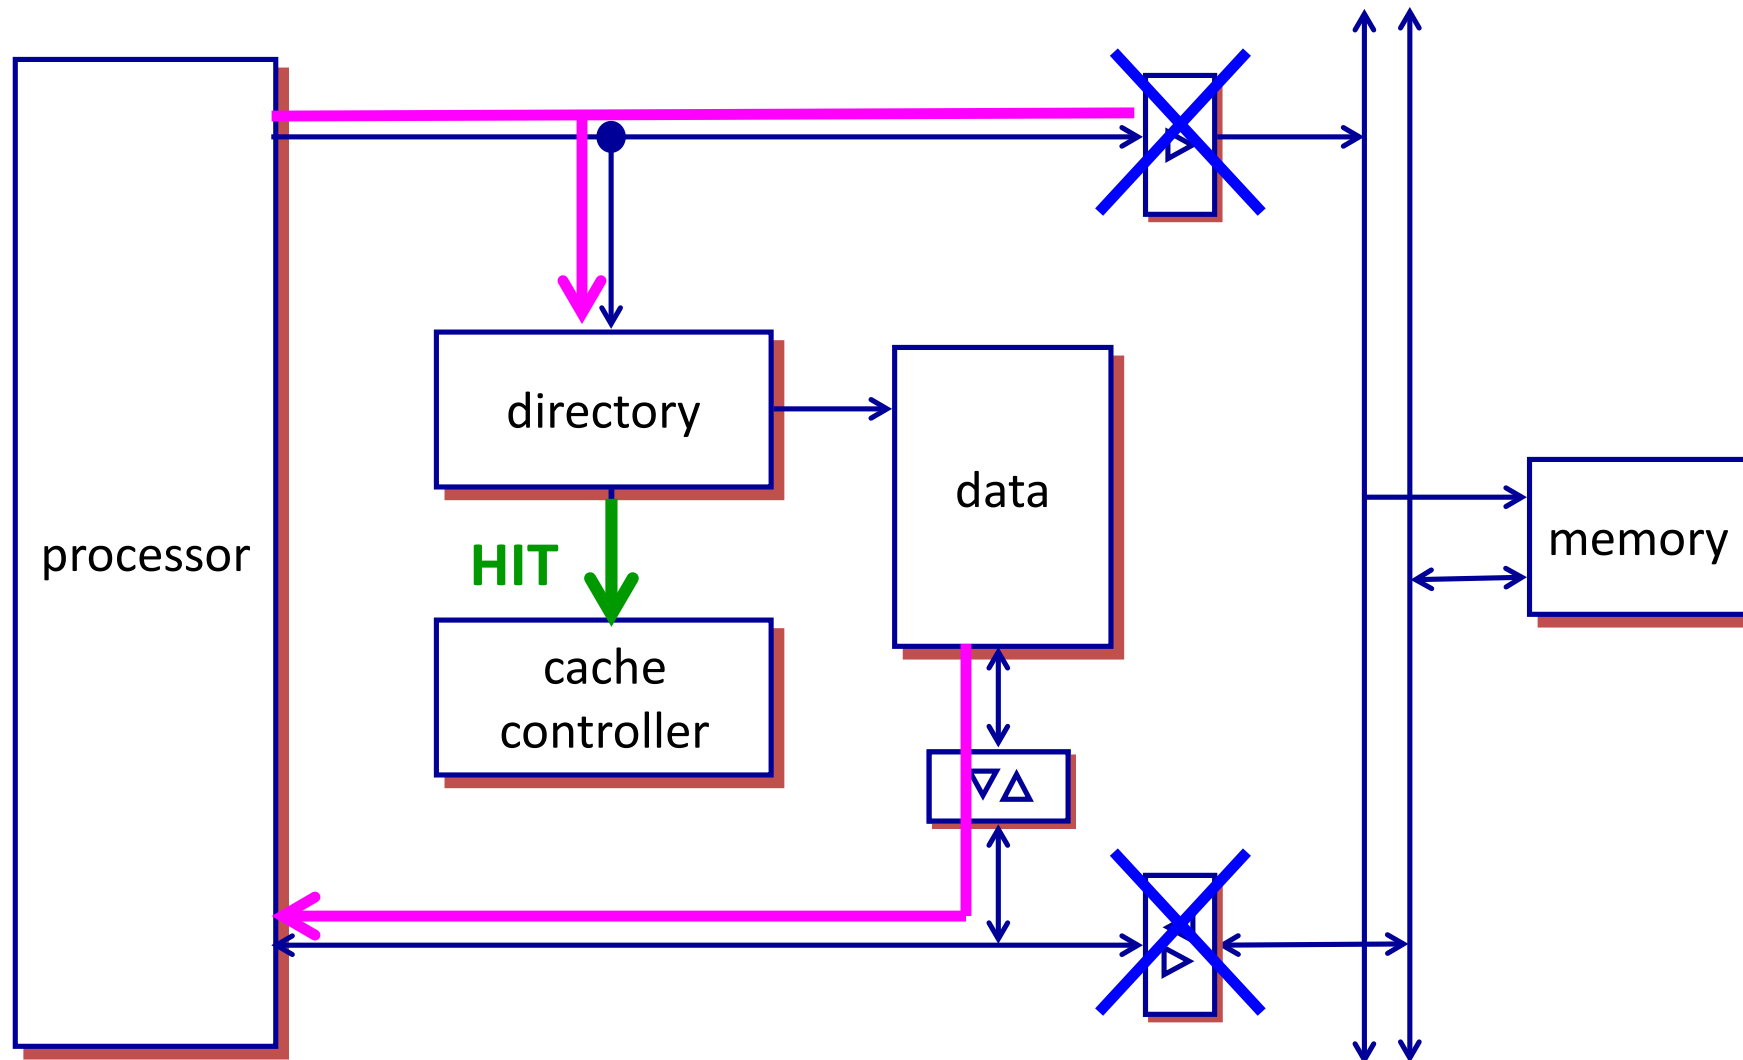
# Cache and Cache Controller



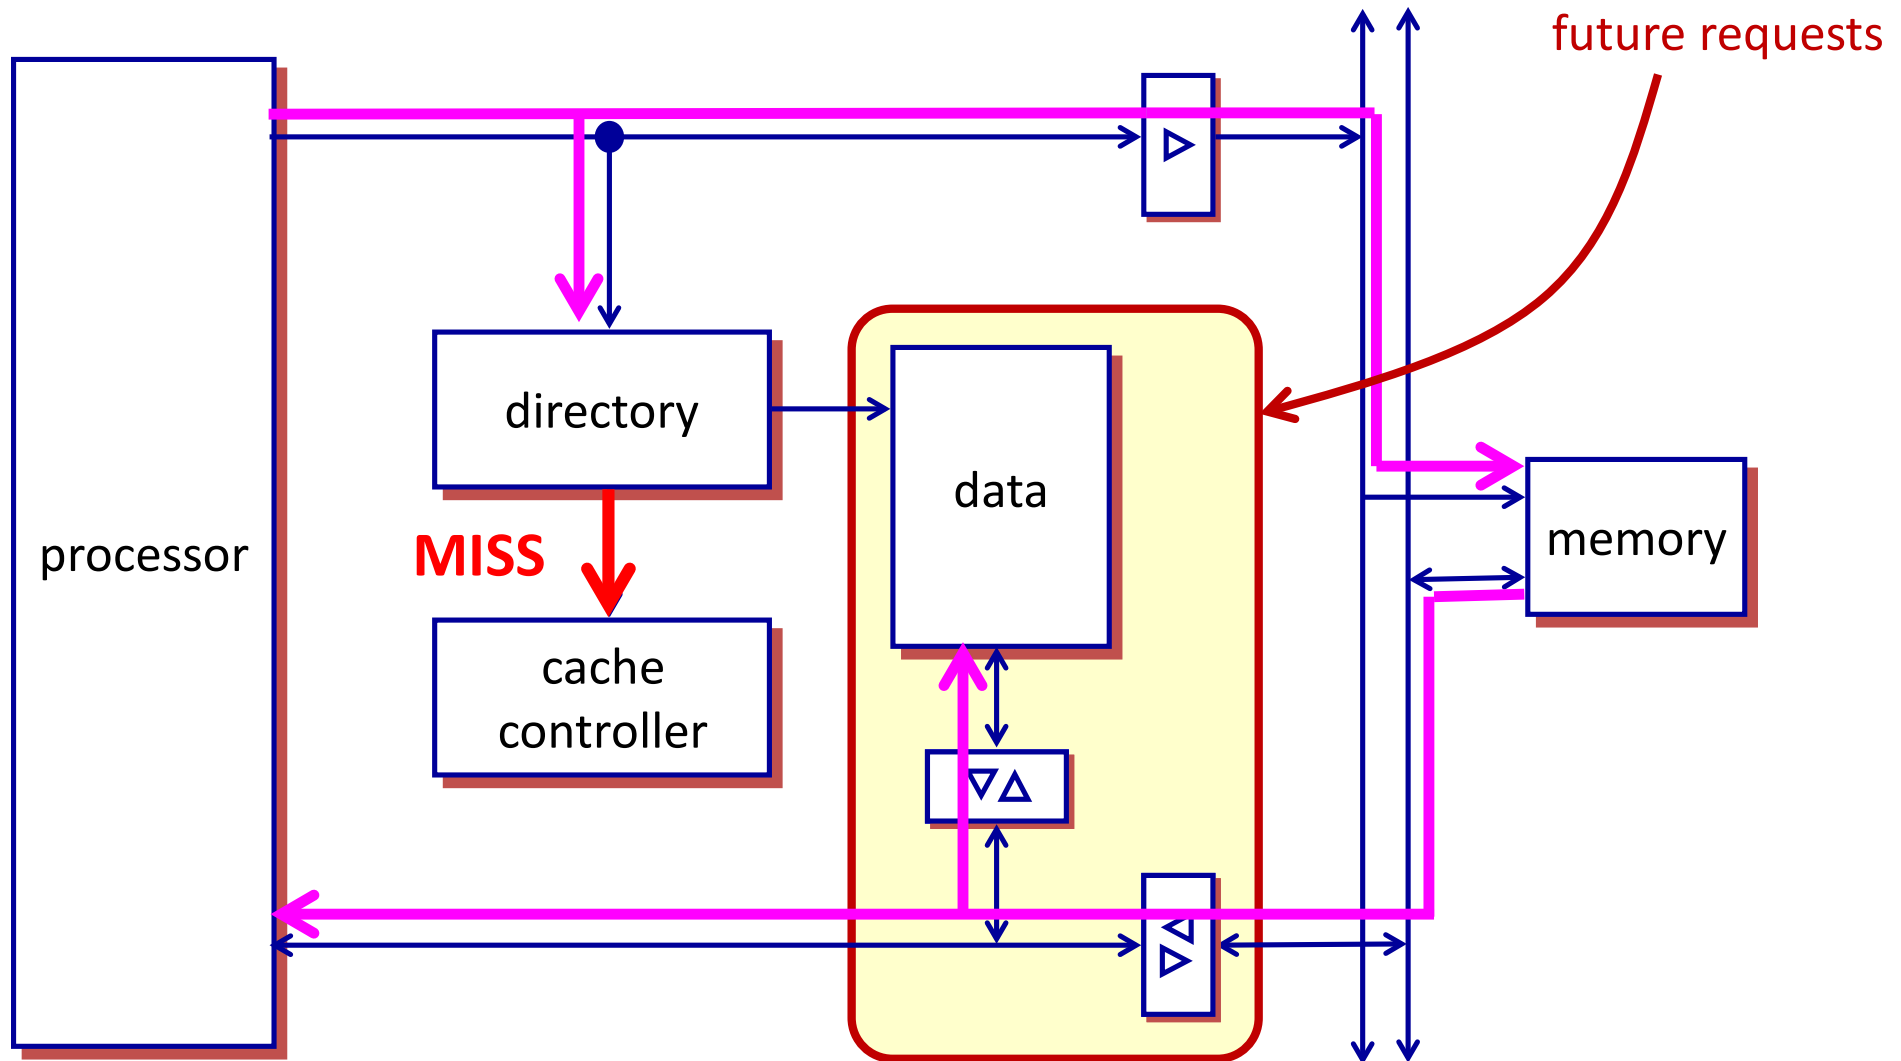
# Cache and Cache Controller




# Cache Hit



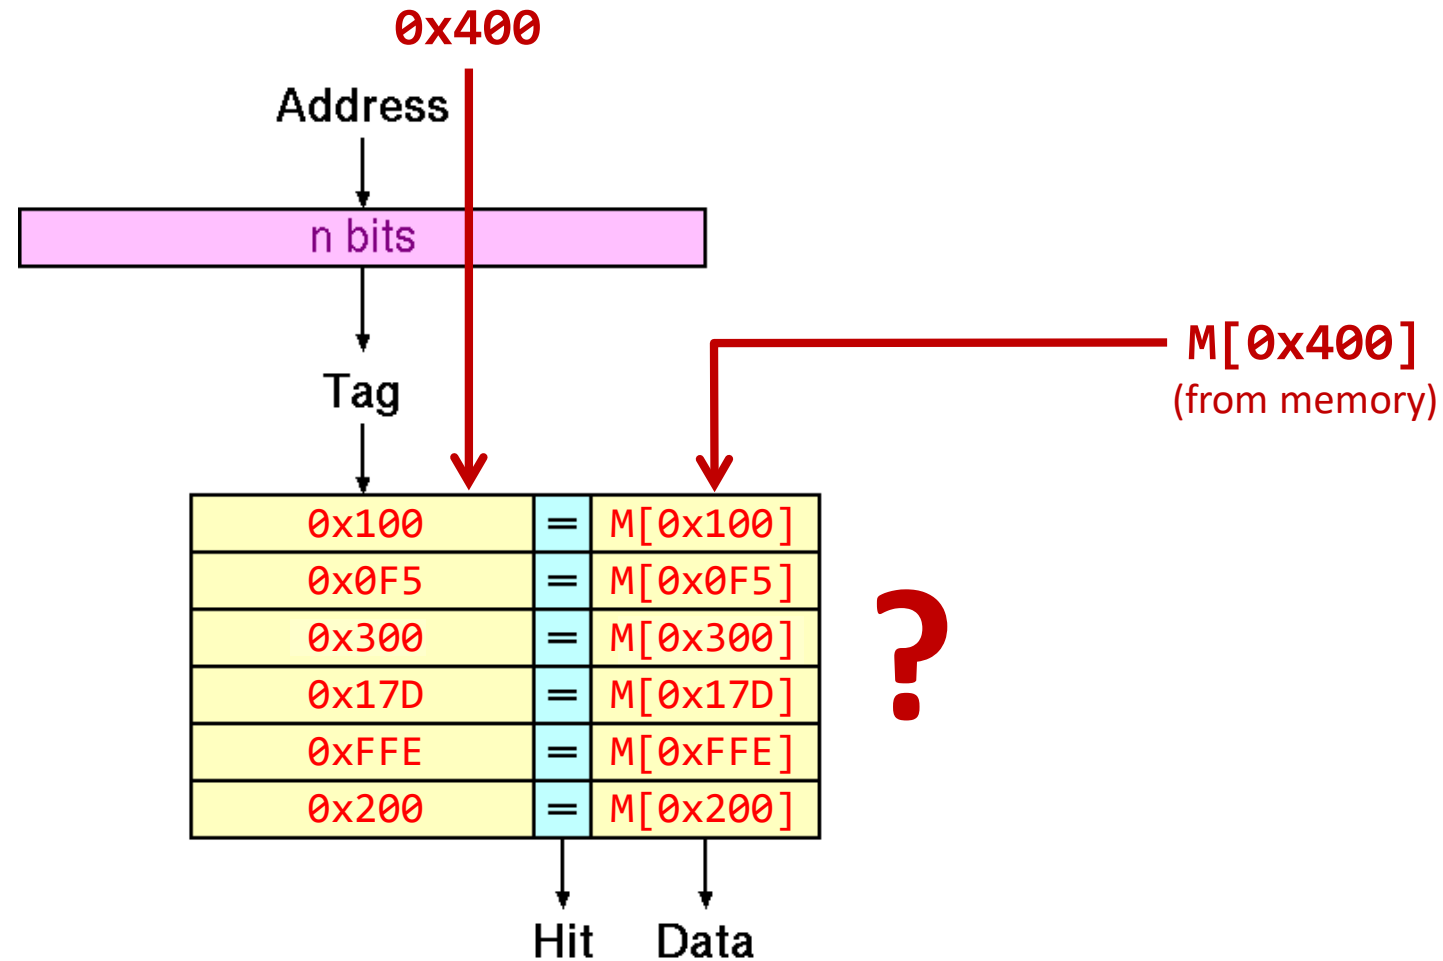
# Cache Miss



# Cache and Cache Controller



# What If the Cache Is Full?




# Eviction Policies

- When there is no appropriate space for a new piece of data, we must overwrite one of the existing lines (**eviction** or **replacement**)
- Several policies to decide what to evict:
  - **Least Recently Used** (LRU)
    - Replace the data that have been unused for the longest period of time
  - **First-In First-Out** (FIFO)
    - Replace the data that came in earliest
  - **Random**
    - Pick one, any one, and throw it away...
  - Approximate schemes, etc...



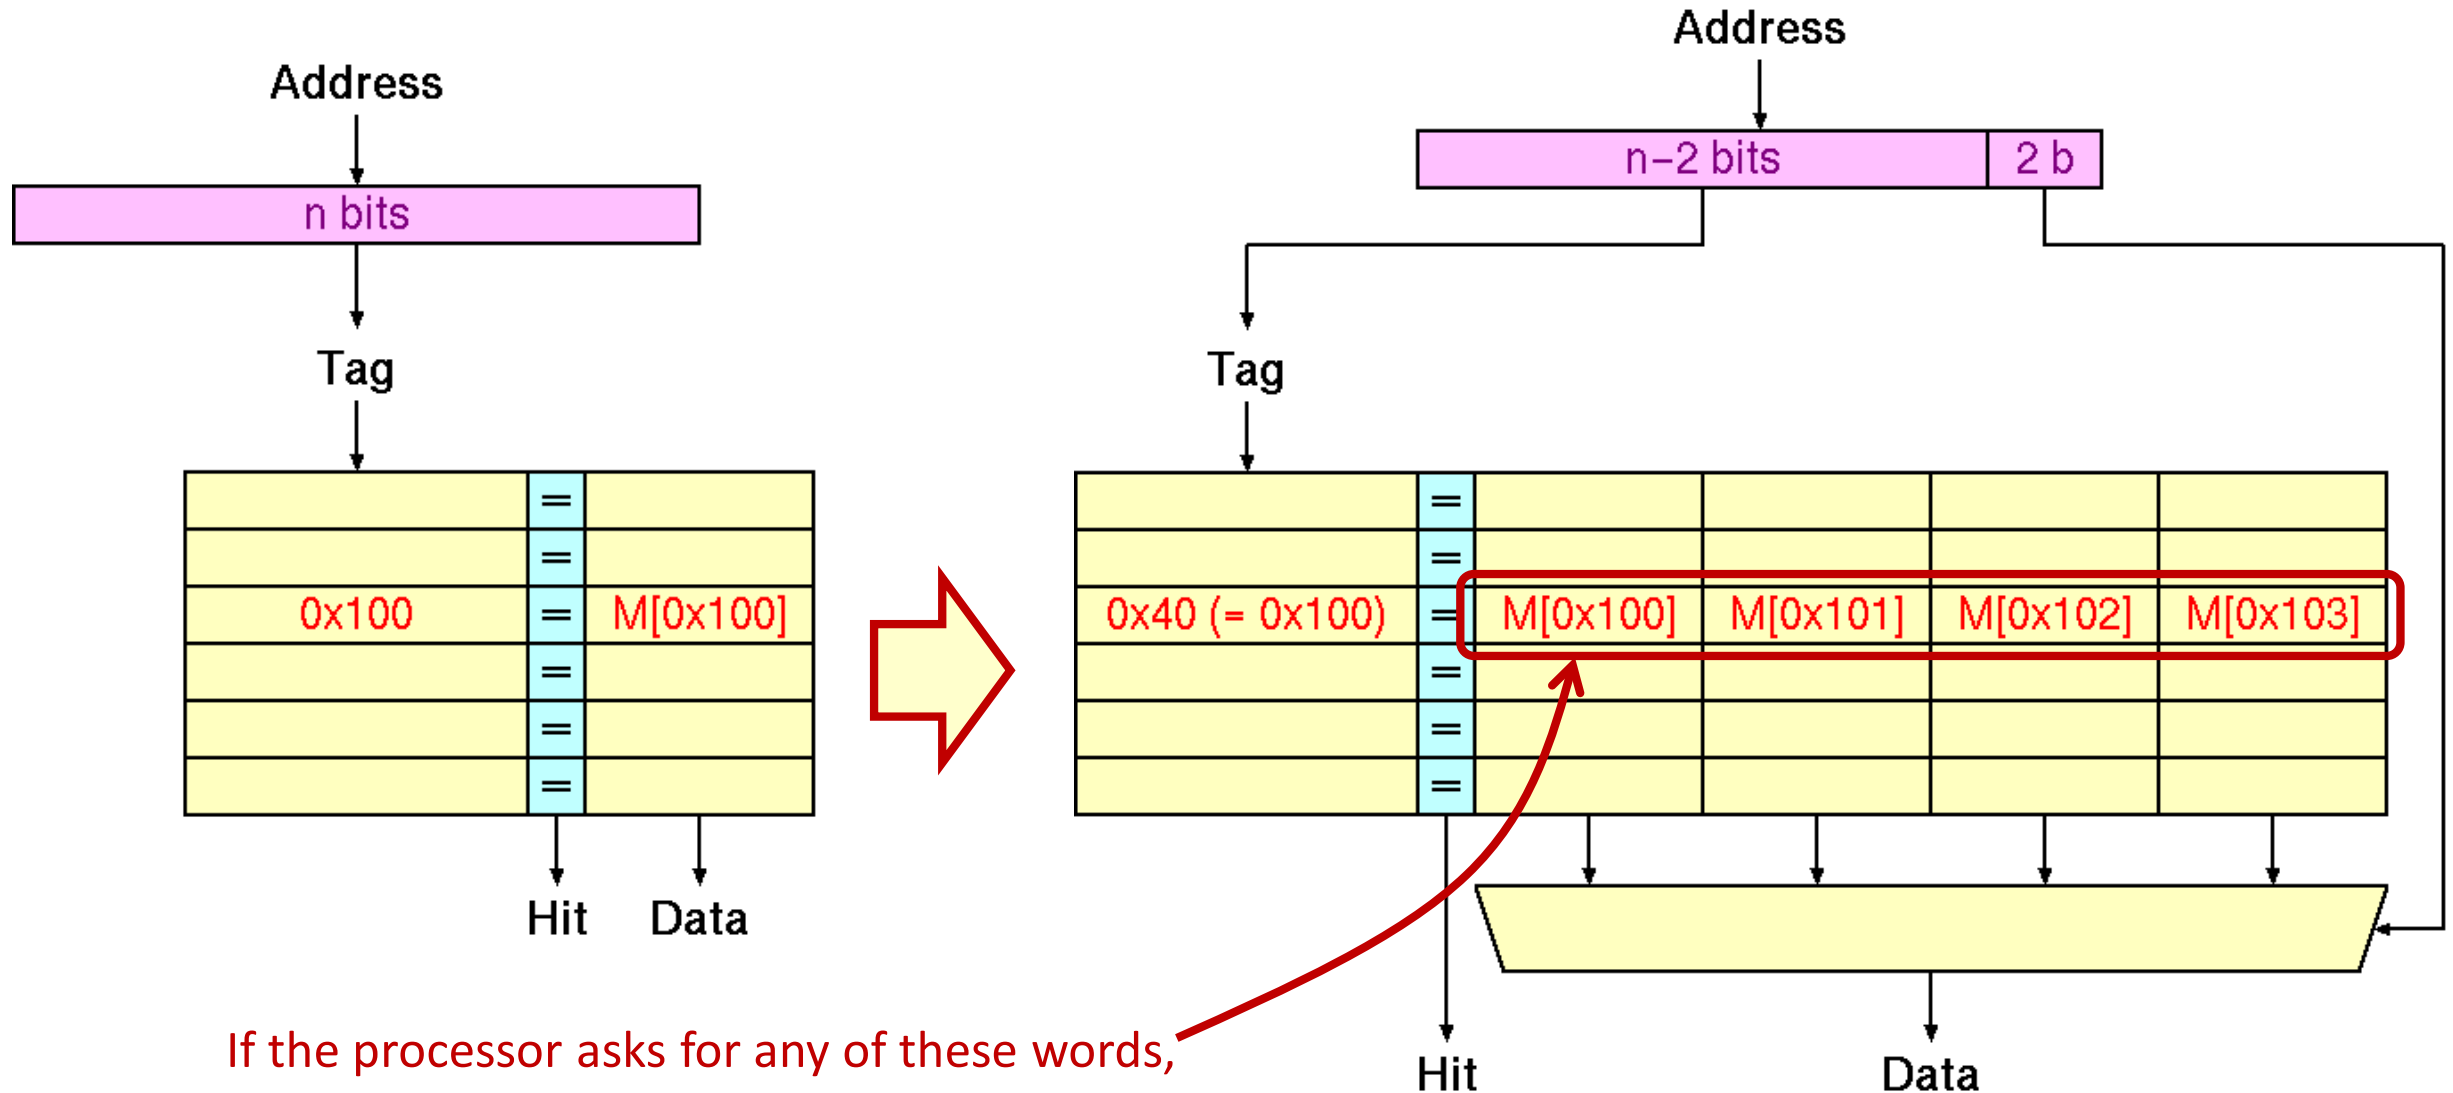
# Only Exploiting Temporal Locality



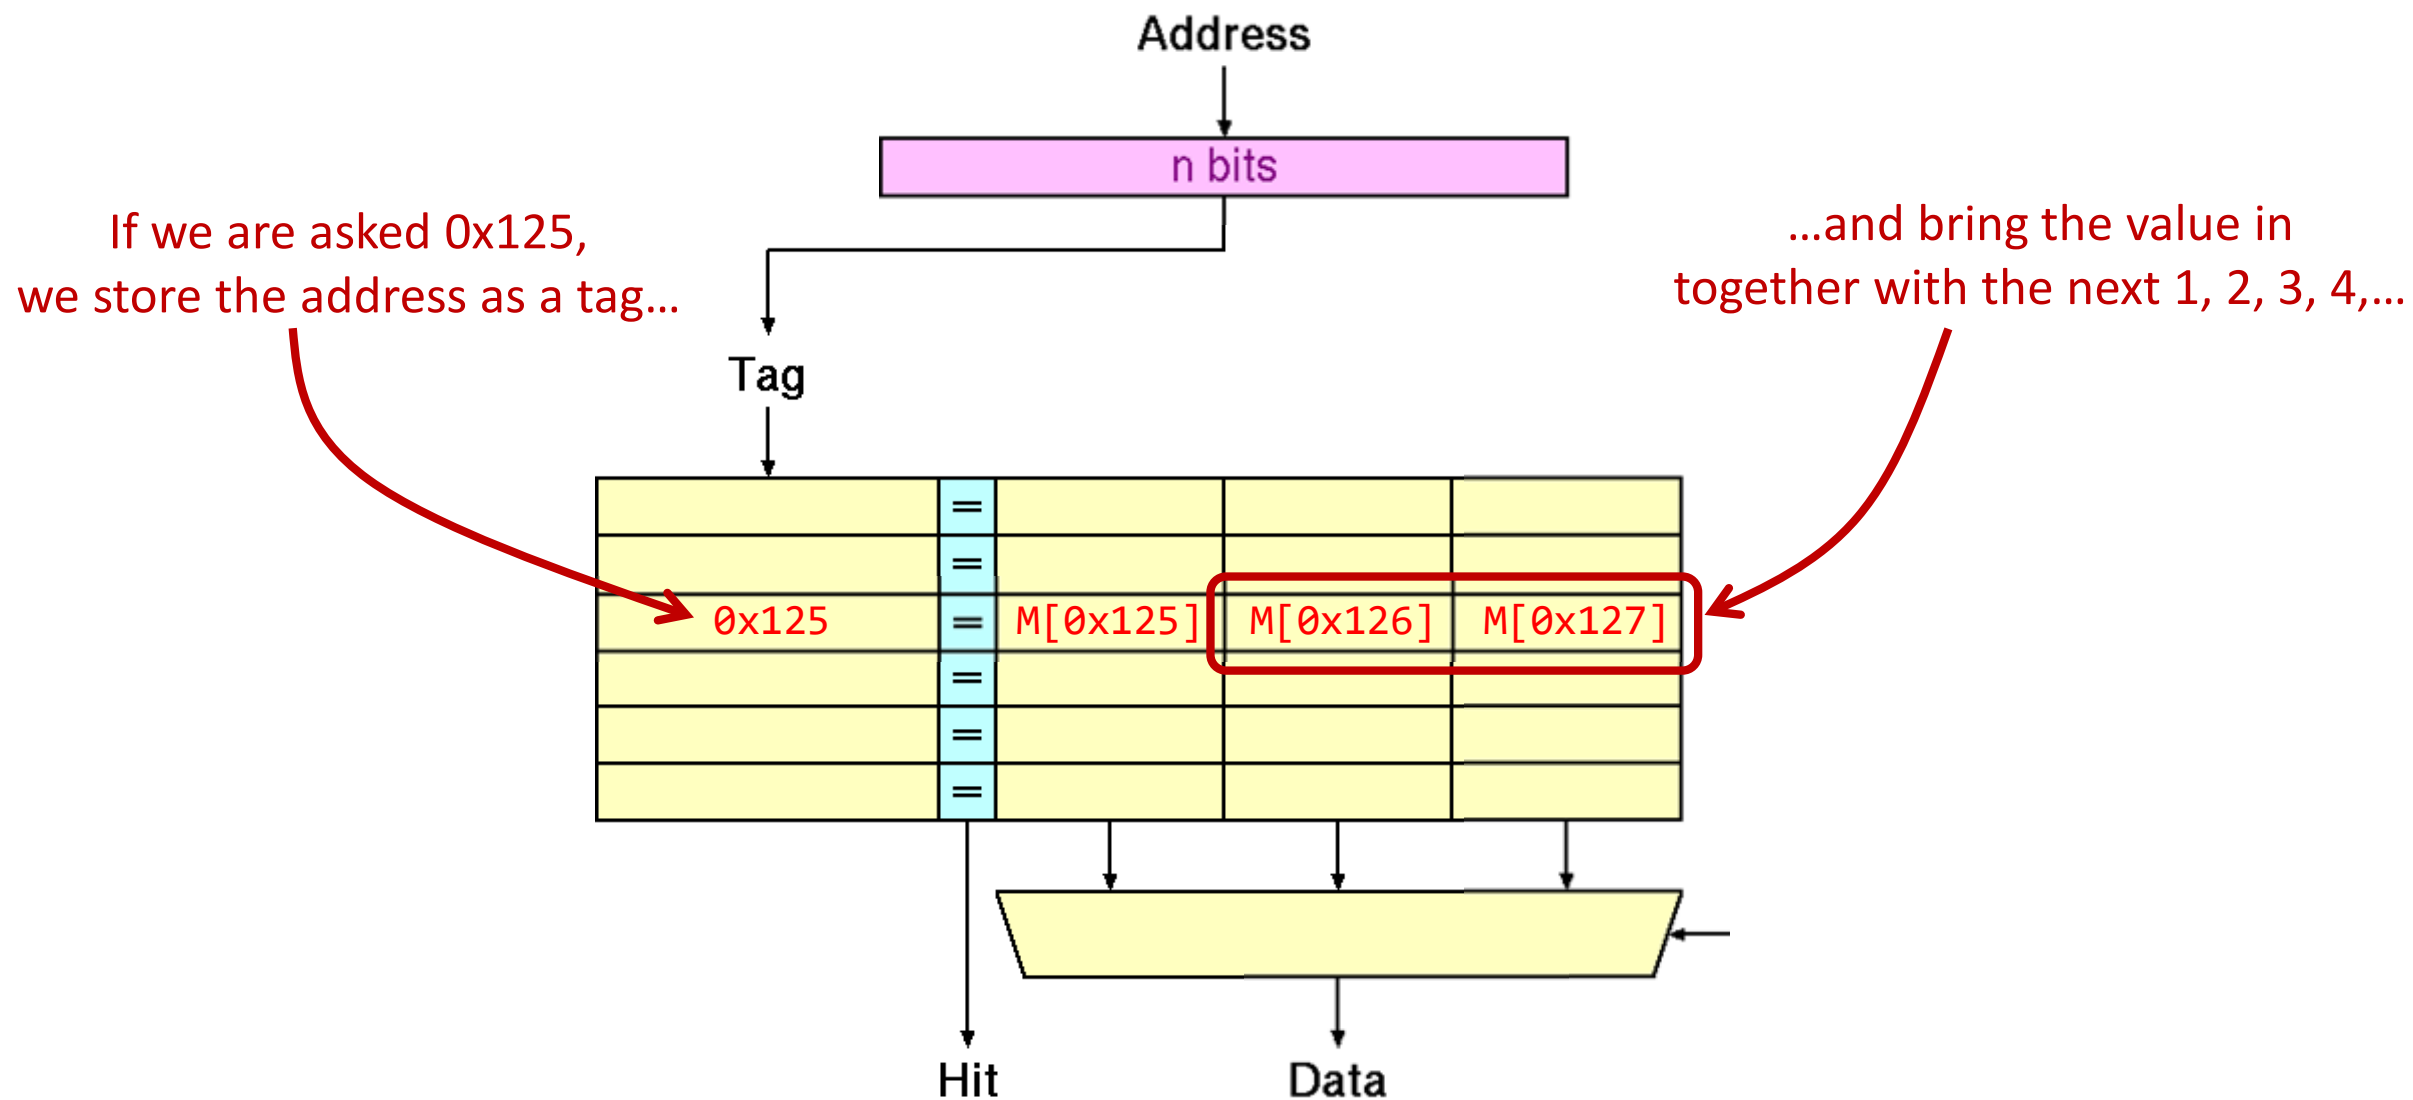
We only bring from main memory to the cache **exclusively the data required and when required**, hoping to use them again

If the processor asks for neighbouring items (e.g., M[0x101]), we will not have it; we are **not able to exploit any spatial locality**

# Exploiting Spatial Locality

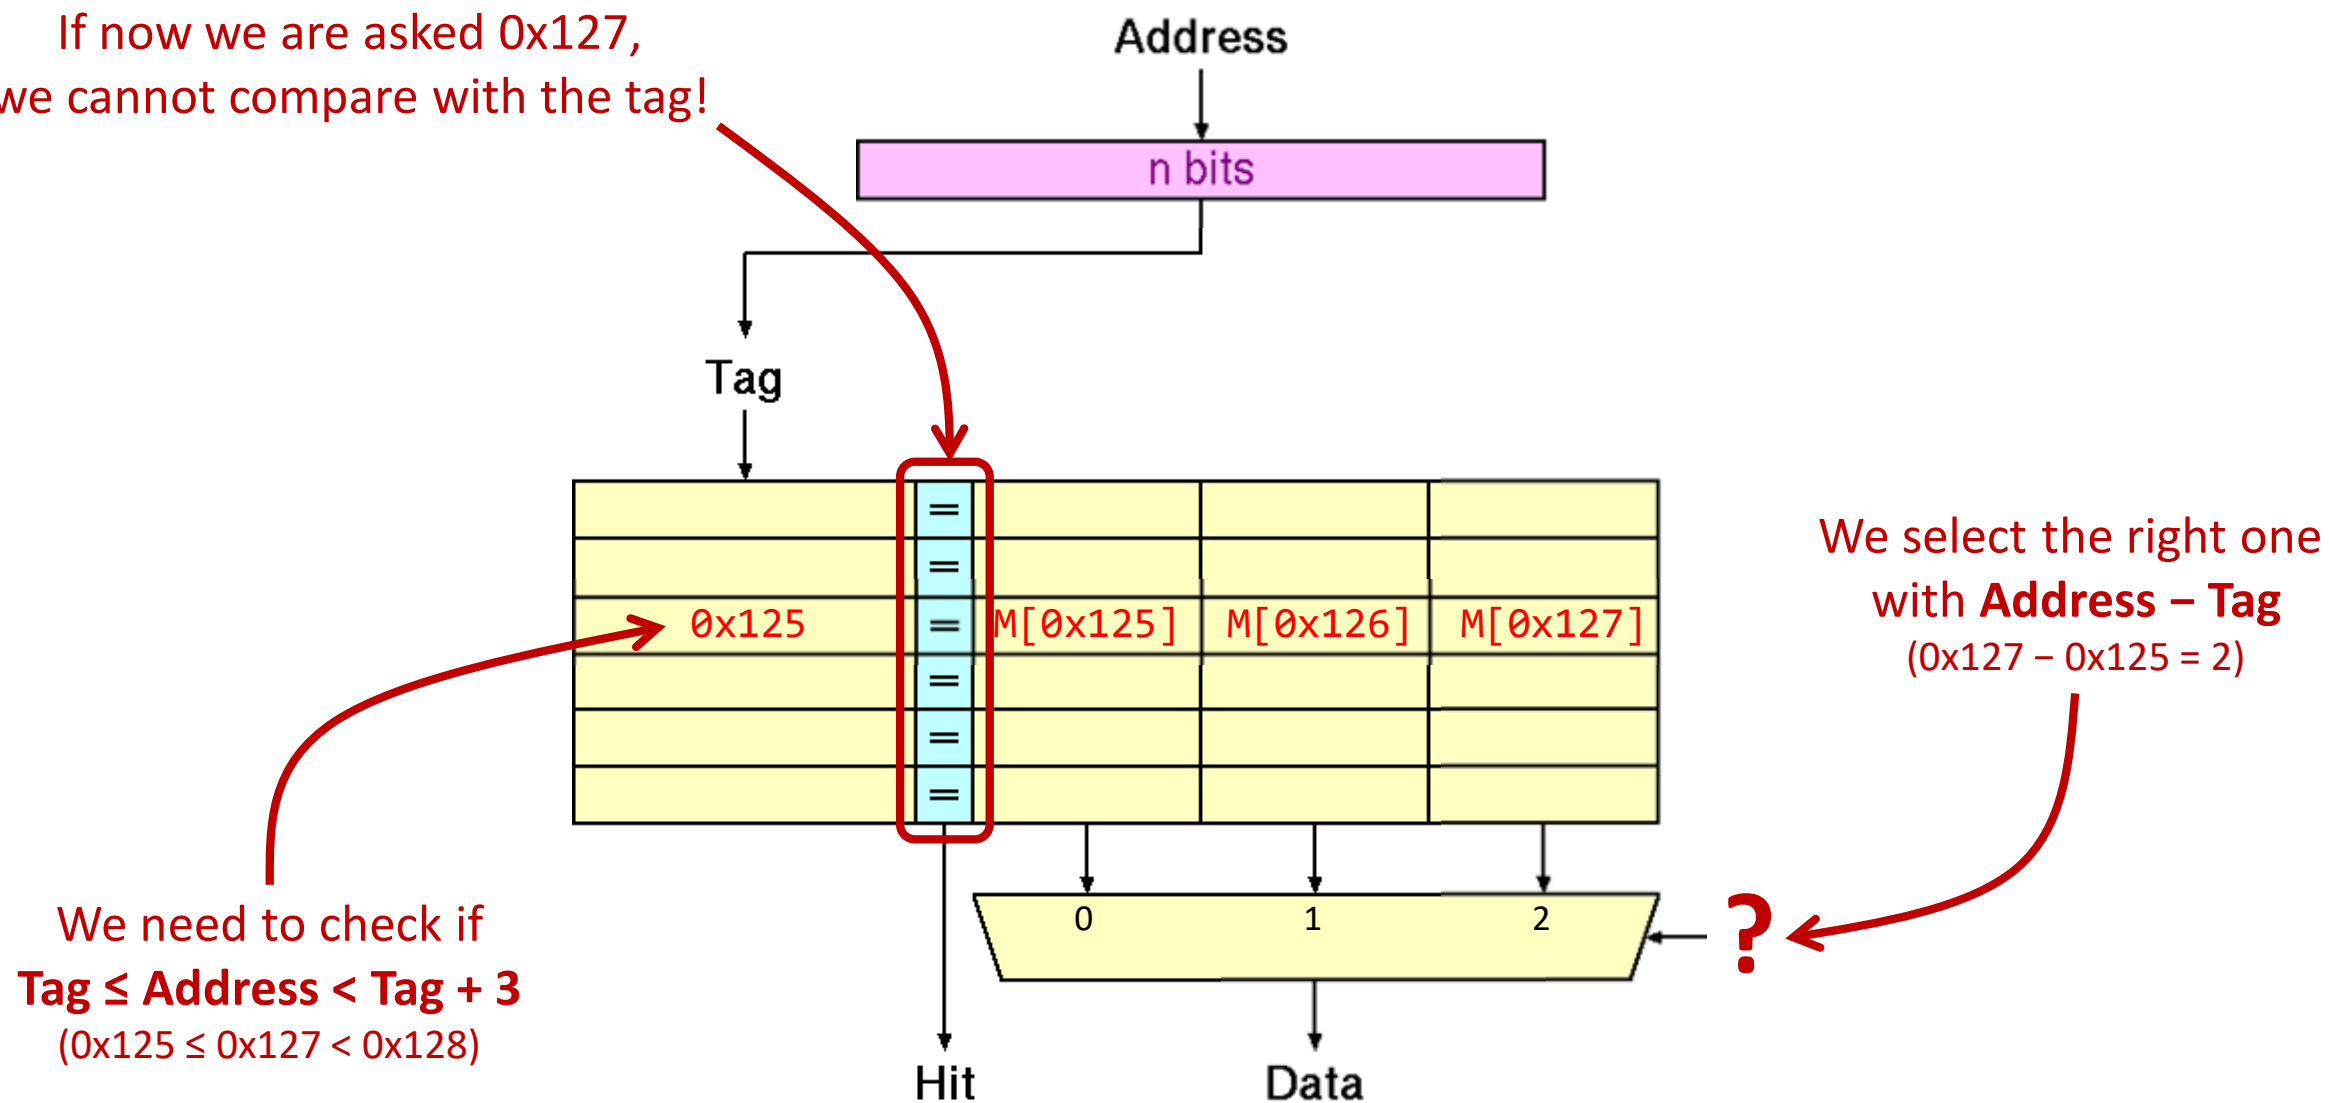


# Why Not This?!

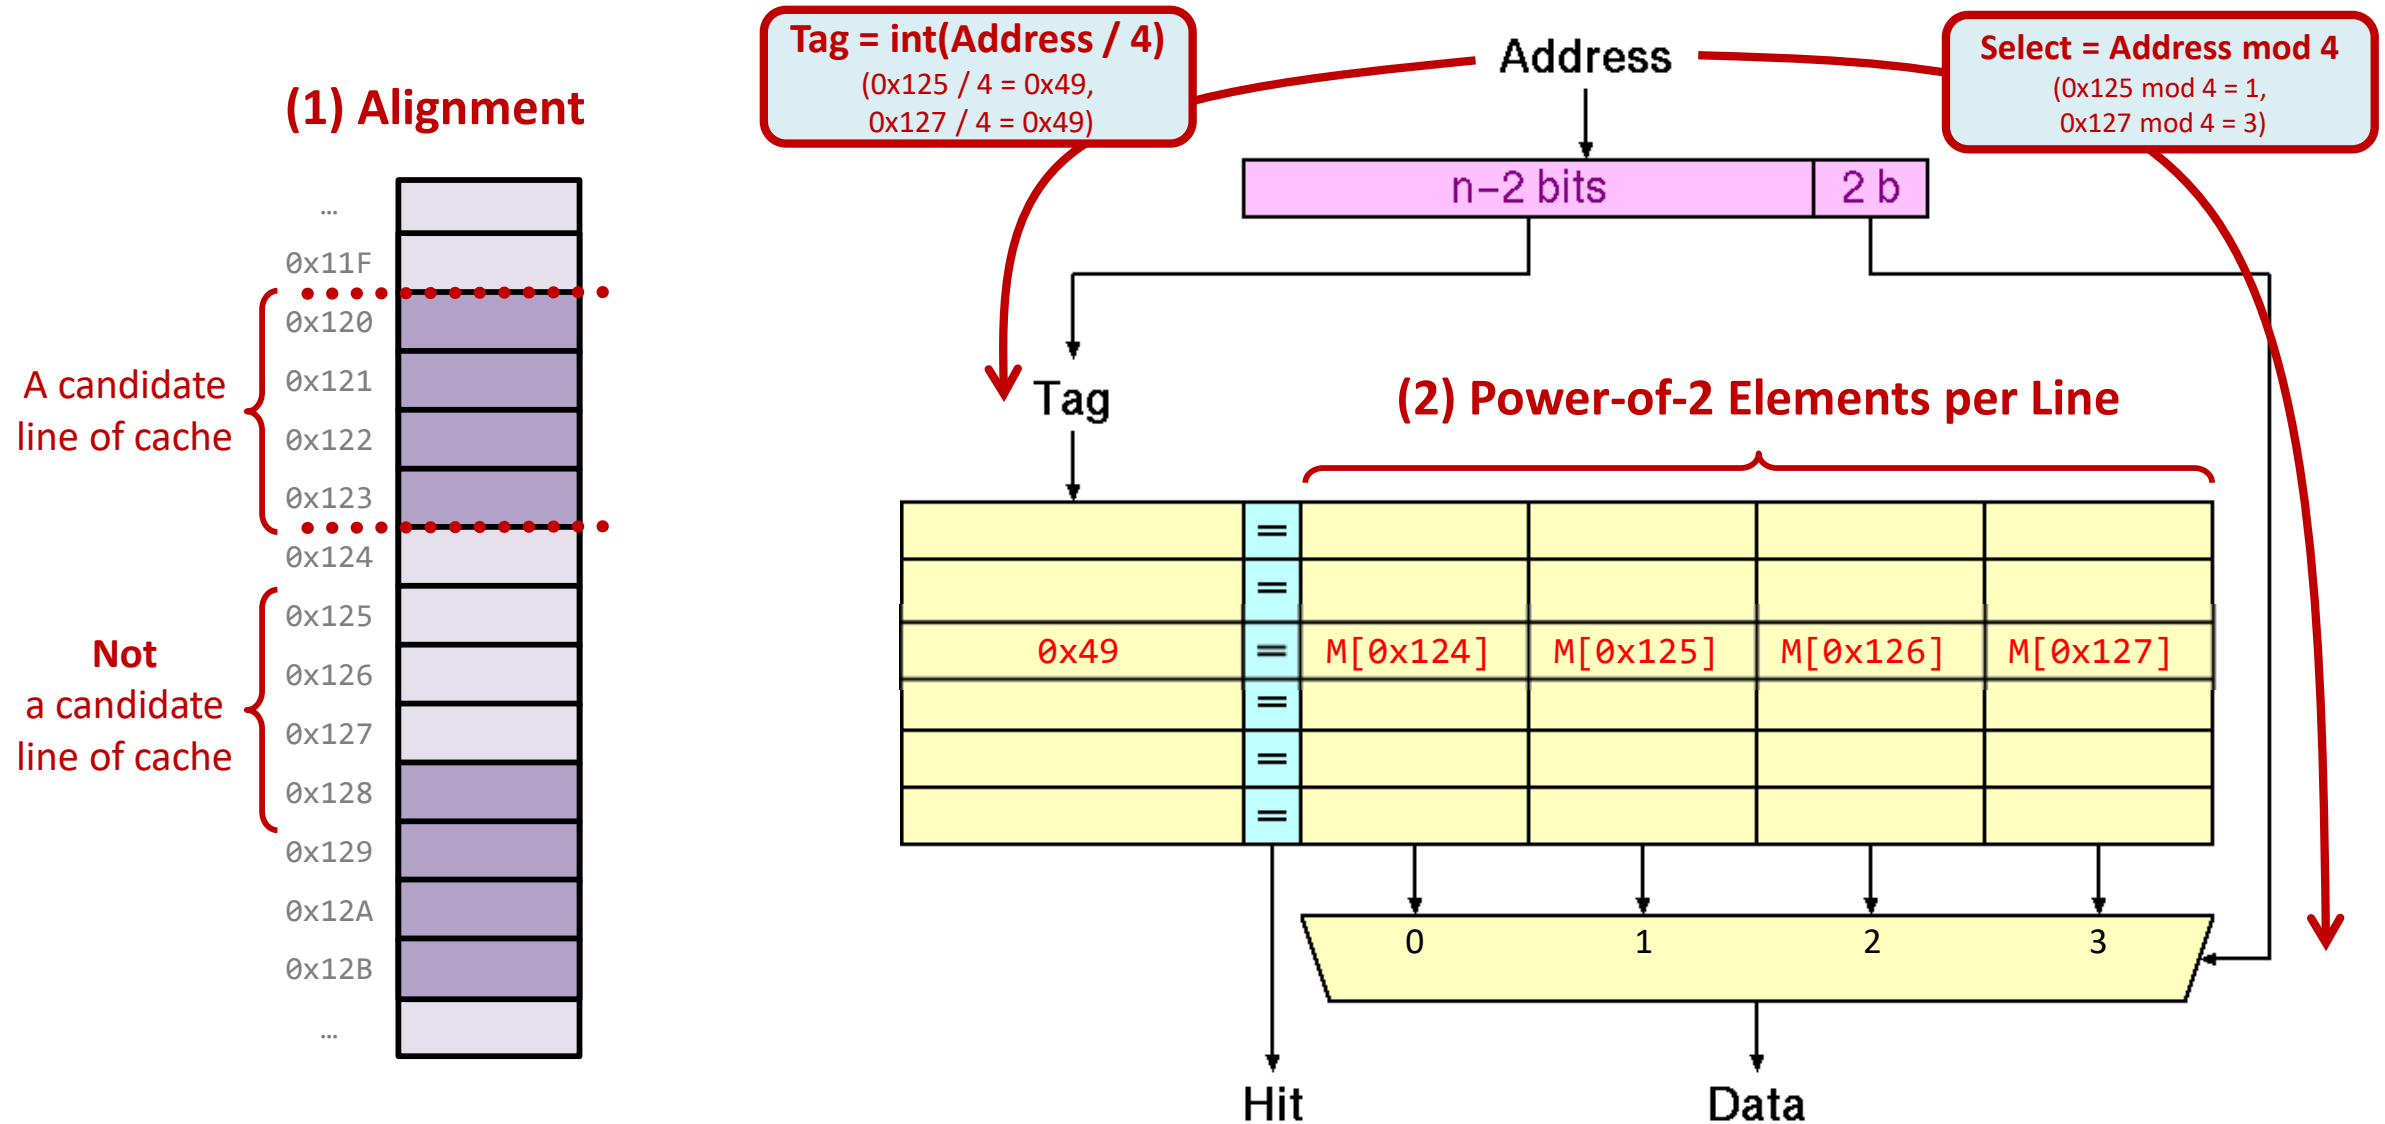


# Why Not This?!


If now we are asked 0x127,  
we cannot compare with the tag!



# This Is Much More Hardware Friendly!



# Hardware Friendliness Is Essential



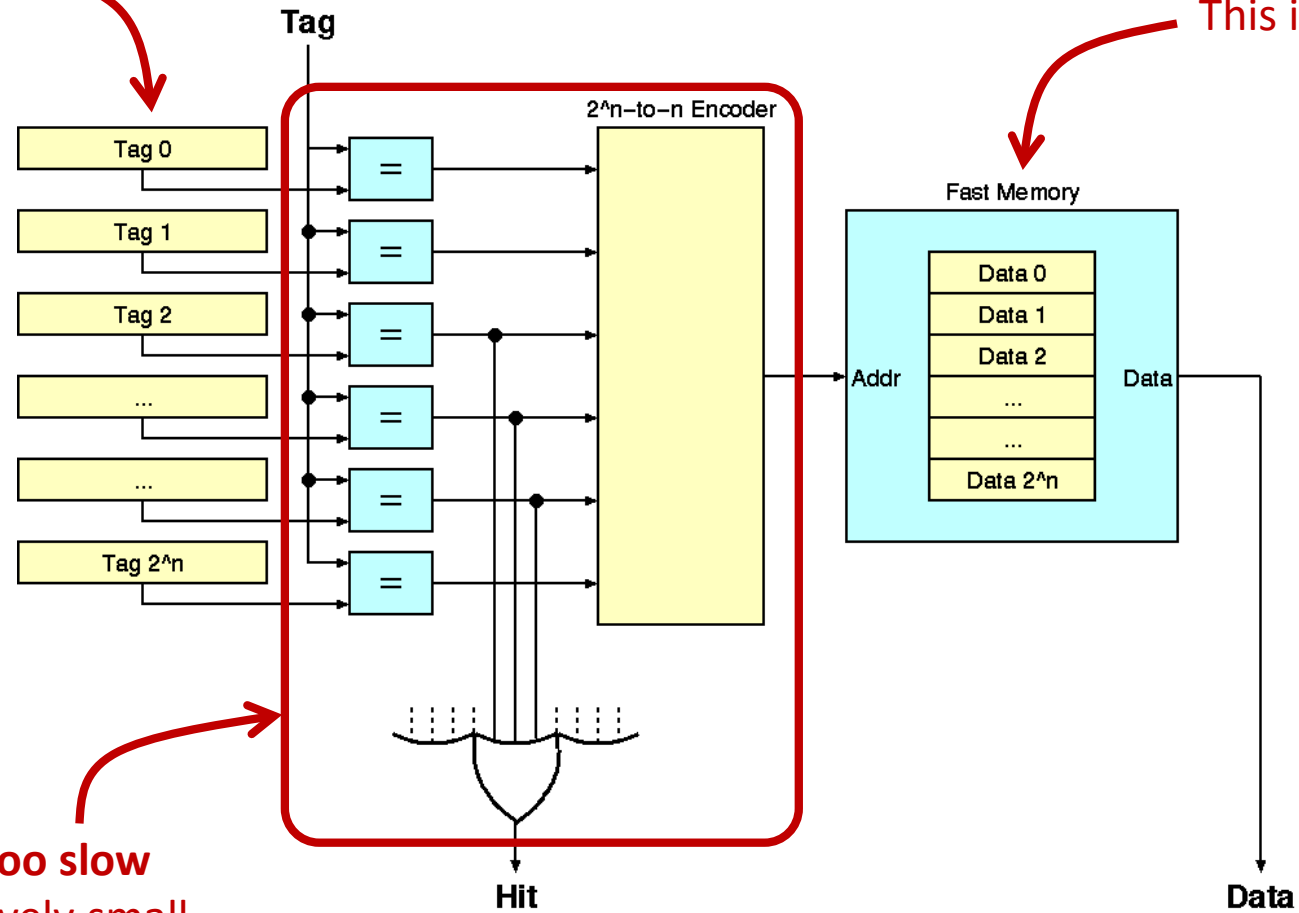
This cache may be better, because if the processor asks 0x127 we bring in potentially **useful** stuff instead of **stale** stuff...

...yet, this cache is **FEASIBLE!**

# Fully-Associative Cache


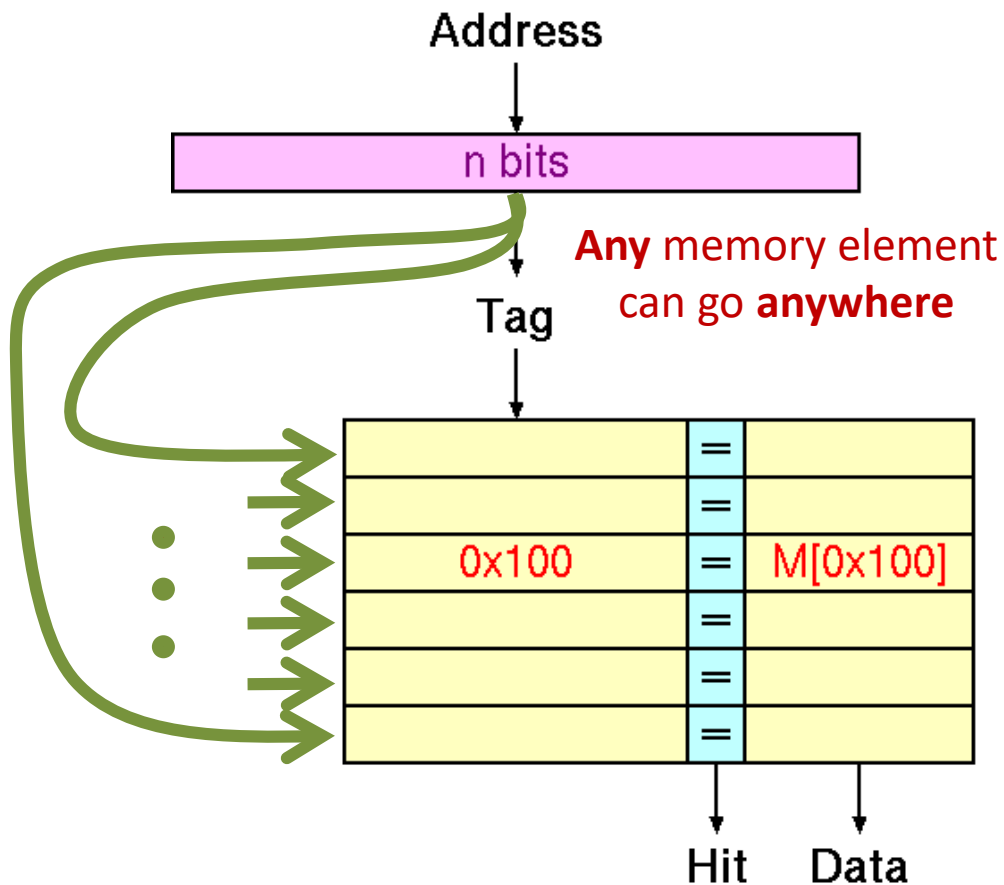
These are registers and relatively cheap

This is a RAM and is cheap



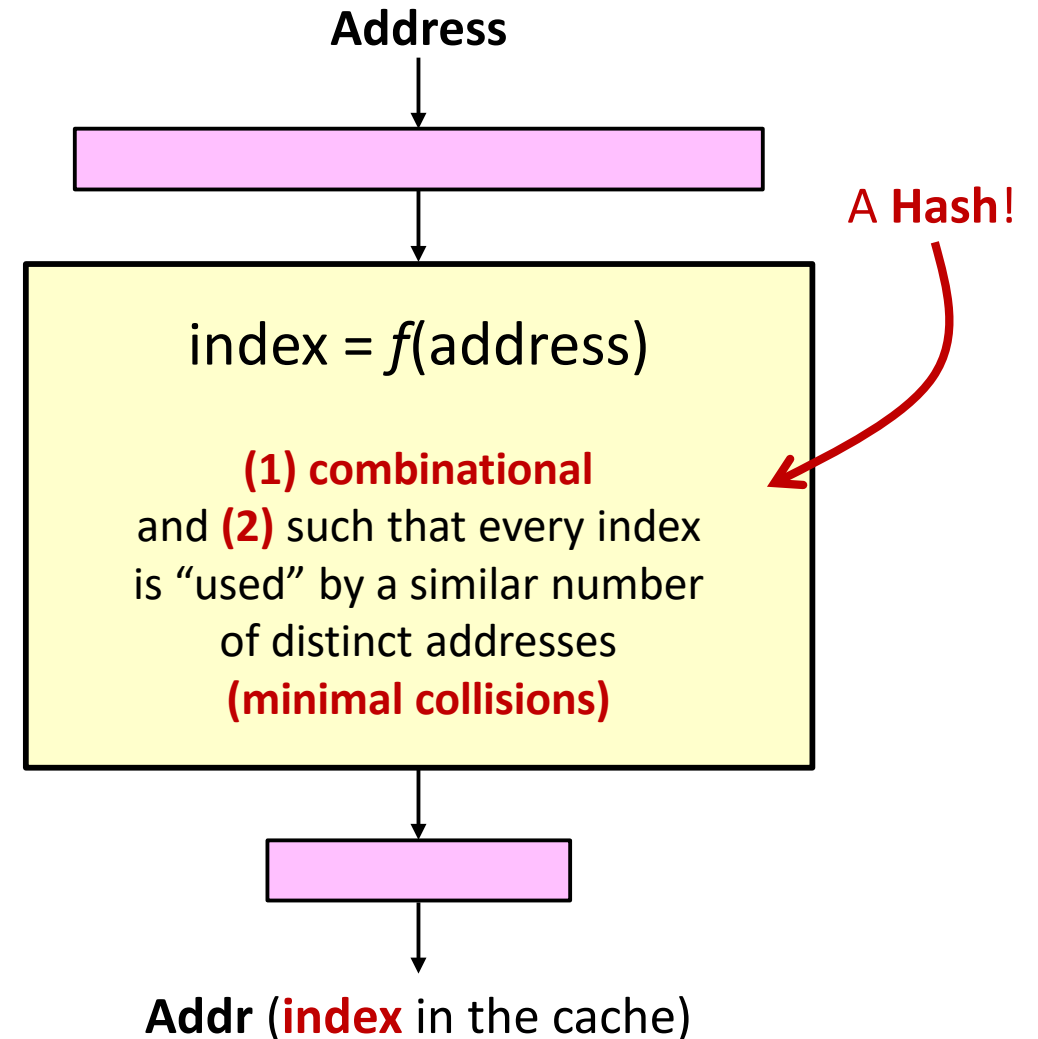

This is **too costly** and **too slow** unless the cache is relatively small

# How Can We Make It Simpler?







# How to Generate Addr and Tag?



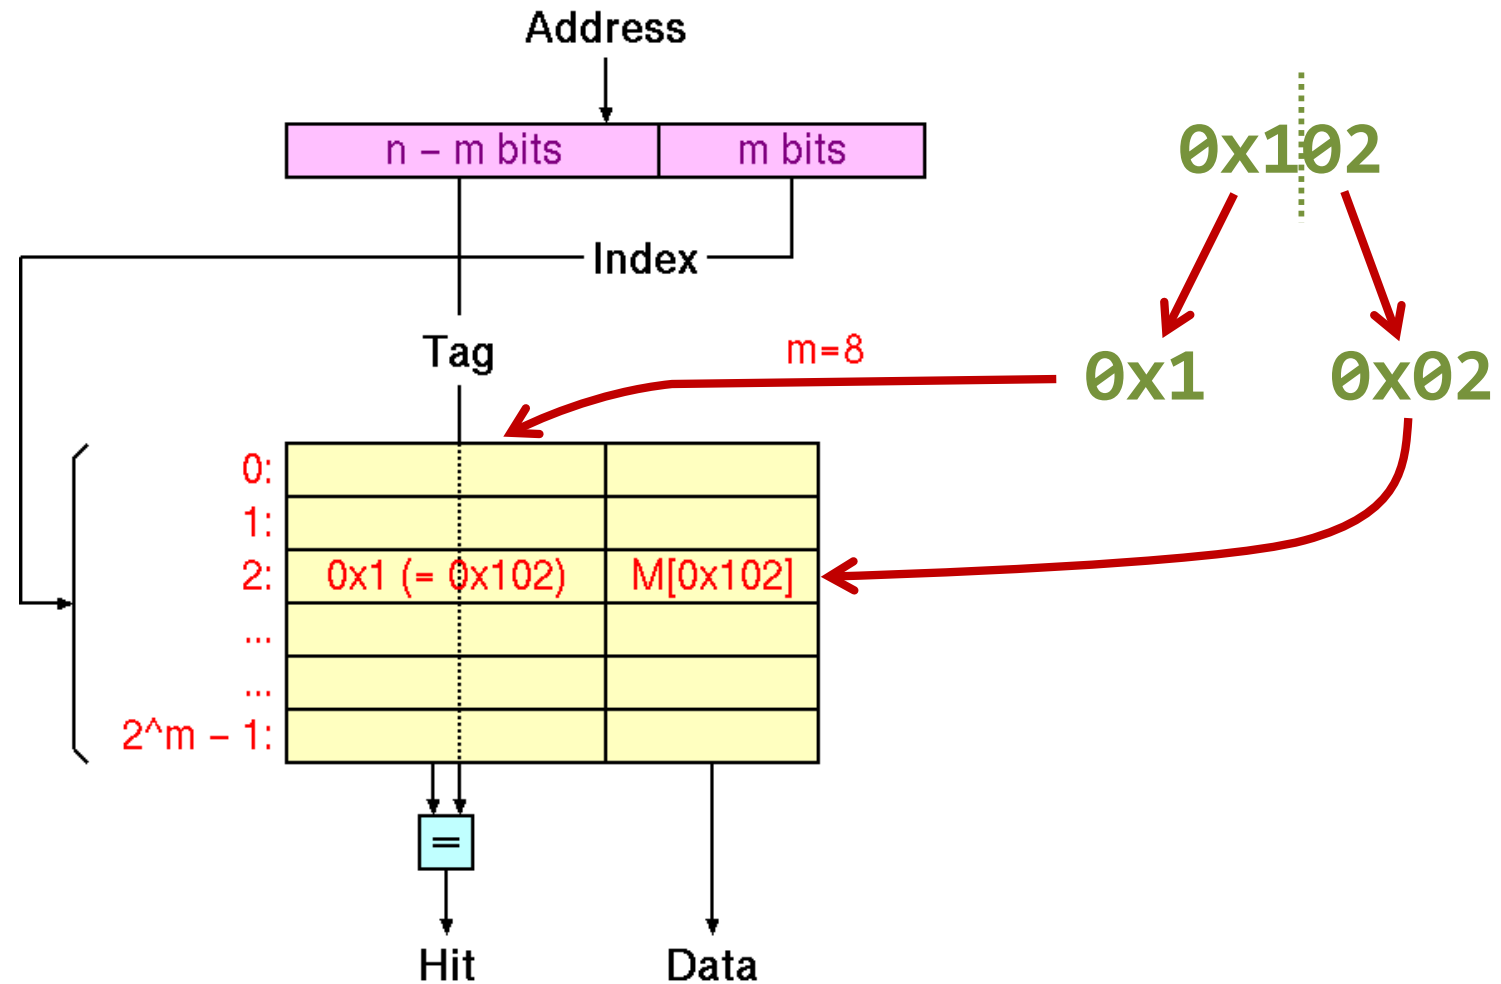
# The Simplest Hashes



# Direct-Mapped Cache


**Much cheaper!**

...but does it have any drawbacks?!

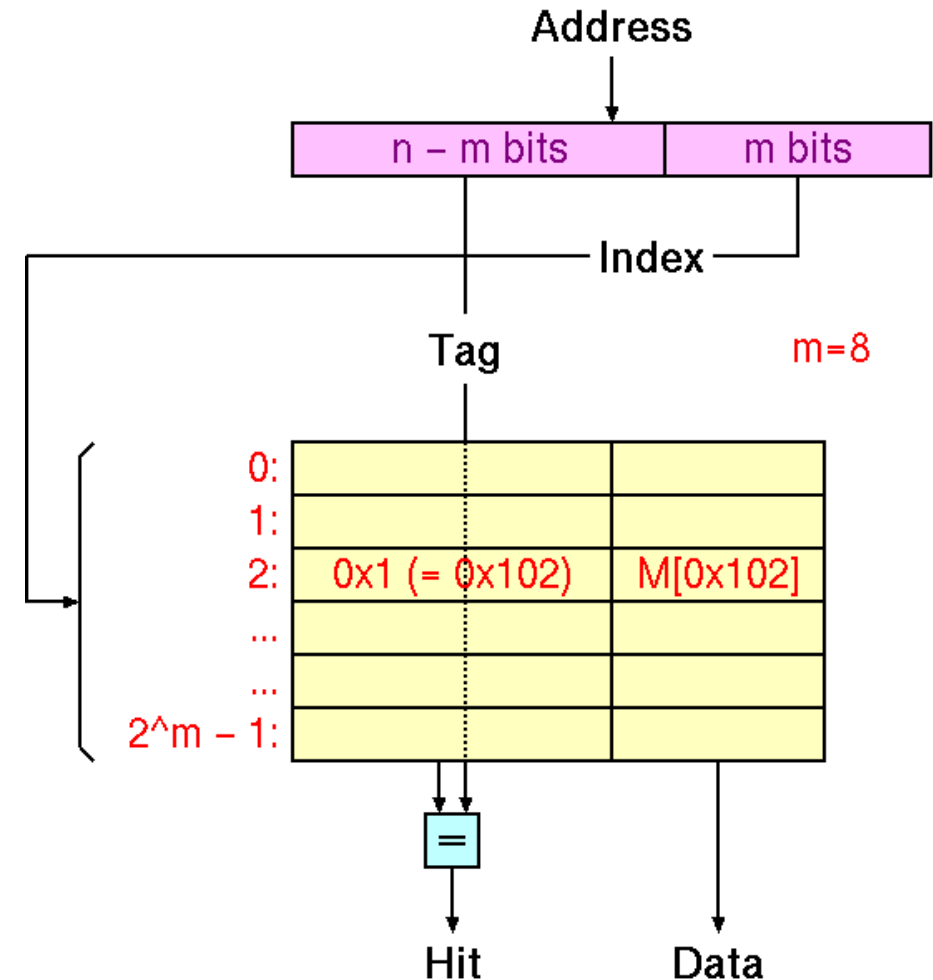


# Which One Is the Best Cache?

## Fully Associative Cache



## Direct-Mapped Cache



# Which One Is the Best Cache?

- Consider a **fully associative** and **direct-mapped** cache, both with **64 lines** with **four words per line** ( $\rightarrow$  256 words per cache)
- Suppose accesses at 0x100, 0x101, 0x200, 0x102, 0x300, 0x103, 0x201, 0x102, 0x301, 0x103,...
- What is the **Hit Rate** of each of these two caches?

Address

	=				
	=				
	=				
	=				
	=				
	=				
	=				
	=				
	=				
⋮					⋮
	=				
	=				

Address

0:				
1:				
2:				
3:				
4:				
5:				
6:				
7:				
	⋮			⋮
62:				
63:				

[illegible]

# Cache Pollution, or Conflict Misses

	0x100	0x101	0x200	0x102	0x300	0x103	0x201	0x102	0x301	0x103
Fully Ass.	M	H	M	H	M	H	H	H	H	H
Direct Mapp.	M	H	M	M	M	M	M	M	M	M

- Addresses 0x1..., 0x2..., and 0x3... use the same line of the direct-cache (they are said to **alias**) → **Cache pollution** or **conflict misses**
- Fully associative cache is **much more robust** than direct mapped

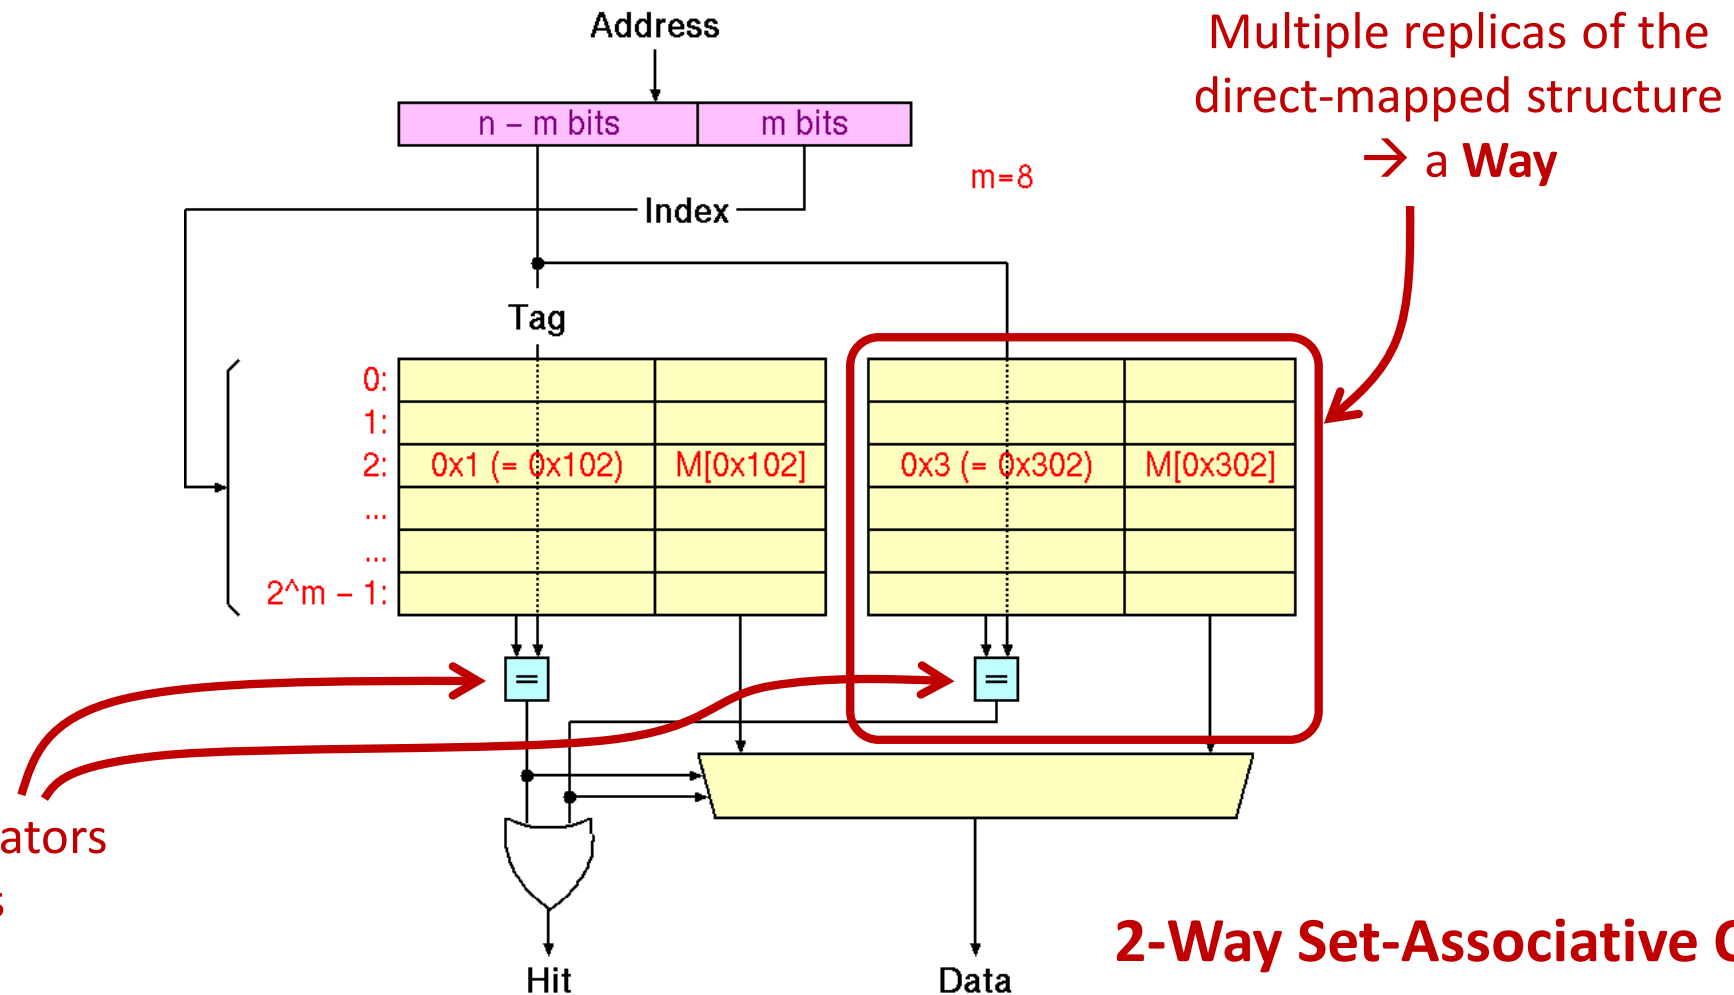
# Associativity

- The probability of aliasing is related to the associativity of caches
- **Associativity** indicates the number of different positions in a cache where one element of data can be placed:
  - **Fully-associative**: every word can go in every line of the cache (hence the “full”) → associativity is the number of lines in the cache
  - **Direct-mapped**: every word is “mapped” to a single line of the cache → associativity is 1

Can we think  
of any **intermediate** possibility?



# Set-Associative Cache

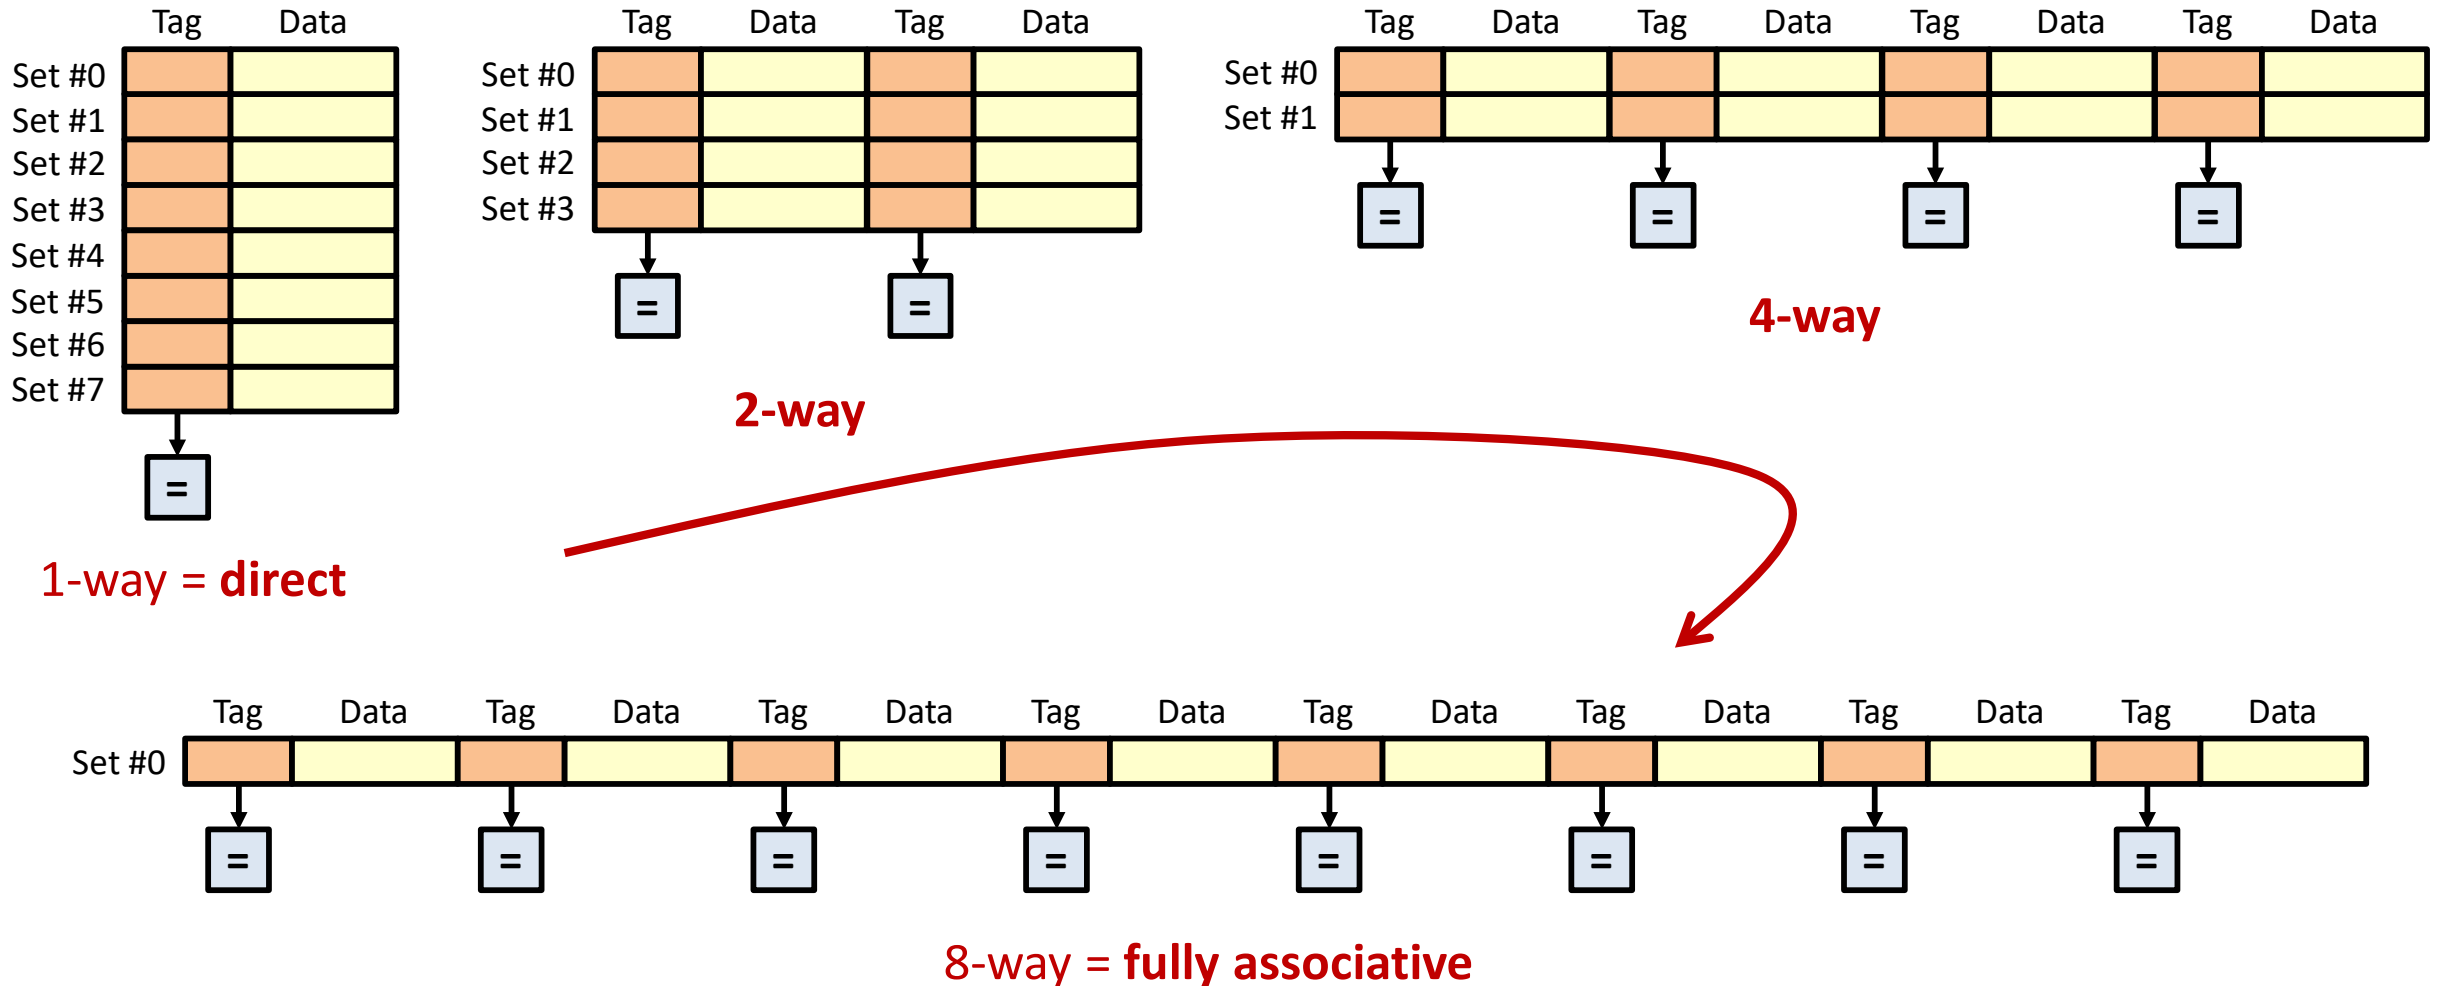


**2-Way Set-Associative Cache**

# comparators = associativity

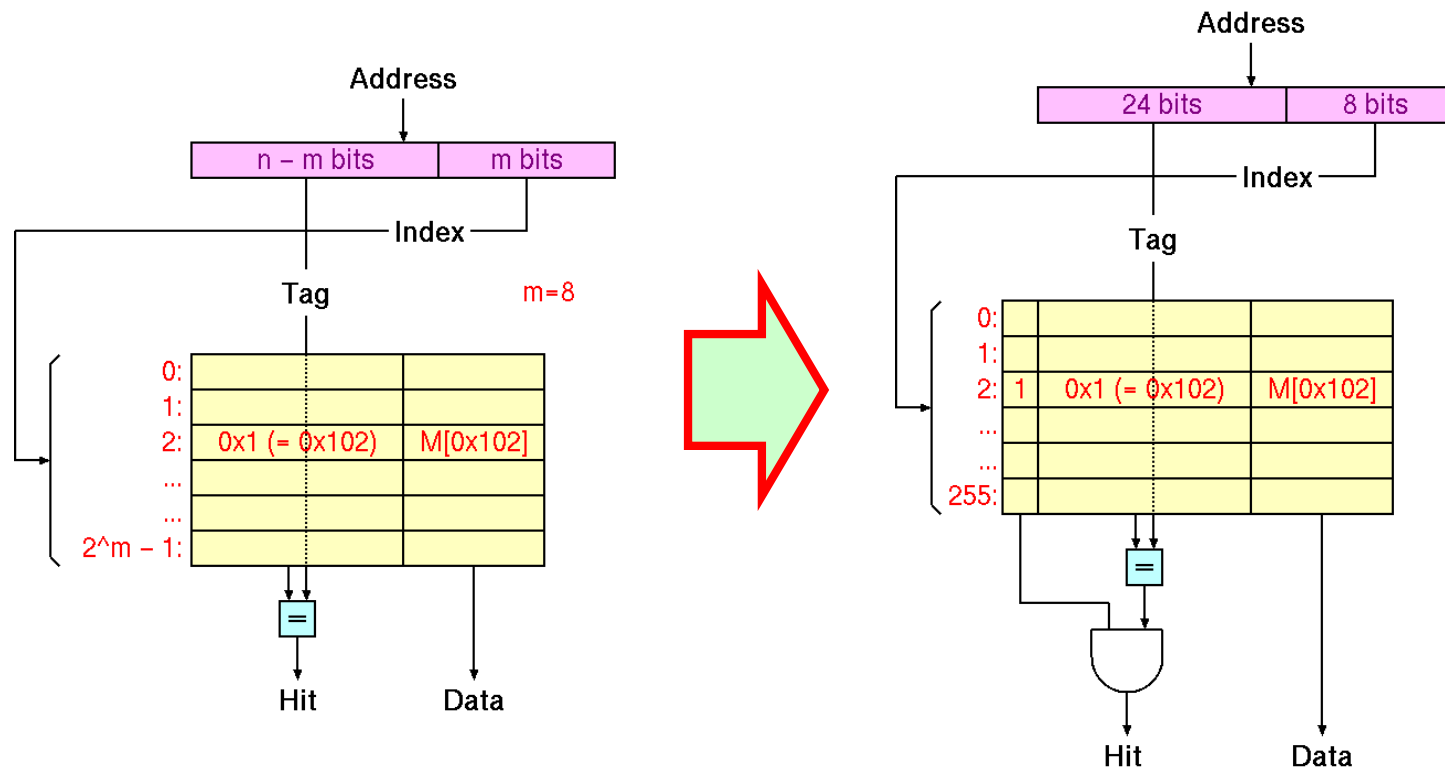


# A Continuum of Possibilities



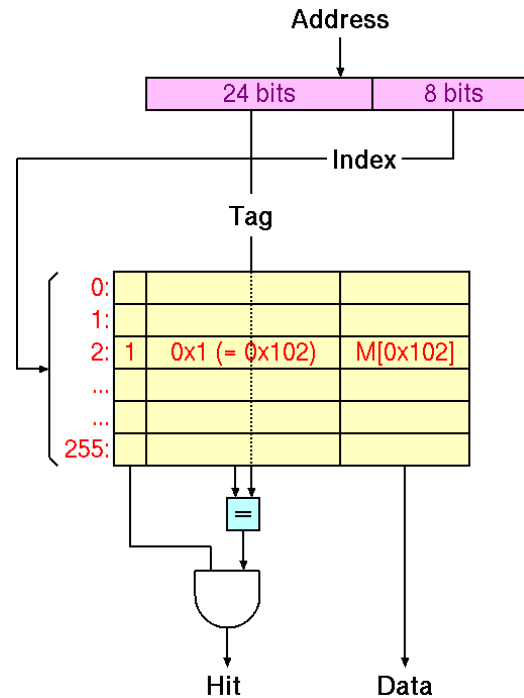
# Validity

- Initial cache content is garbage
- All caches need a special bit (**Valid Bit**) in each cache line to indicate whether something meaningful is in the specific cache line ('0' at reset)

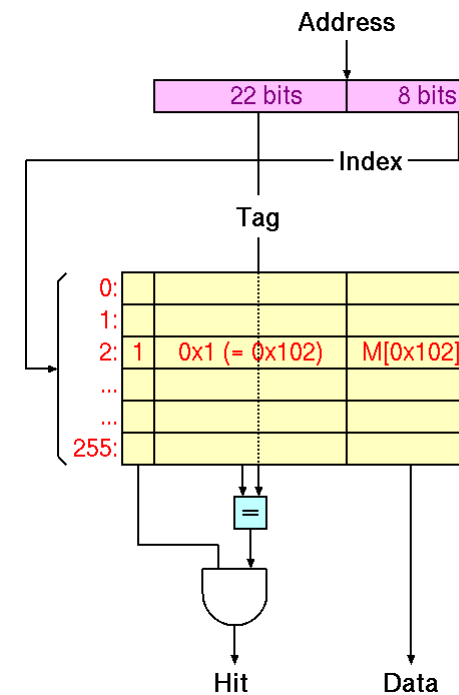


# Addressing by Byte

- If addressing is **by byte** and **word size is  $2^n$  bytes**, the  **$n$  least-significant bits** of the address represent the byte offset and are thus **irrelevant**



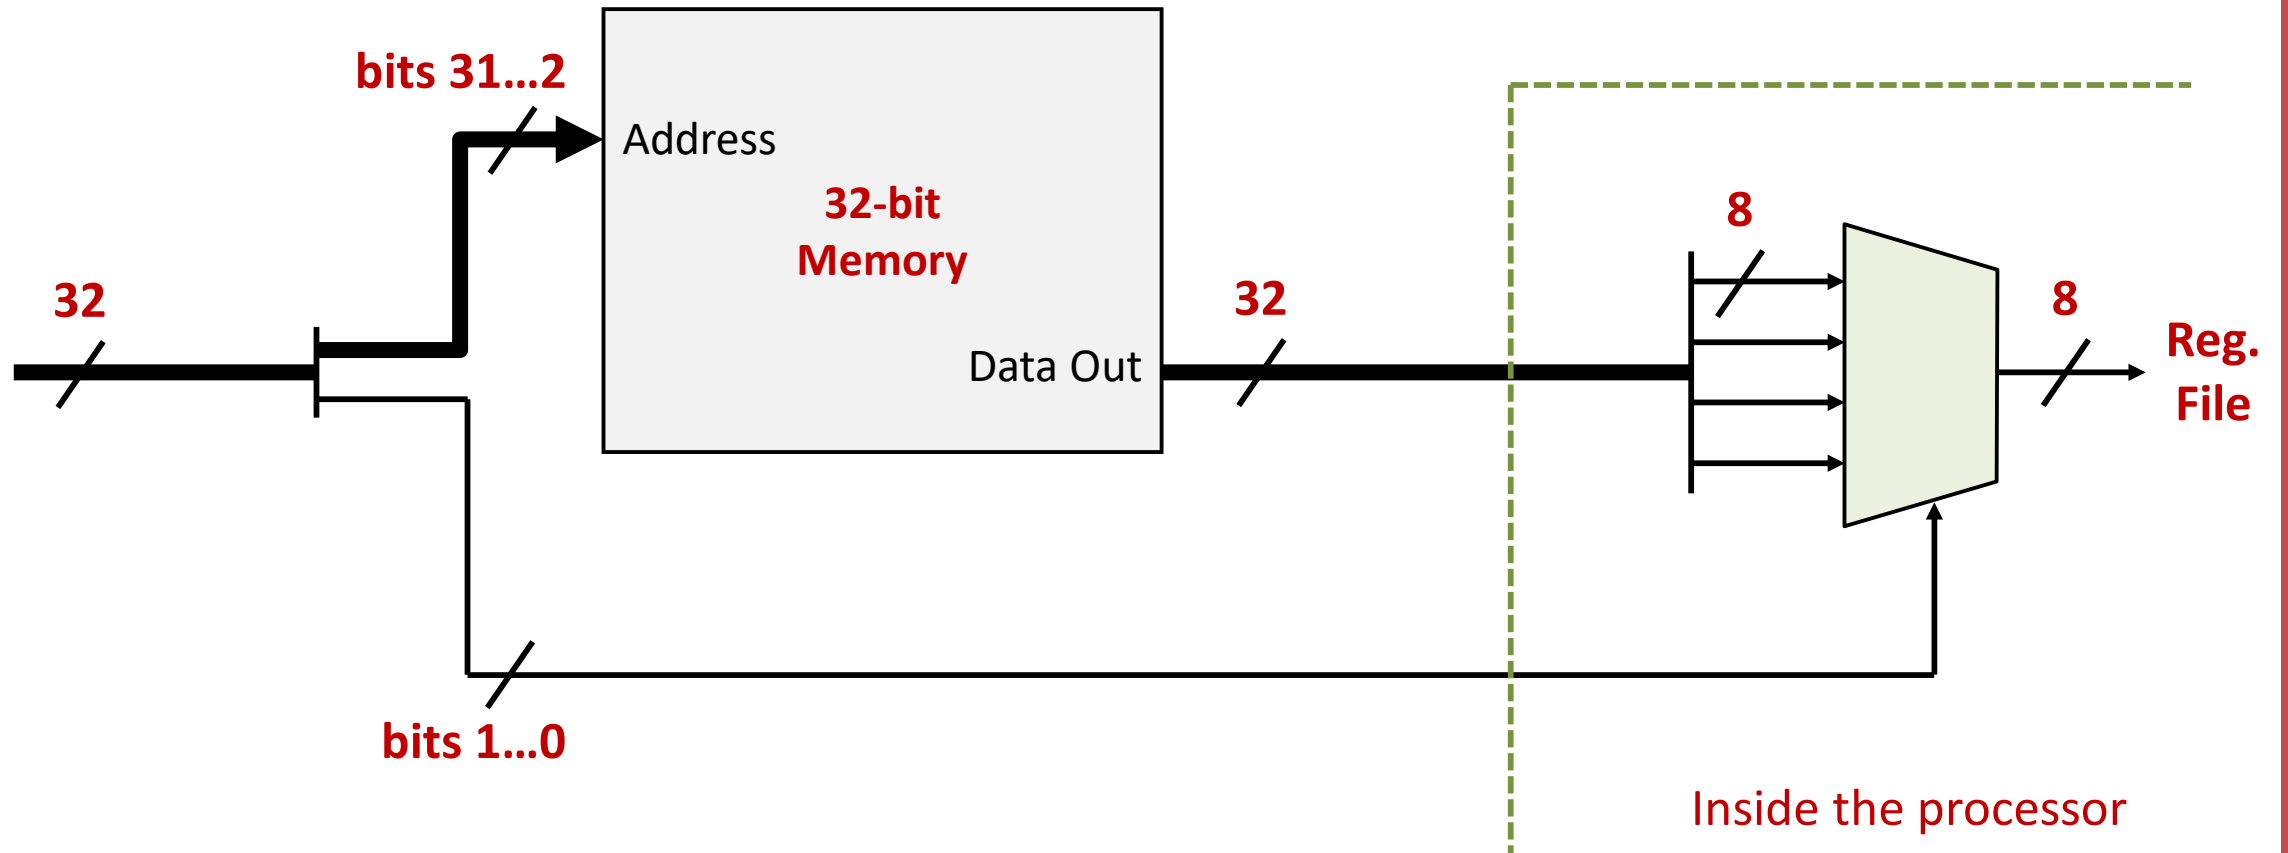
Addressing **by word**



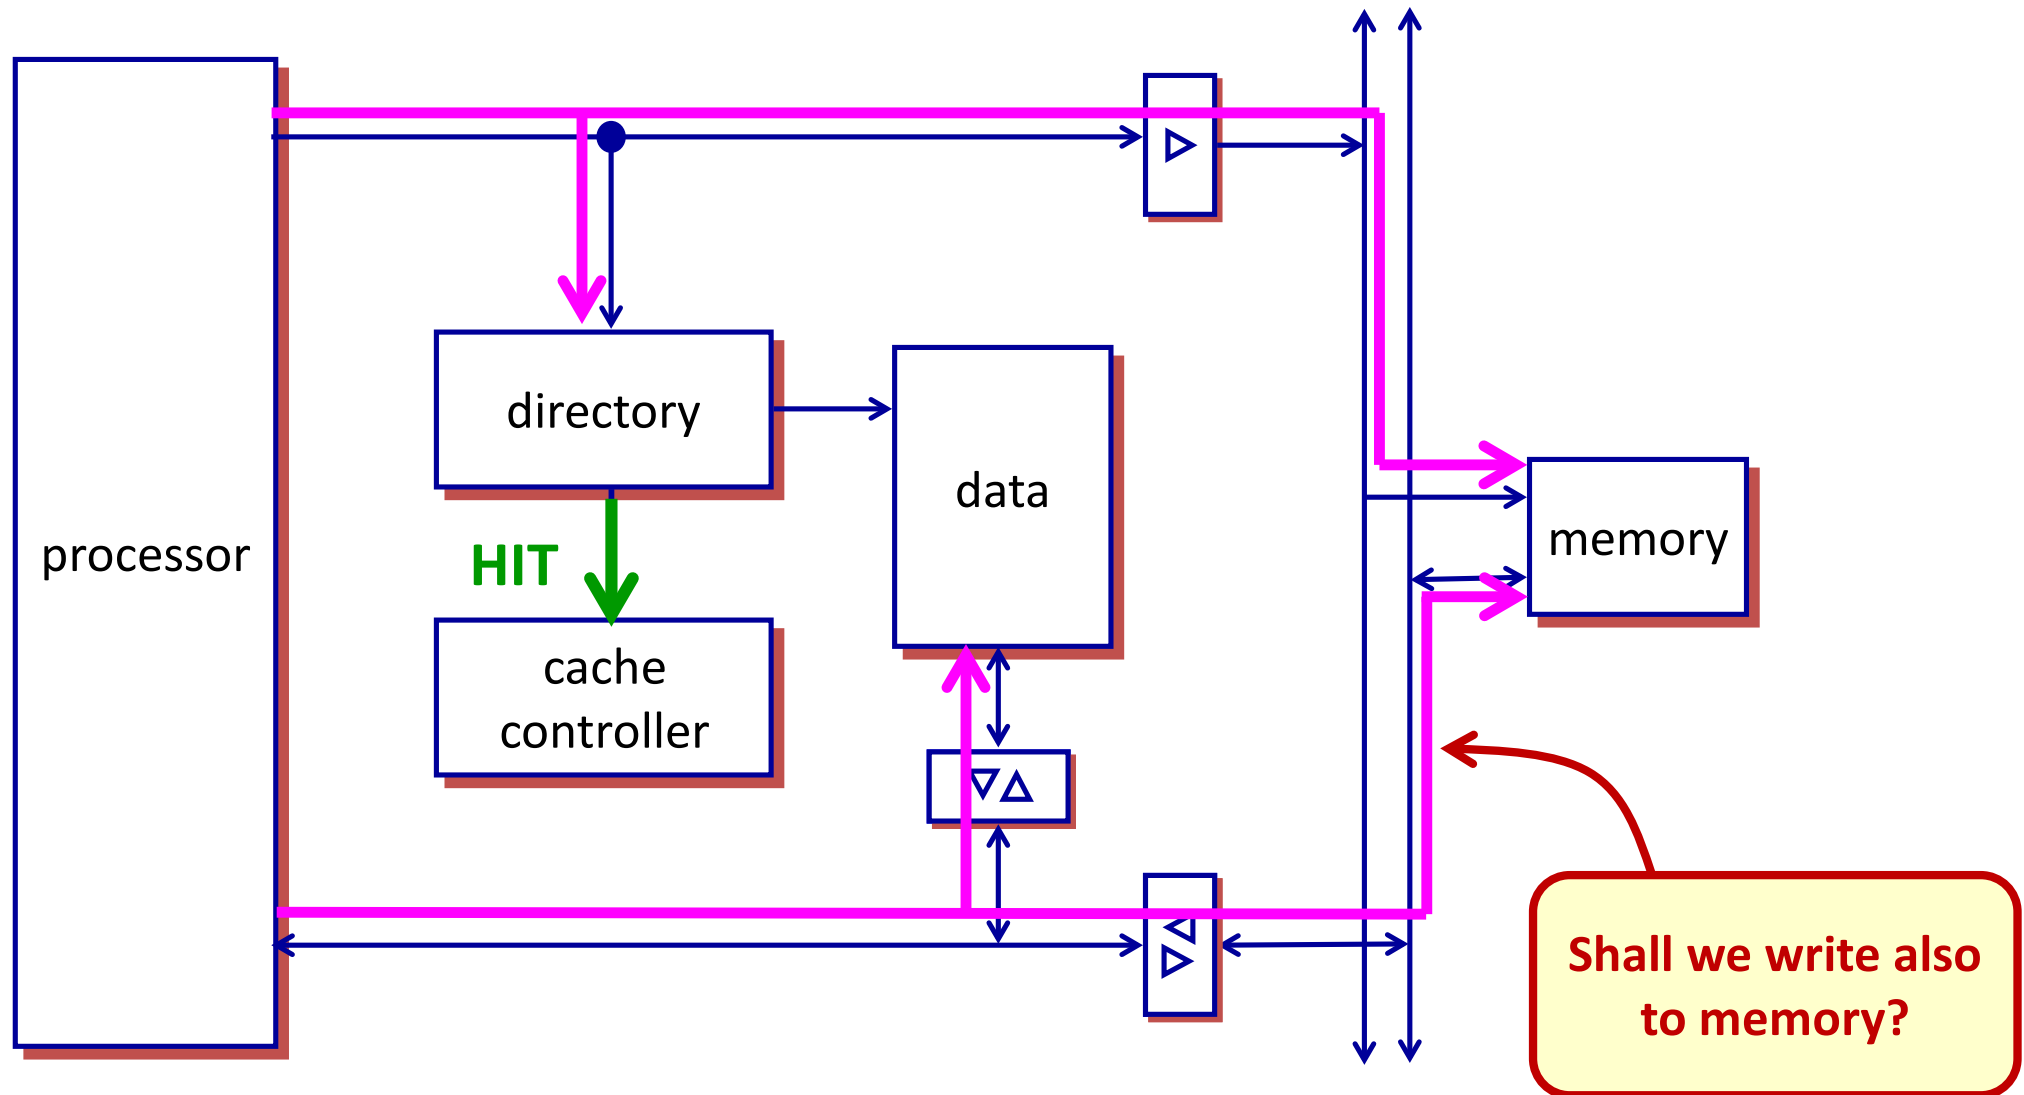
Addressing **by byte**

These bits are **internal to the processor** and do not even get to the memory system!

# Loading Bytes (1b)



# Write Hit

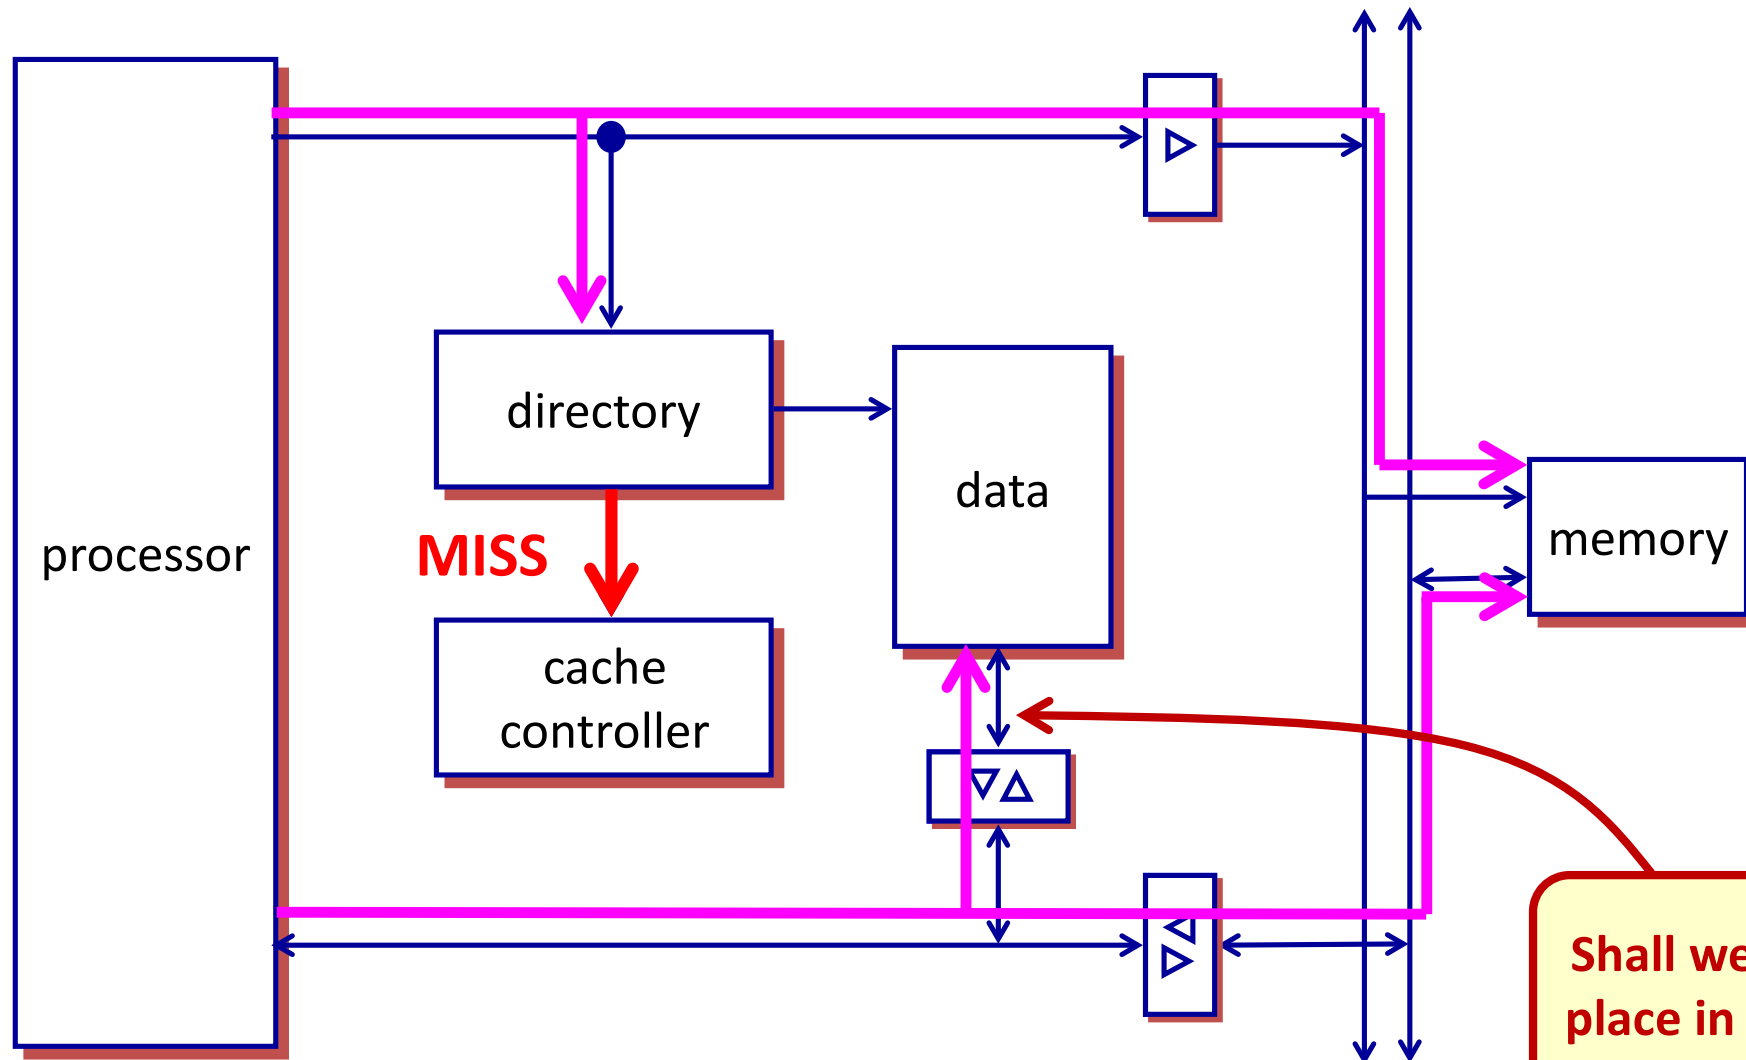


# Write Policies

- **Write-through**: on a write, data are always immediately written into main memory
  - Simpler policy
  - May keep the memory/buses busy for nothing
- **Write-back** or **Copy-back**: on a write, data are only updated in the cache (hence, main memory data will become wrong/obsolete)
  - Needs a **Dirty Bit** to remember that cache **data are incoherent with memory**
  - When a dirty line is **evicted**, first it must be **copied back** to main memory



# Write Miss



Shall we allocate a place in the cache?

# Allocation Policies

- **Write-allocate**: on a write miss, data are also placed in the cache
  - Simple and straightforward
  - Need to **fetch the block of data** from memory **first**
  - If the processor writes a lot of data that it will never read back, it may **unnecessarily pollute the cache**
- **Write-around** or **Write-no-allocate**: on a write miss, data are only written to memory
  - If the processor will load from the same address, it will be a **Read Miss**

# The “3 Cs” of Caches

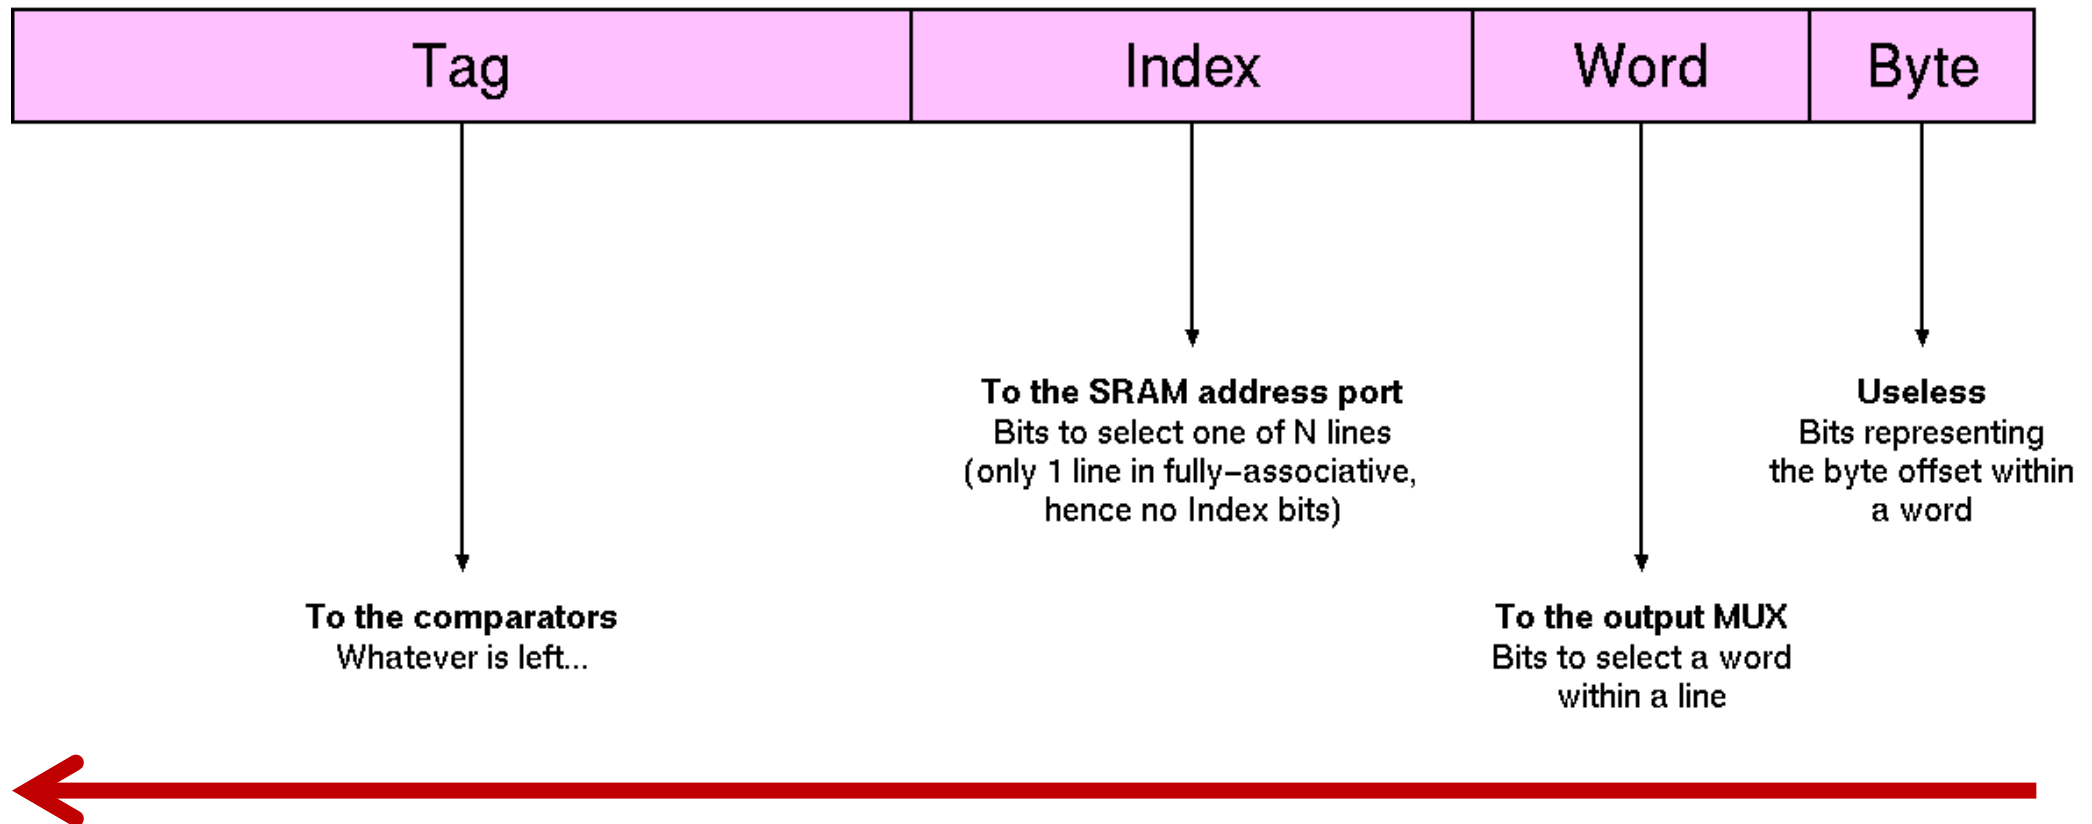
- Three types of **cache miss**
  - **Compulsory** → Misses that would happen in an infinitely large fully-associative cache with the same blocks (also called **cold-start** misses or **first-reference** misses)
  - **Capacity** → Additional misses that occur because the corresponding block has been evicted due to the **limited capacity of the real cache**
  - **Conflict** → Further misses that occur because the corresponding set is full and the corresponding block has been evicted due to the **limited associativity of the cache**
- Useful to understand the **source of the limited performance**

# Summary of Cache Features

- **Cache size**: total data storage (usually excluding tags, valid bits, dirty bits, etc.)
- **Addressing**: by byte or word
- **Line** or **block size**: bytes or words per line
- **Associativity**: fully-associative,  $k$ -way set-associative, direct-mapped
- **Replacement policy** (except for direct mapped): LRU, FIFO, random, etc.
- **Write policy**: write-through or write-back
- **Allocation policy**: write-allocate or write-around

# Summary of Cache Addressing

## Address



Allocate bits as required, from LSB to MSB

# References

- Patterson & Hennessy, COD – RISC-V Edition
  - **Sections 5.3** and **5.4**
  - **Sections 5.9** and **5.15** for more info