

CS-200

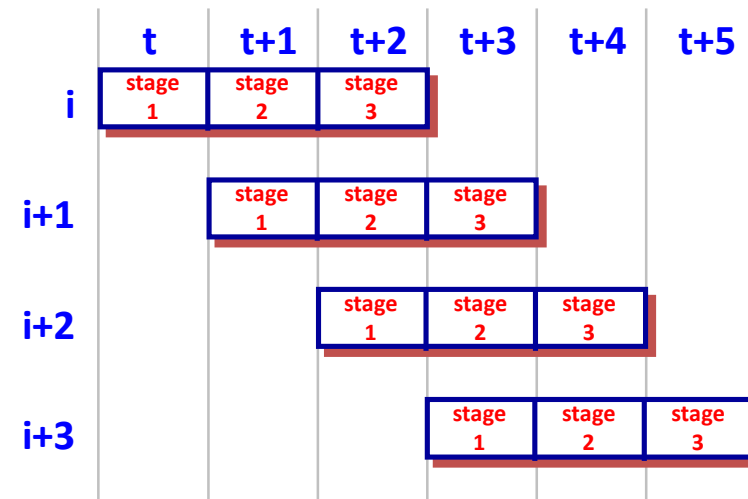


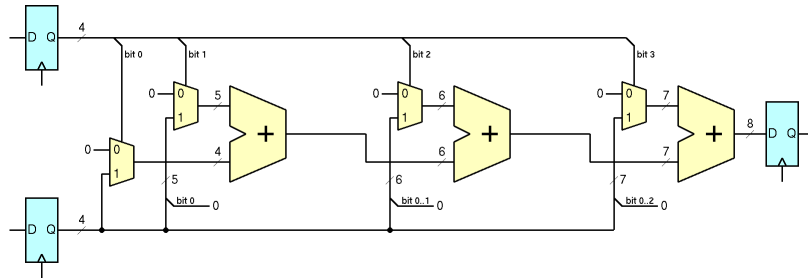
Computer Architecture

Part 4c. Instruction Level Parallelism


Pipelining

Paolo Ienne
<paolo.ienne@epfl.ch>

Pipelining



Practical Pipelining

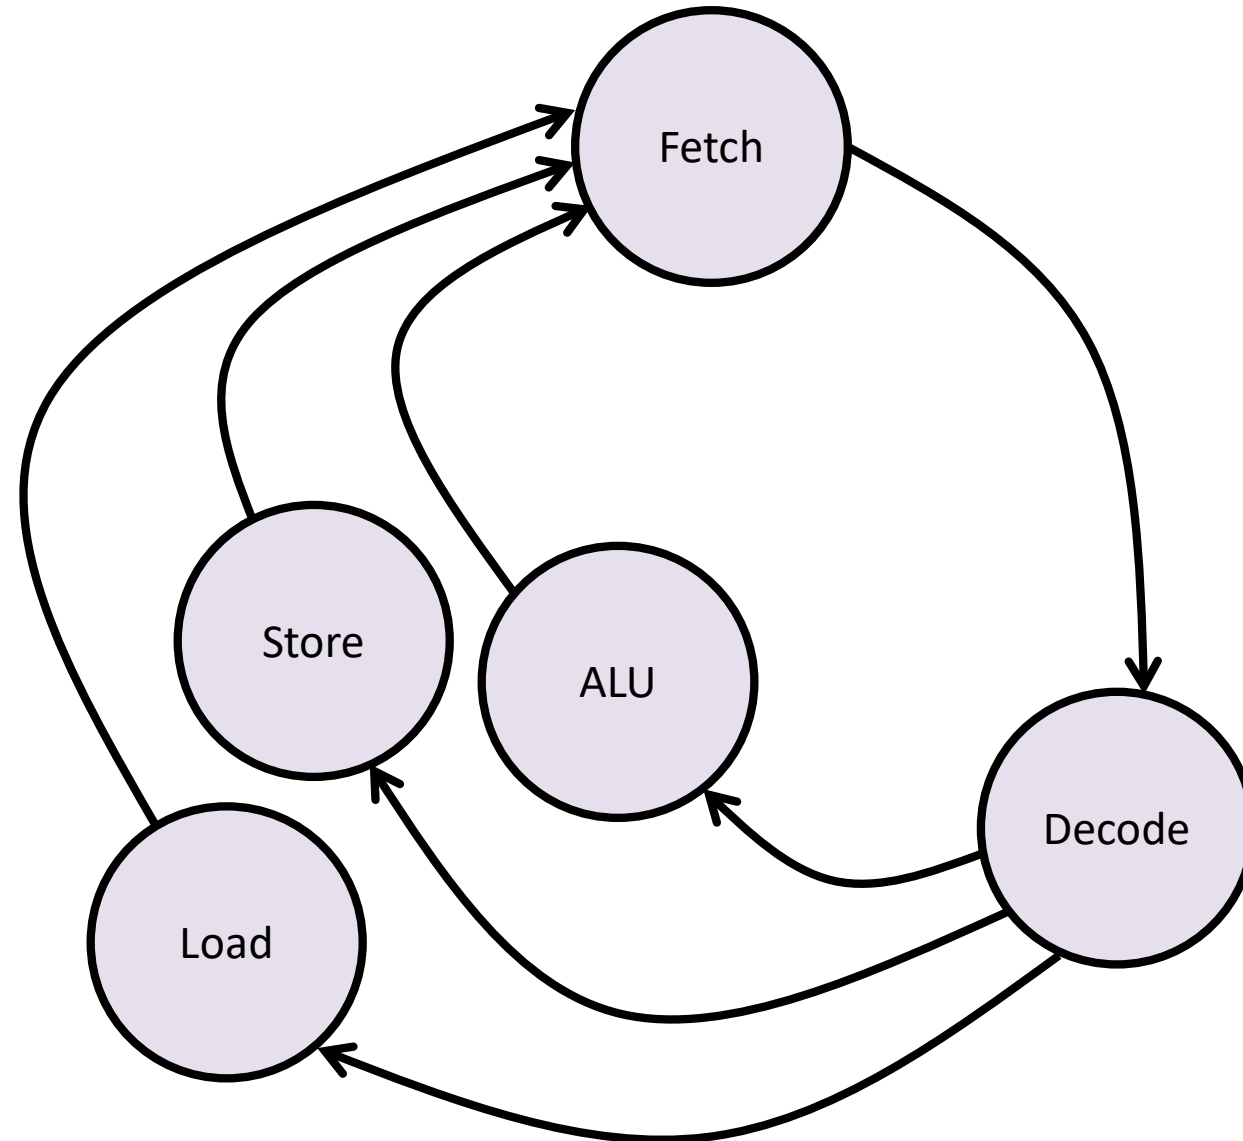


- **Latency** $\lambda_{pipe} = N \cdot \max_{i=0..N-1} (T_i + T_{FF}) = N \cdot T_{CLK,pipe} = N / f_{pipe} > \lambda_{orig}$
- **Throughput** $\phi_{pipe} = 1 / \max_{i=0..N-1} (T_i + T_{FF}) = f_{pipe}$

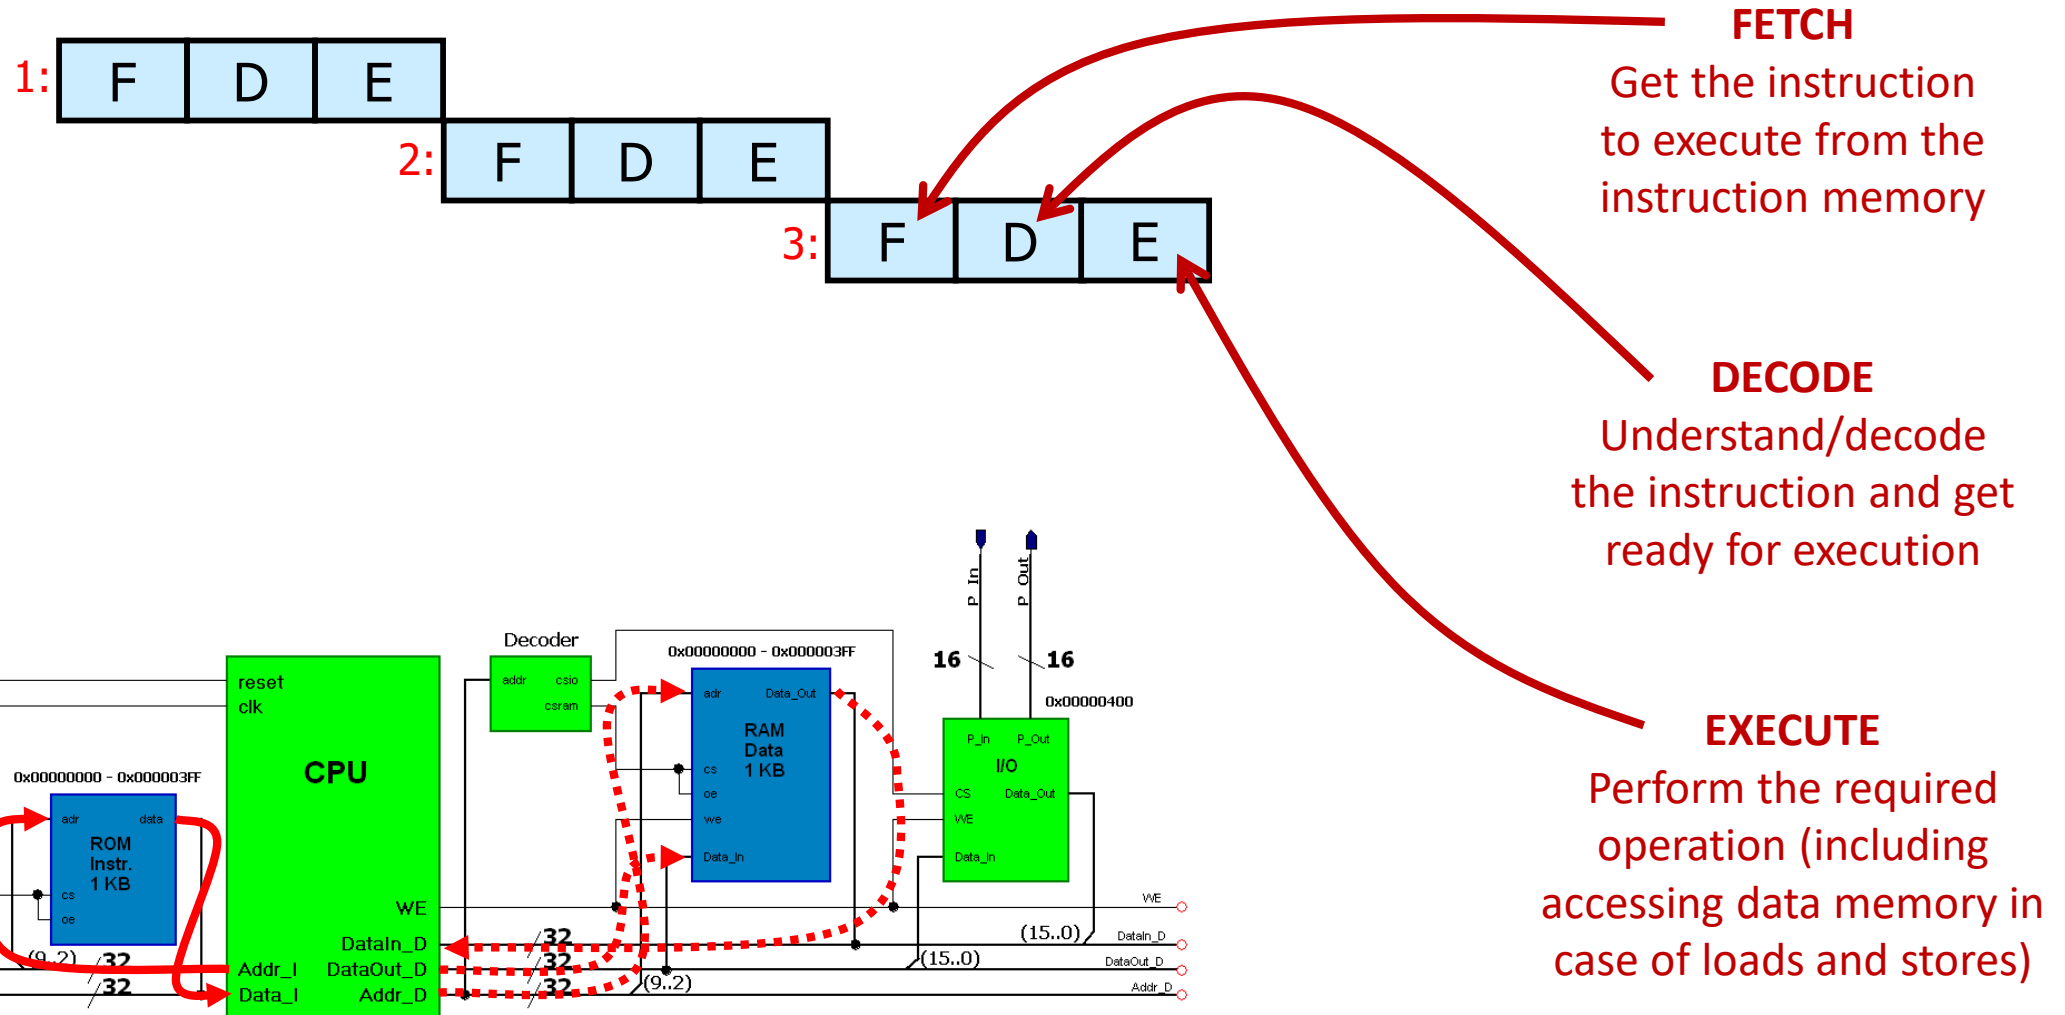
Pipelining for Processors?

- Pipelining useful only if the activity in object needs to be repeated many time
 - We have plenty of **instructions** to execute!
- Pipelining needs to split the single activity in object into many subactivities
 - We have a **good logical split** into *fetch, decode, execute*, etc.





Example: A Simple Multi-Cycle Processor



Example: A Simple Schedule



Pipelining the Processor?



No Hardware Reuse across Stages

- In a multicycle processor, some hardware components may be shared across **states**:
 - **FETCH** typically requires an **adder** to increment the program counter
 - **EXECUTE** naturally needs an **ALU**
 - They are **never used at the same time**, so the ALU can be used to increment the program counter
- In a pipelined processor, there cannot be sharing across **stages**, in general:
 - All stages **active all the time**!
 - Hardware needs to be **replicated** where appropriate

Two Main Problems

1. CISC vs. **RISC**



- Can we **build equally well a pipeline** for a Complex Instruction Set Computer as for a Reduced Instruction Set Computer?
- What does that mean?!

2. Instructions are **not** independent

- Can we execute code **correctly**?


FSM vs. Pipeline

- **Any path** through our FSM represents the **sequence of necessary steps** for the execution of **an** instruction
- **The ordered path** through the pipeline is the **sequence of all possible steps** for the execution of **any** instruction



Adding Instructions to a Multi-Cycle Processor

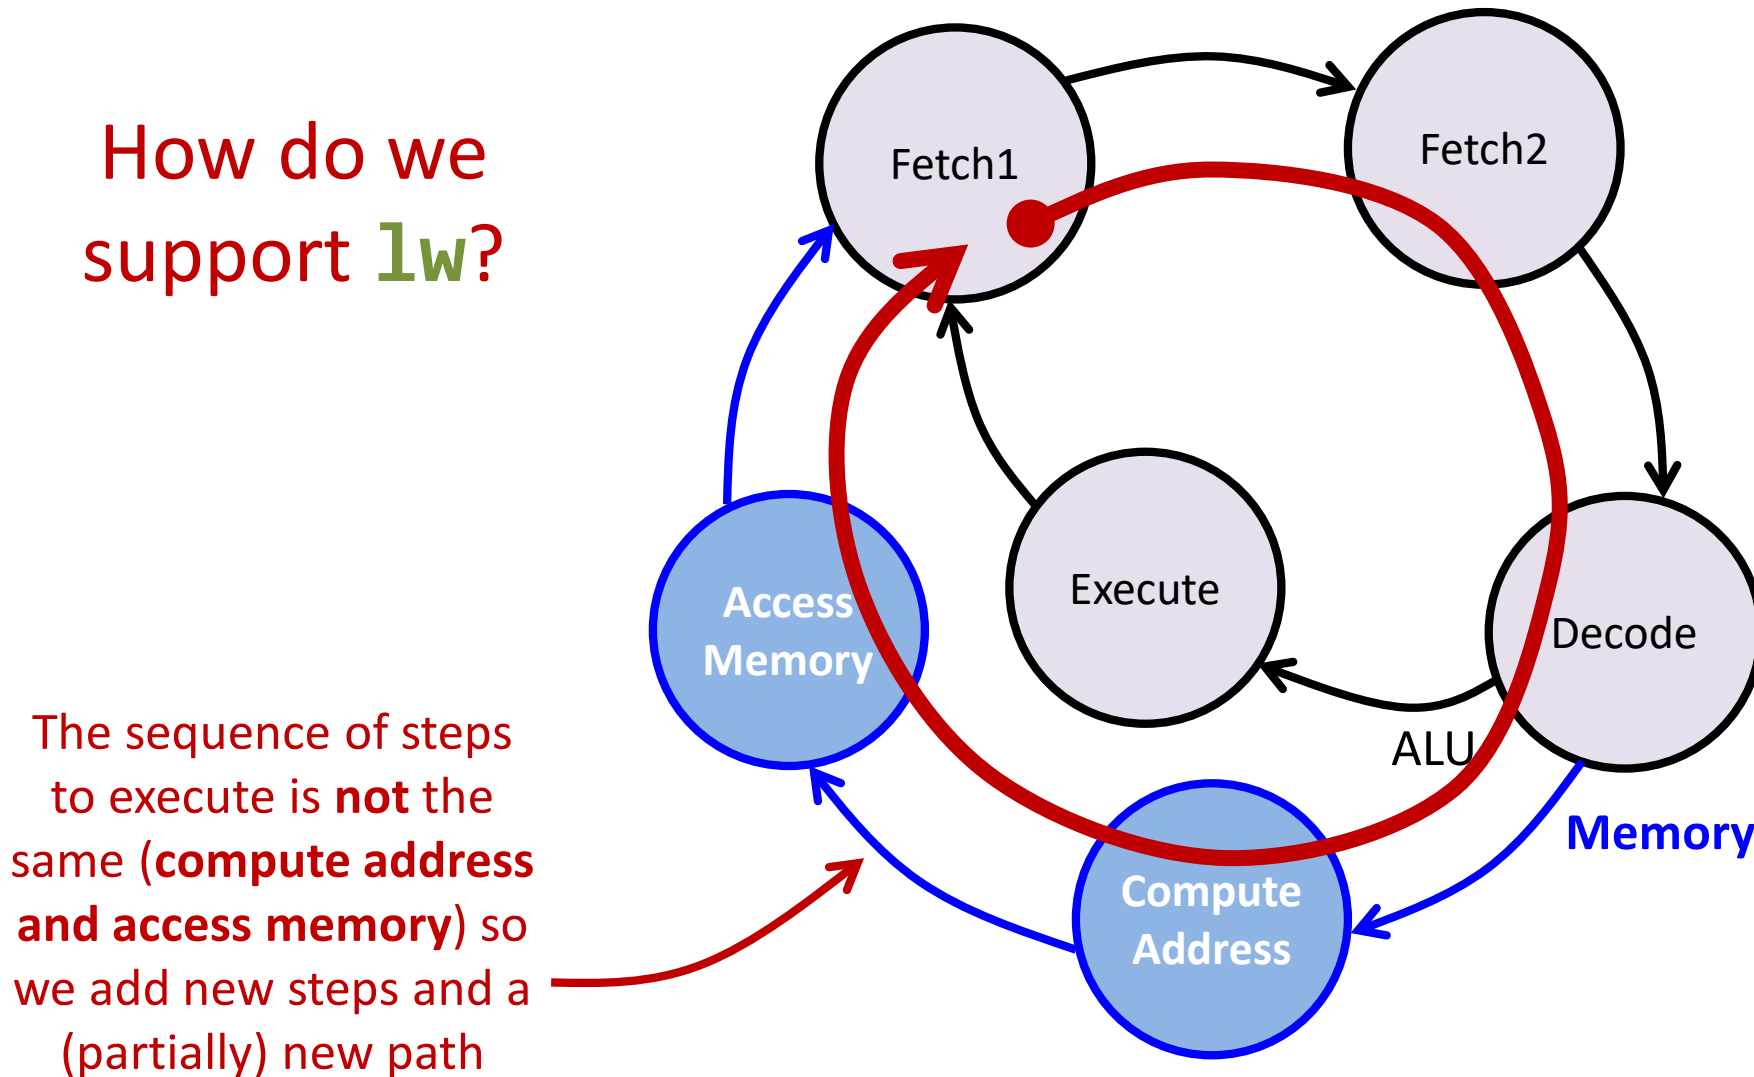
How do we support **add**?



If the **sequence of steps to execute is the same** (fetch the instruction, read registers, use the ALU, save result in the register), we just need to **an ALU that can perform additions**


Adding Instructions to a Multi-Cycle Processor

How do we support **lw**?




Adding Instructions to a Pipelined Processor

How do we support **lw**?



We need to accommodate the steps **compute address** and **access memory**

How?!



Latency of **add** is now **five** cycles!

The Importance of the ISA

Imagine that we want to have an instruction

`sub 8(t4), 0(t1), 0(t2)`

A CISC instruction!

Complex
Instruction-Set
Computer

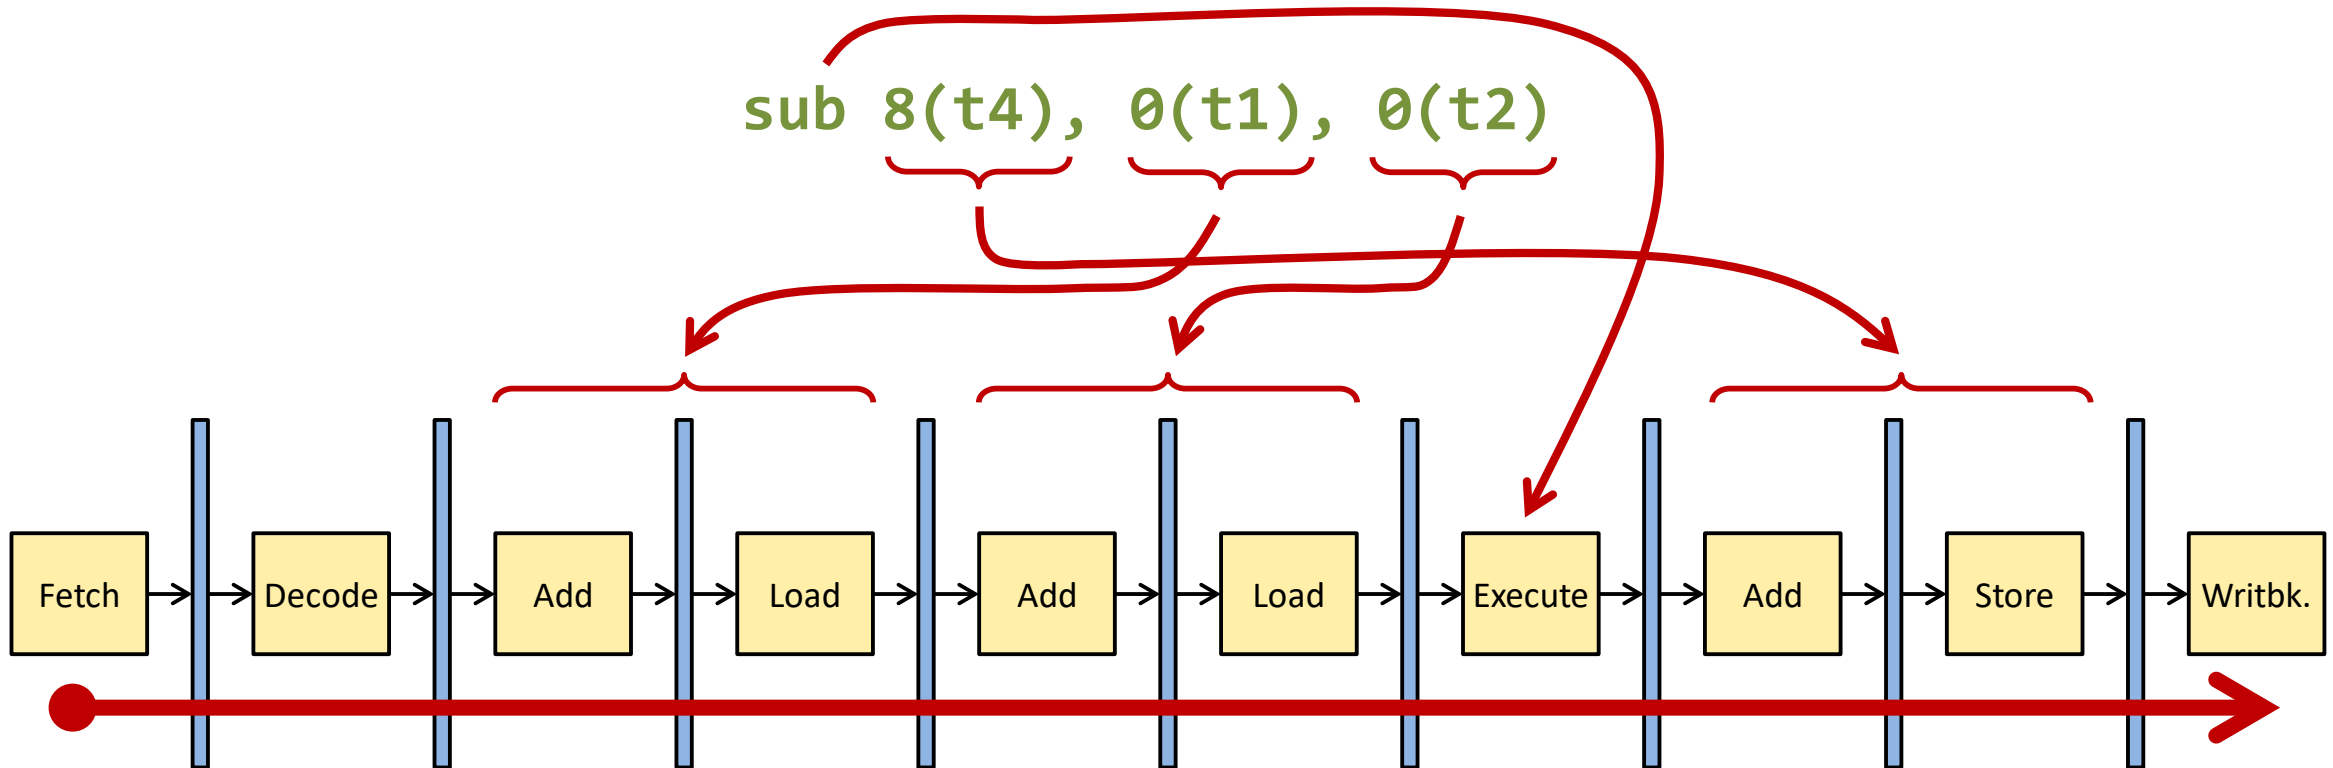
Read the value in memory
at the address $t2 + 0$

Subtract the value in memory at the address $t2 + 0$
from the value in memory at the address $t1 + 0$
and store the result in memory at the address $t4 + 8$

The Importance of the ISA

Imagine that we want to have an instruction

`sub 8(t4), 0(t1), 0(t2)`




`sub t4, t1, t2`

The Importance of the ISA

Imagine that we want to have an instruction


`sub 8(t4), 0(t1), 0(t2)`



Reduced Instruction-Set Computer

- Instead of imposing a **huge penalty to every simple instruction** by making complex instructions possible, let's **have only similarly simple instructions** and build our programs with those:

`sub 8(t4), 0(t1), 0(t2)`



```
lw    t3, 0(t1)
lw    t5, 0(t2)
sub   t3, t3, t5
sw    t3, 8(t4)
```

- It turns out that it is not the only way to go, but it is a **good one** and we will follow it...

Two Main Problems

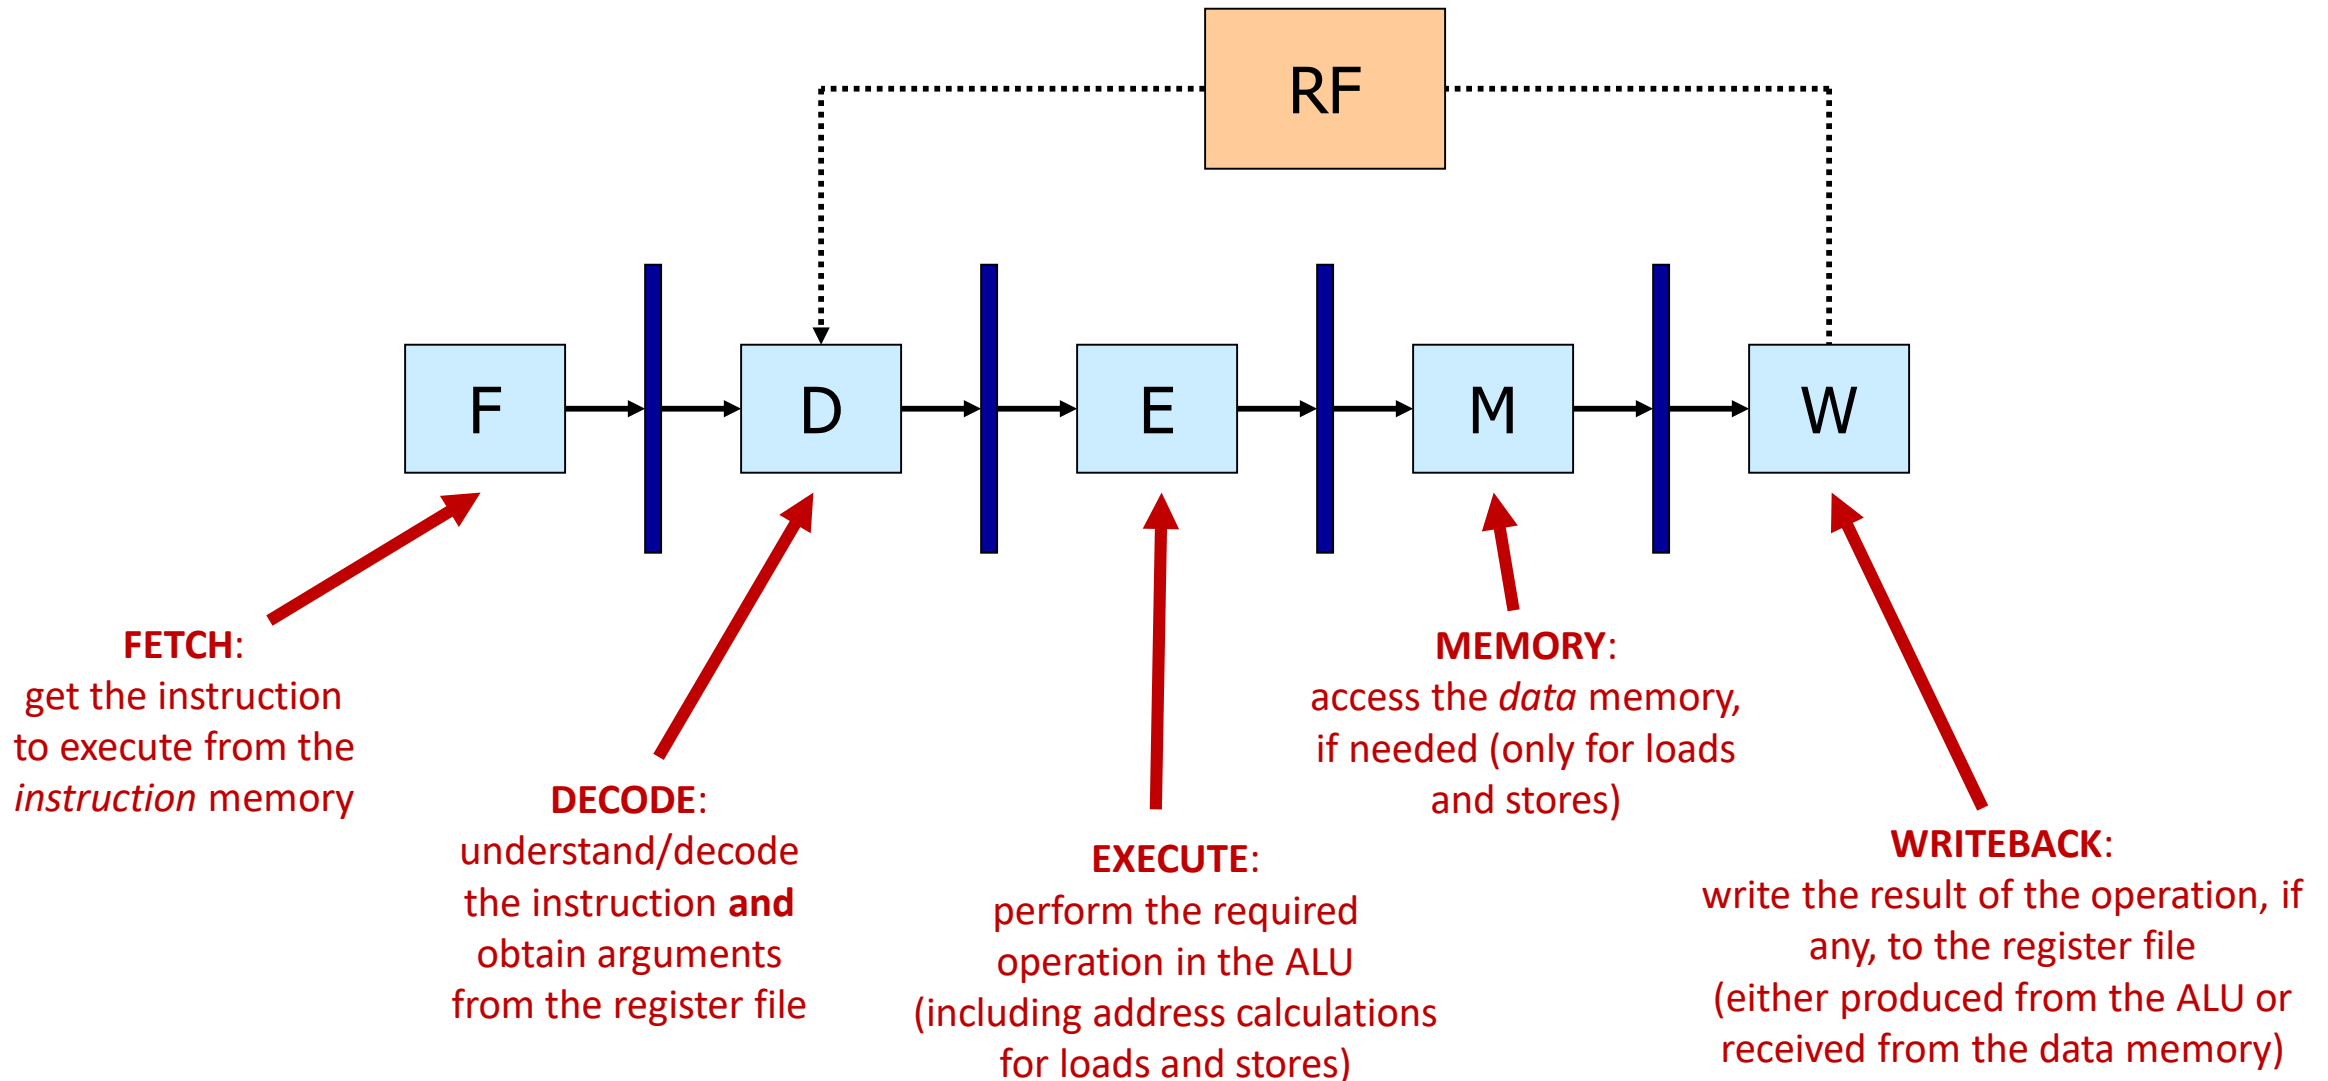
1. CISC vs. **RISC**

- Can we **build equally well a pipeline** for a Complex Instruction Set Computer as for a Reduced Instruction Set Computer?
- What does that mean?!


2. Instructions are **not** independent

- Can we execute code **correctly**?

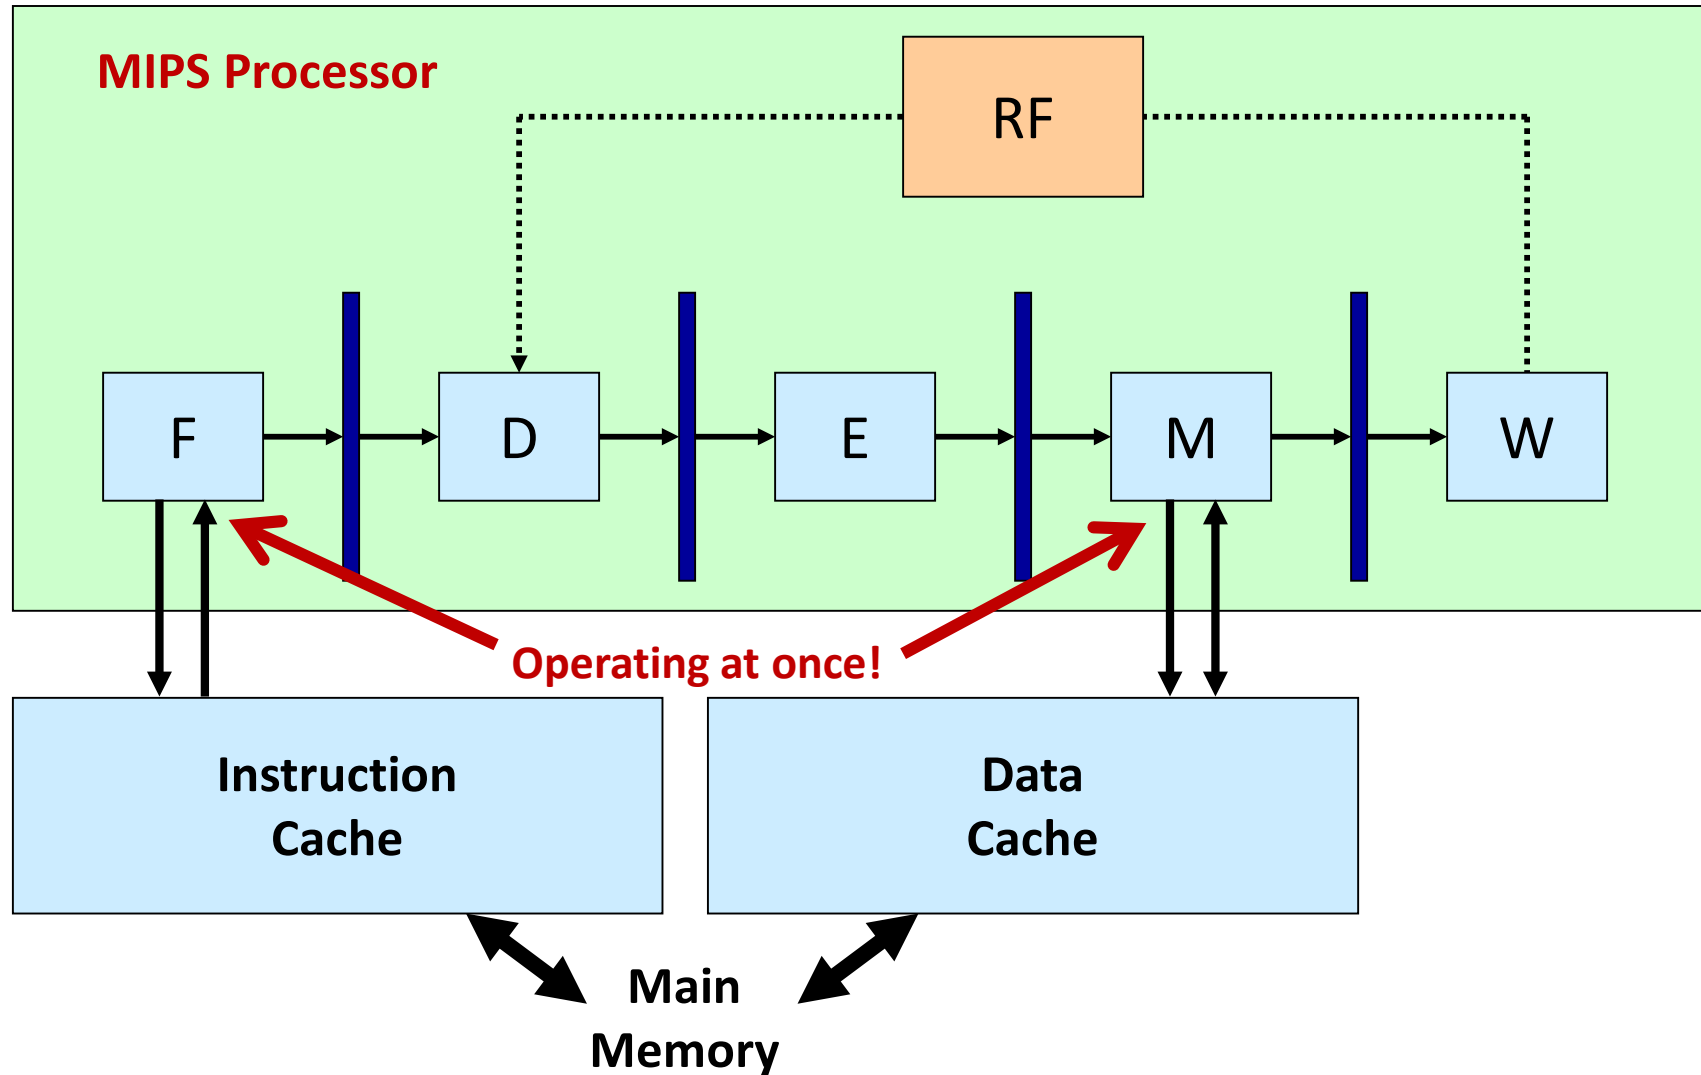
Simple 5-Stage MIPS Pipeline



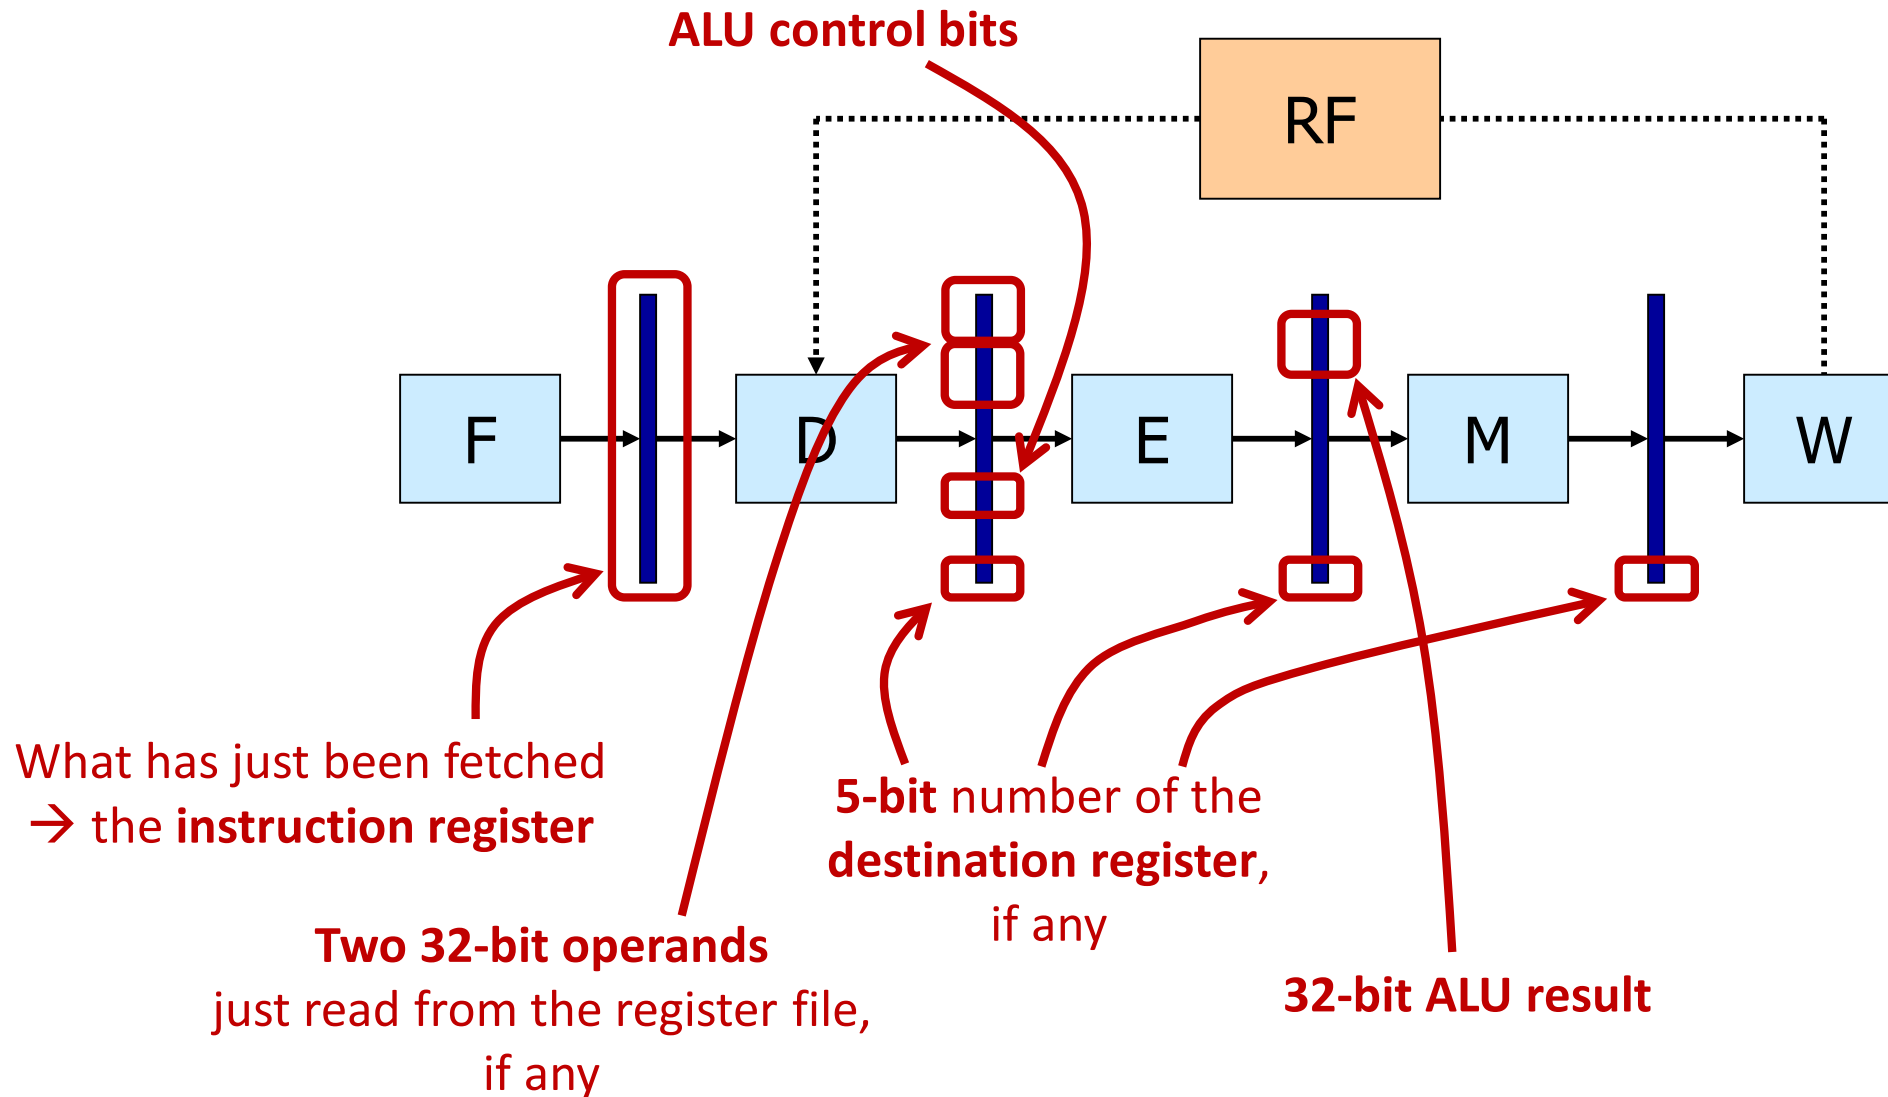
The Laundry Metaphor



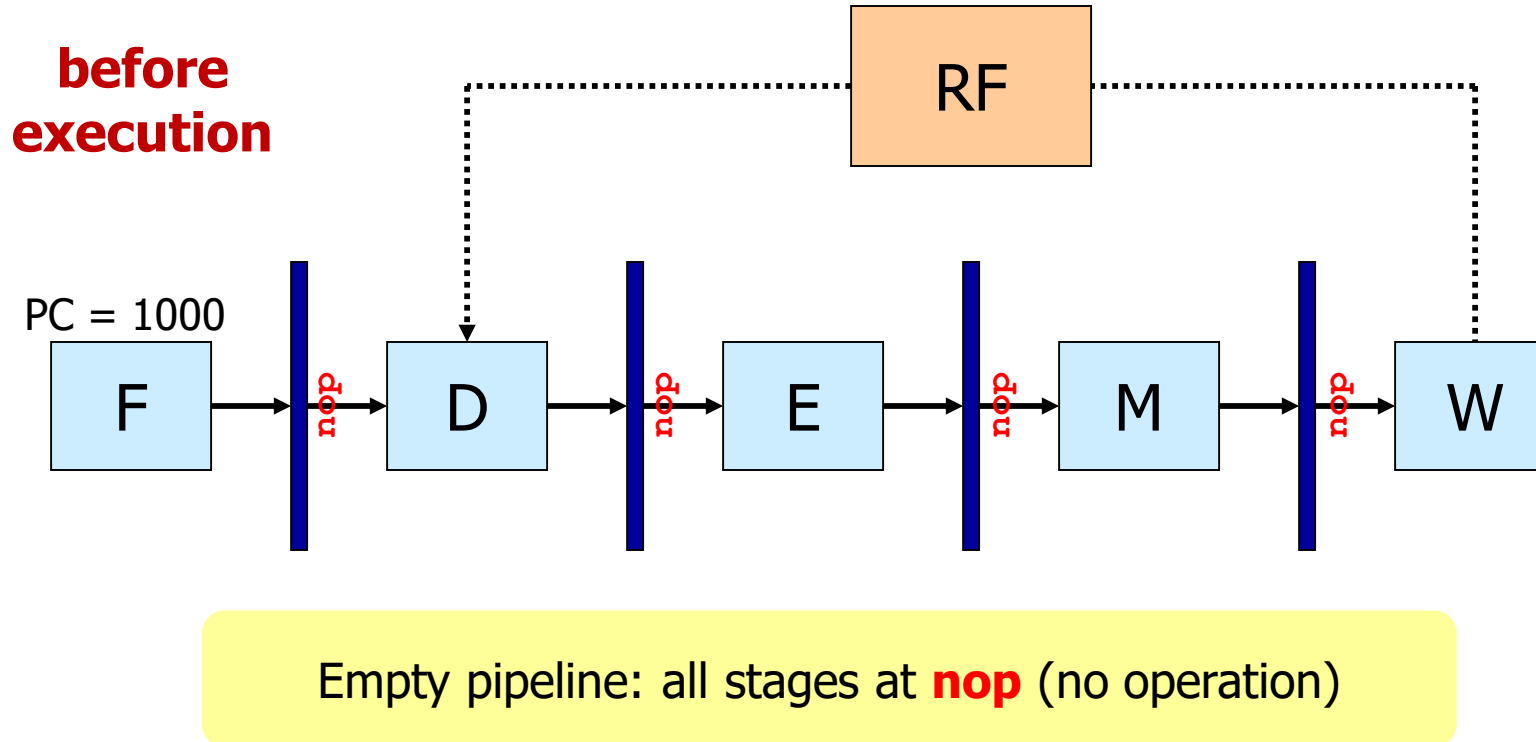
Two Distinct Memory Interfaces



What Is in the Pipeline Registers?

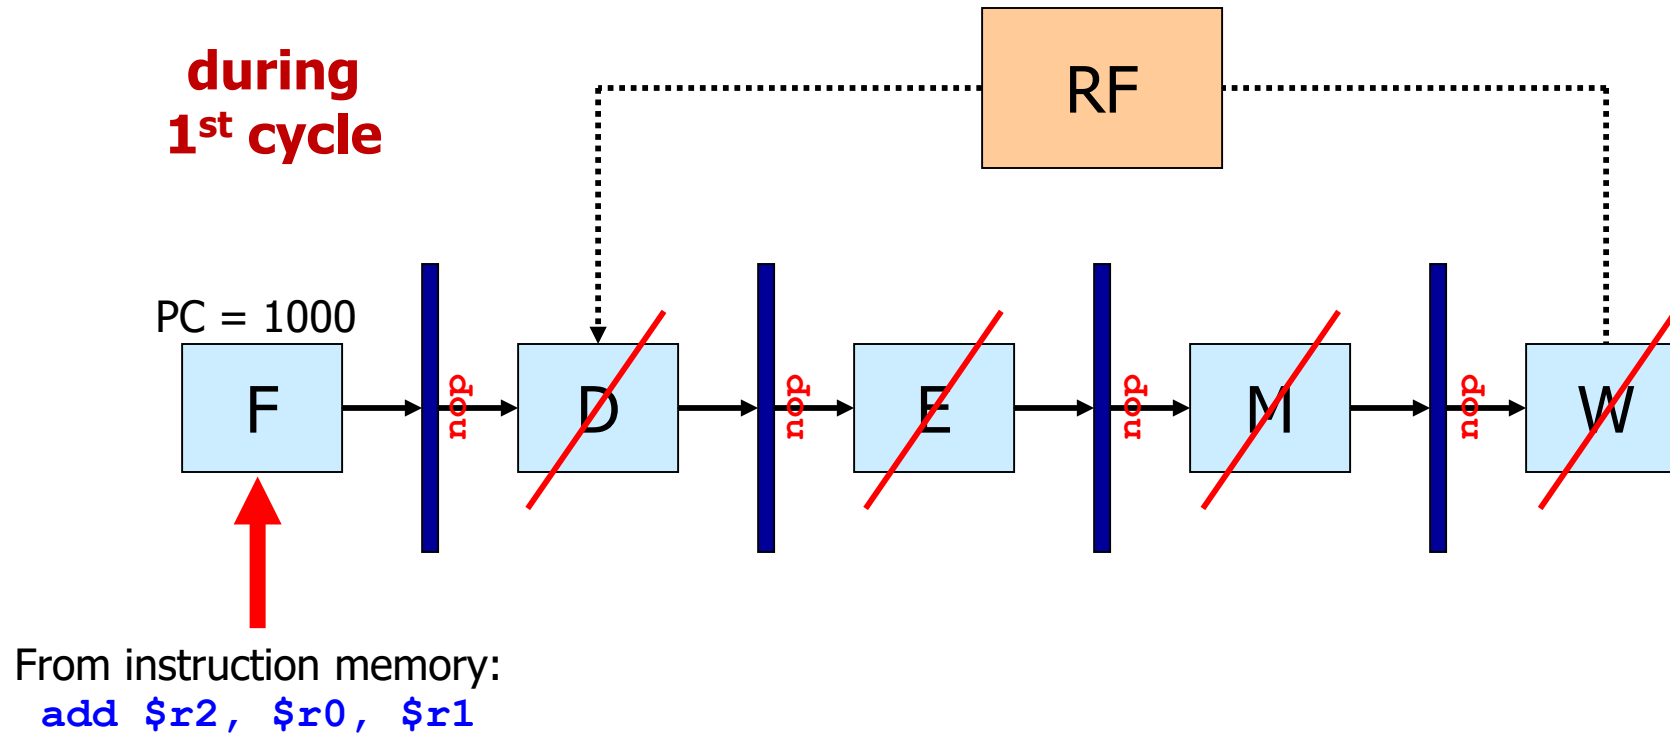


Example of Pipelined Execution



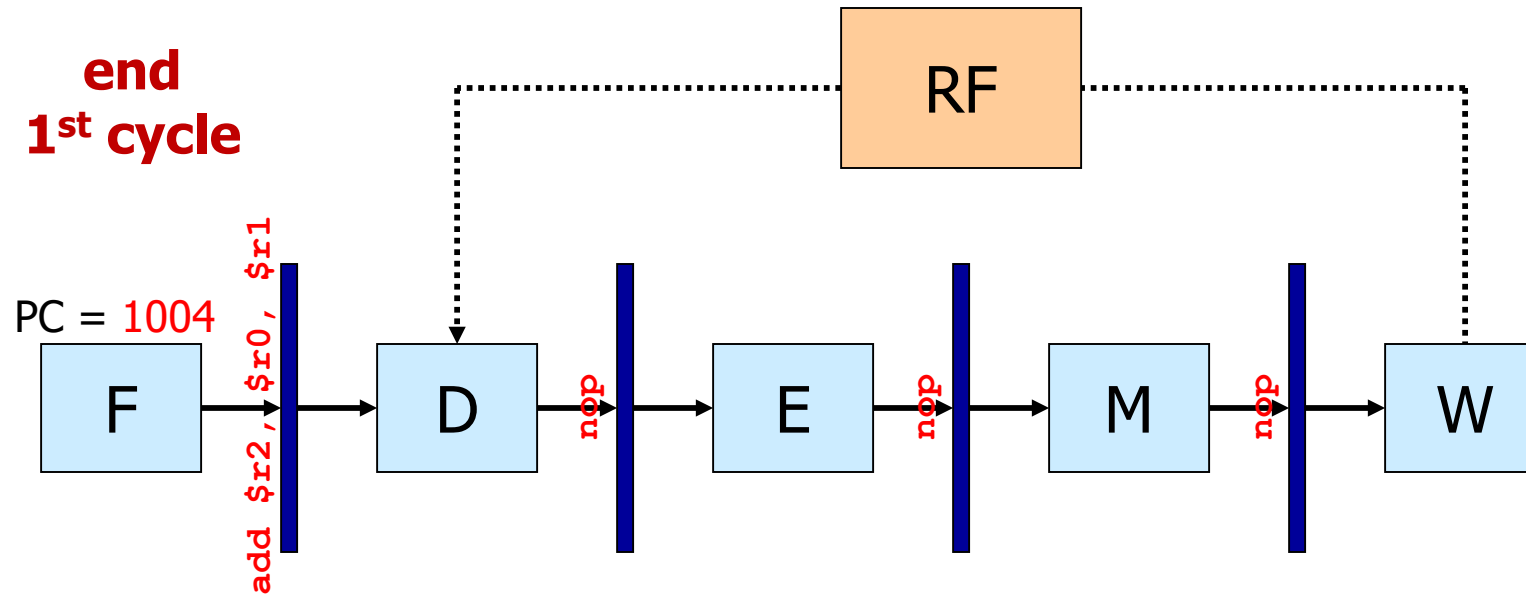
```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

Example of Pipelined Execution



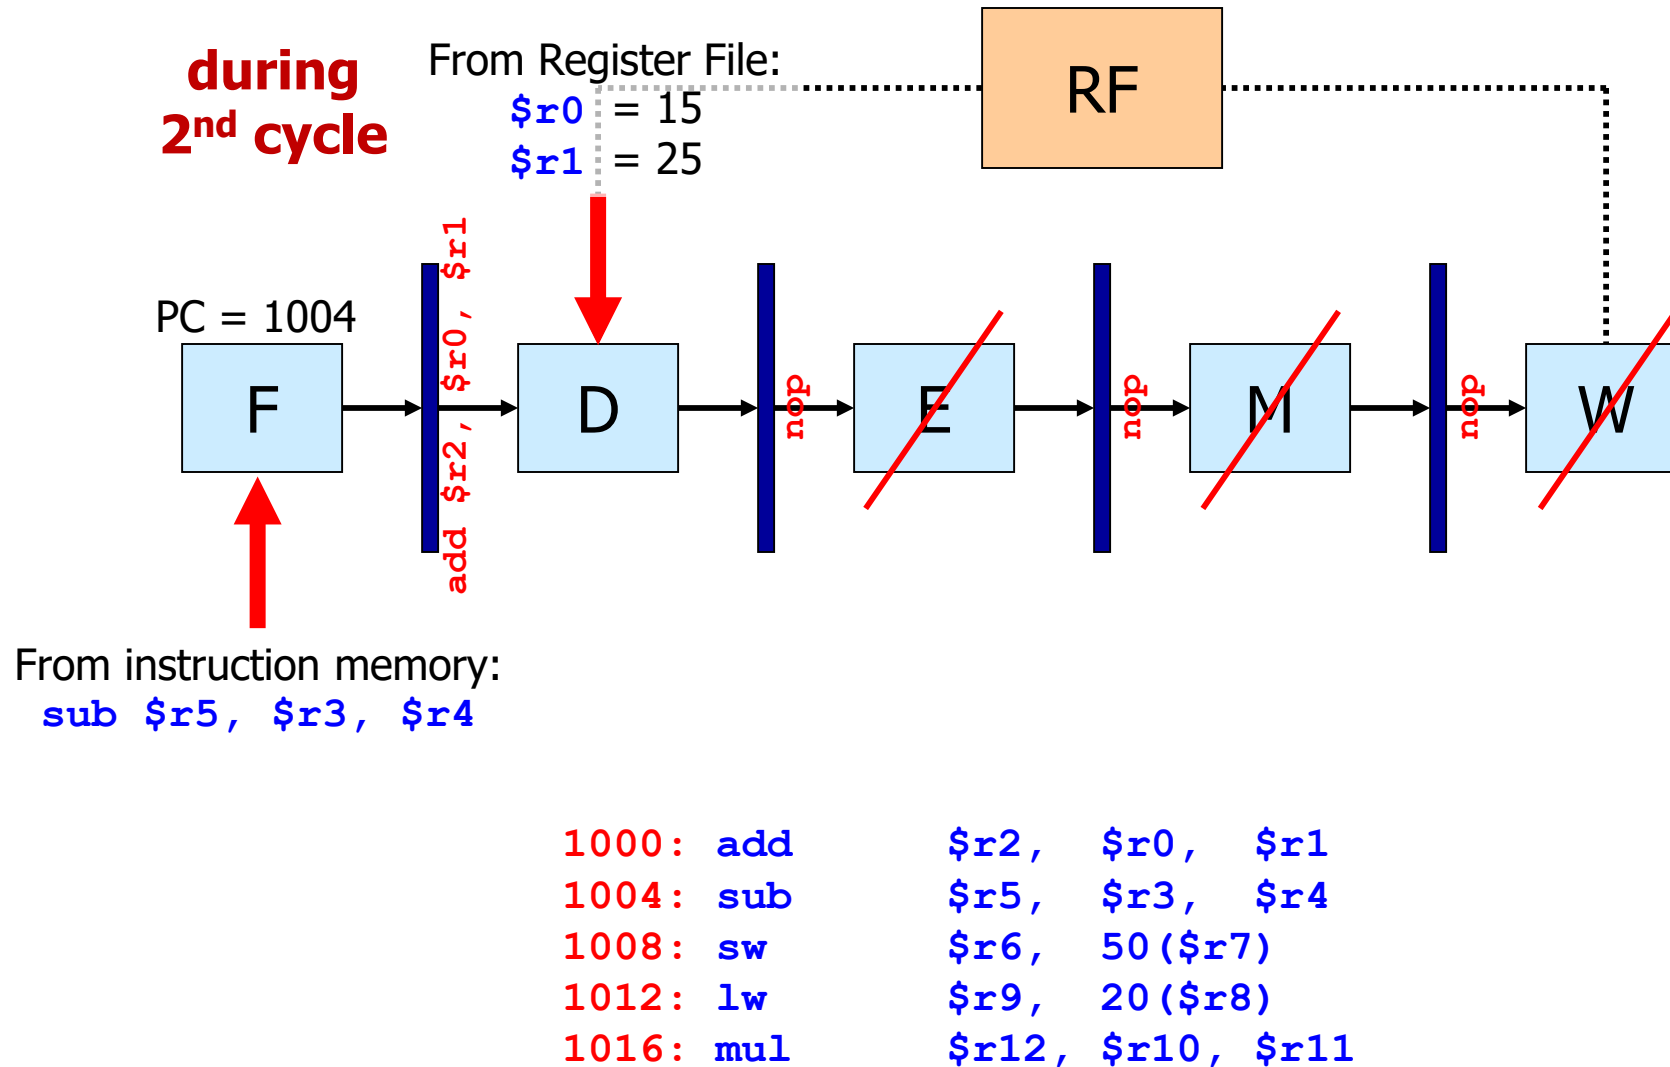
```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```


Example of Pipelined Execution

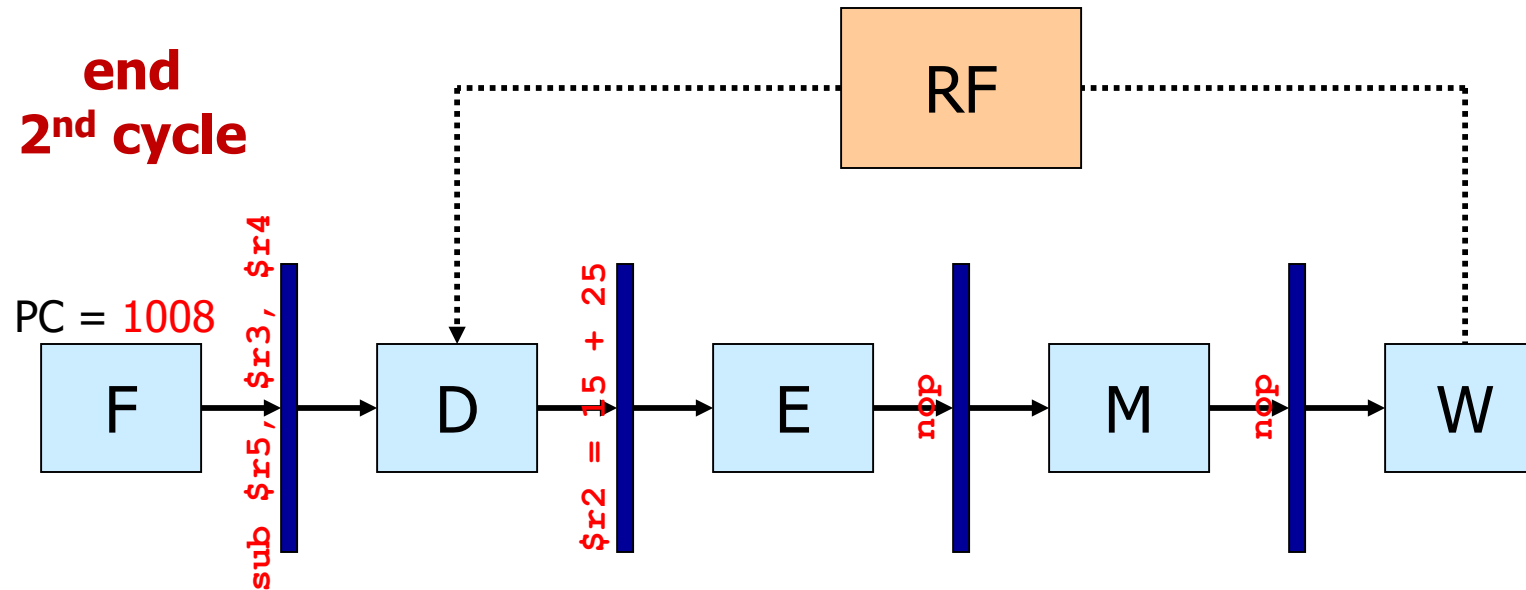


```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

Example of Pipelined Execution

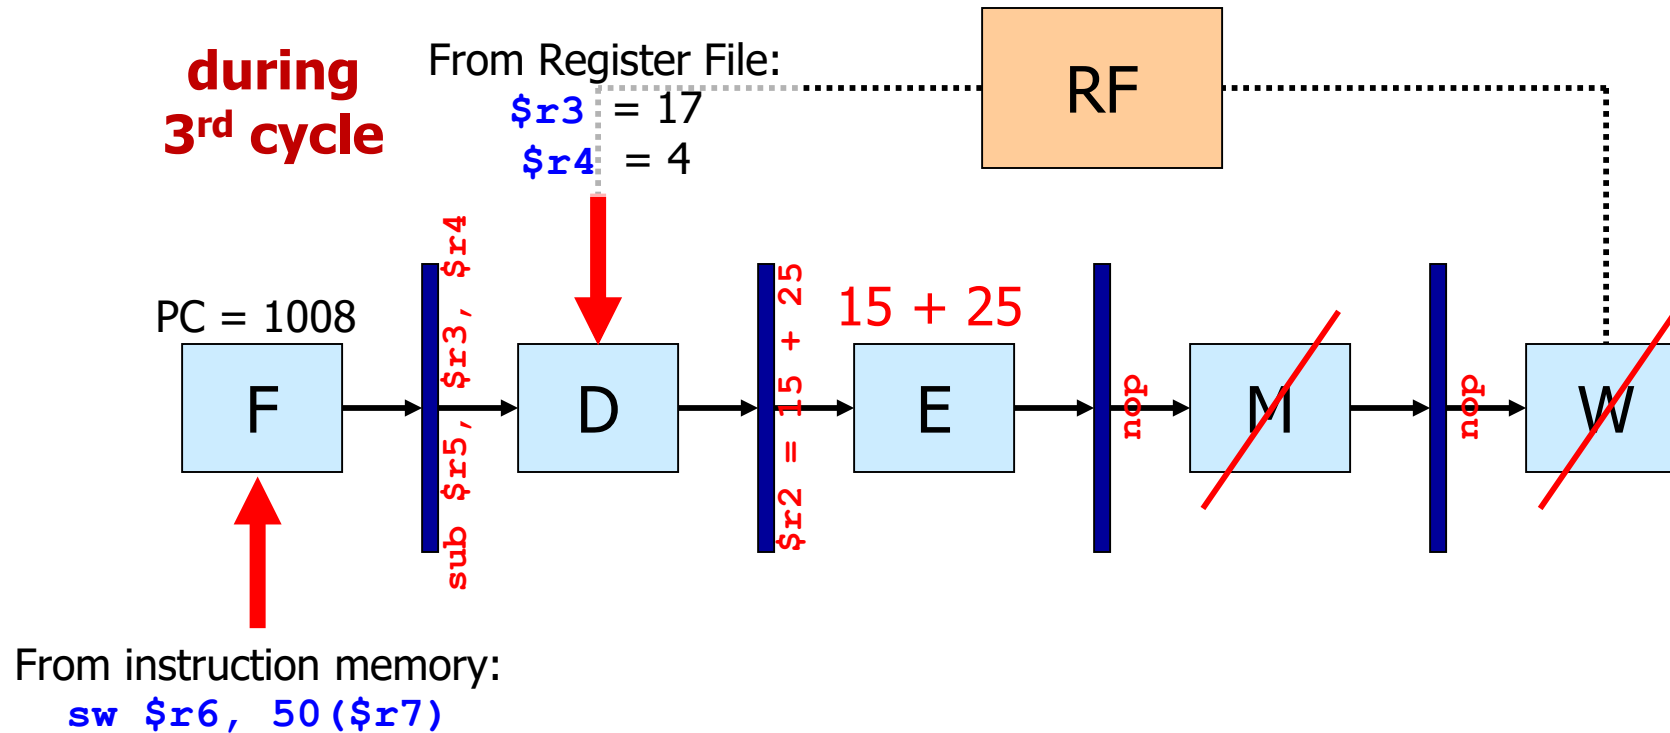


Example of Pipelined Execution



```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

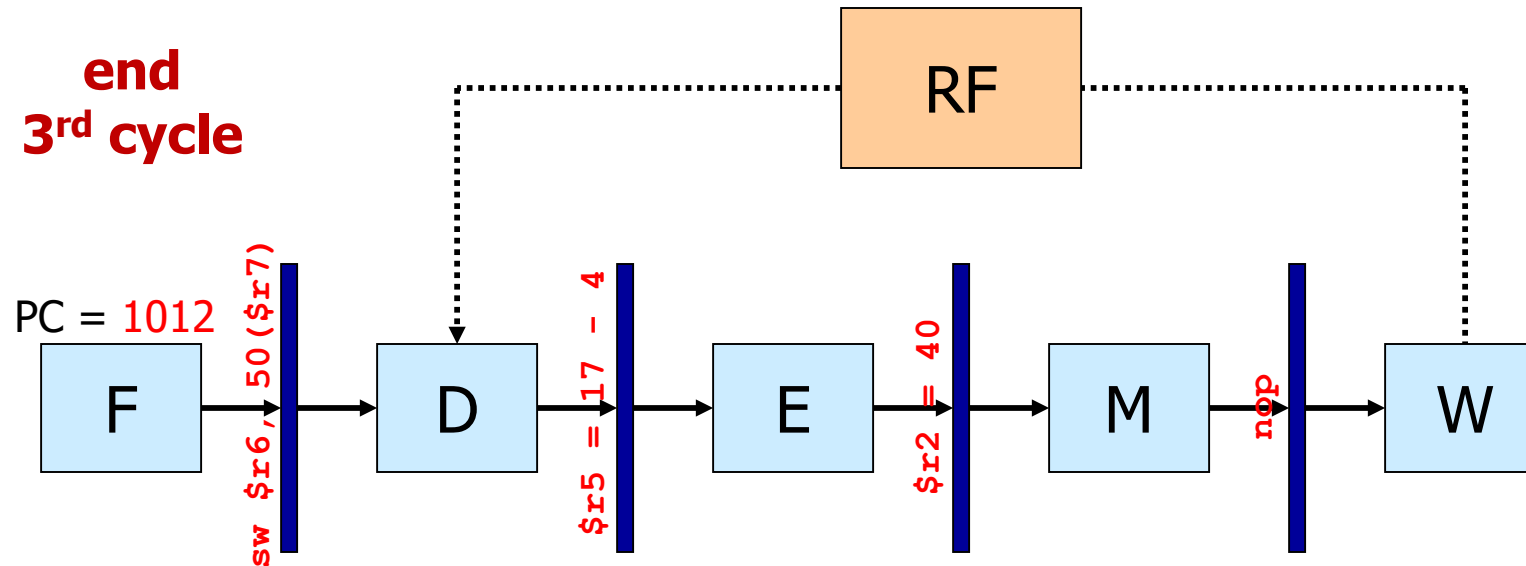
Example of Pipelined Execution



```

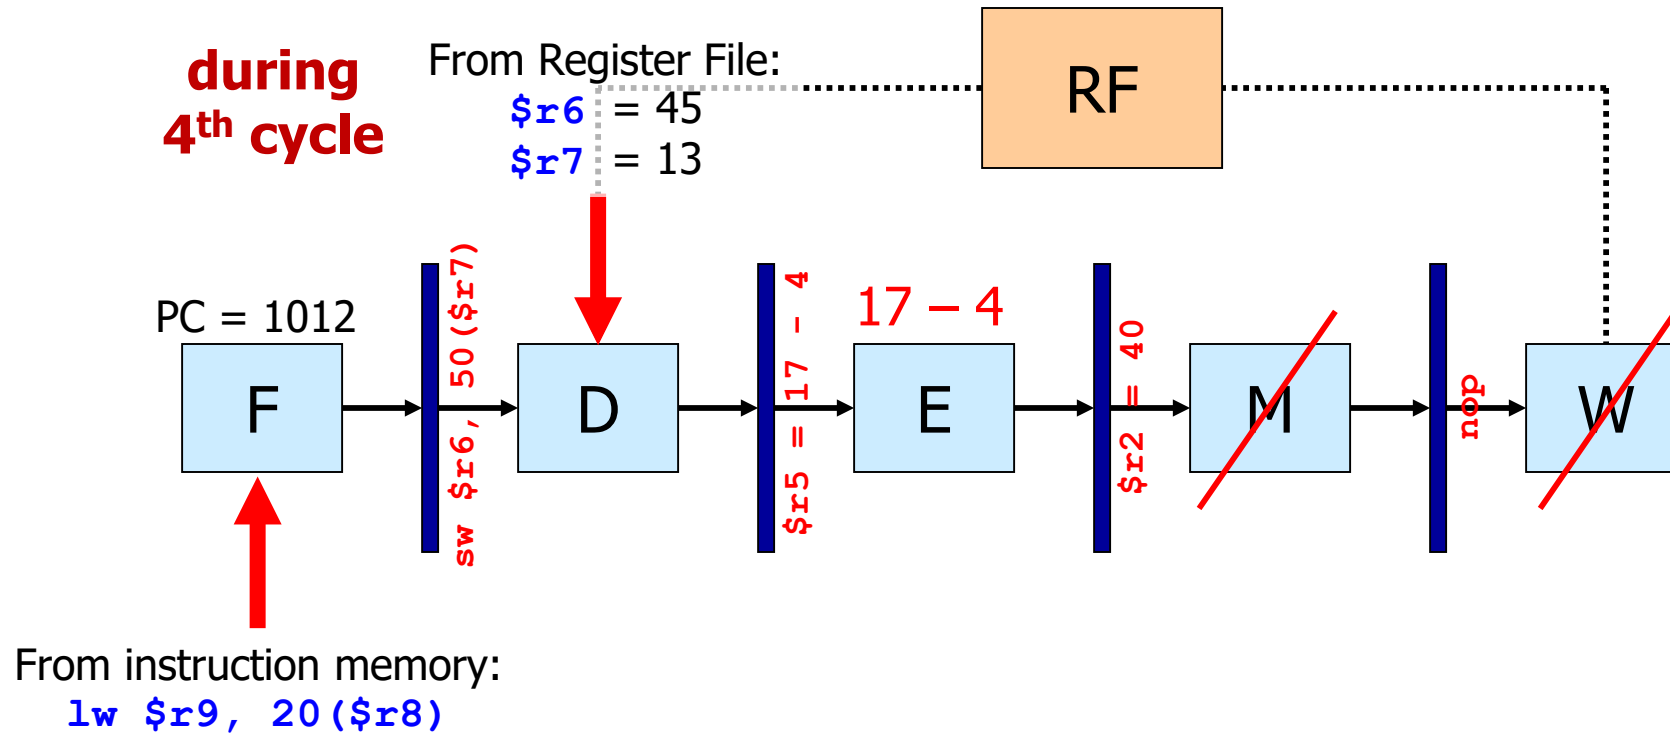
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
    
```

Example of Pipelined Execution



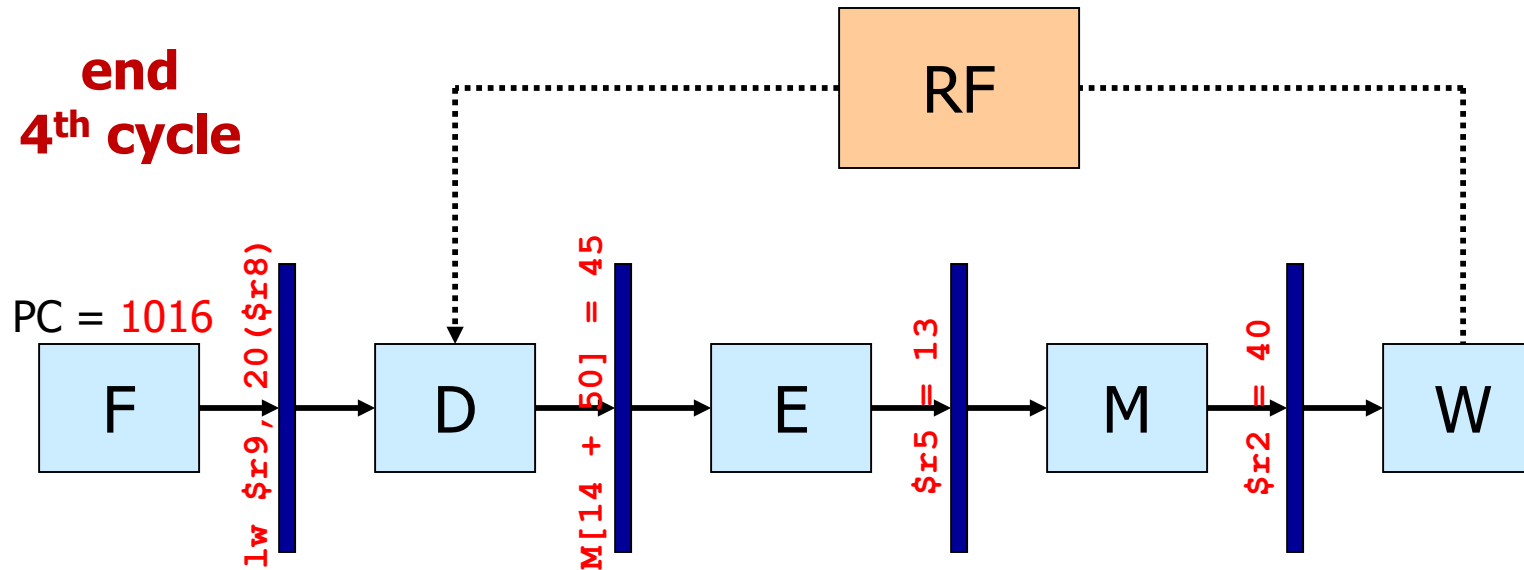
```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

Example of Pipelined Execution



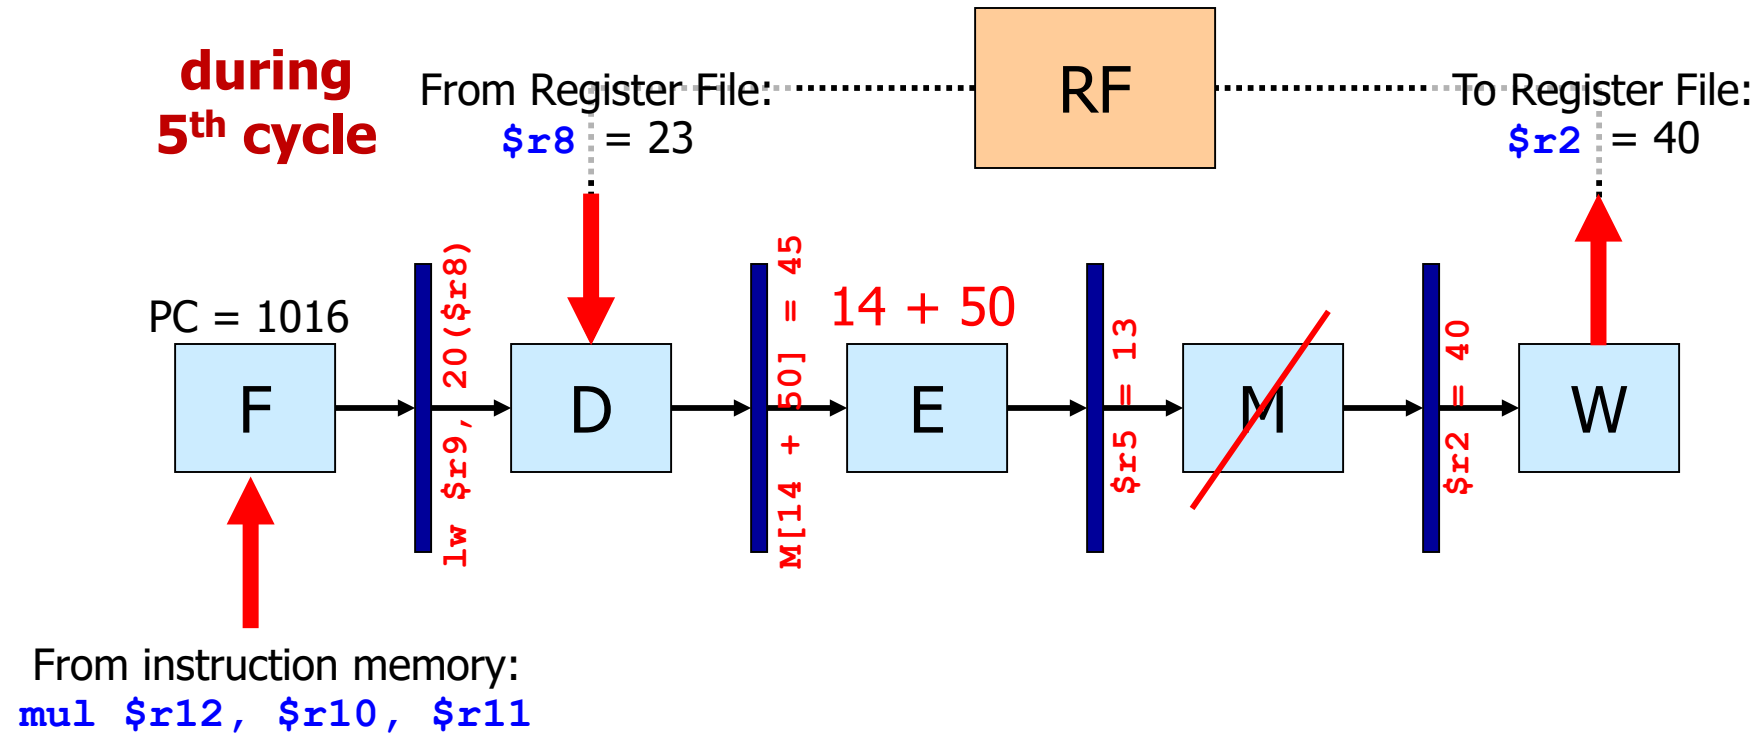
1000:	add	\$r2,	\$r0,	\$r1
1004:	sub	\$r5,	\$r3,	\$r4
1008:	sw	\$r6,	50(\$r7)	
1012:	lw	\$r9,	20(\$r8)	
1016:	mul	\$r12,	\$r10,	\$r11

Example of Pipelined Execution



```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

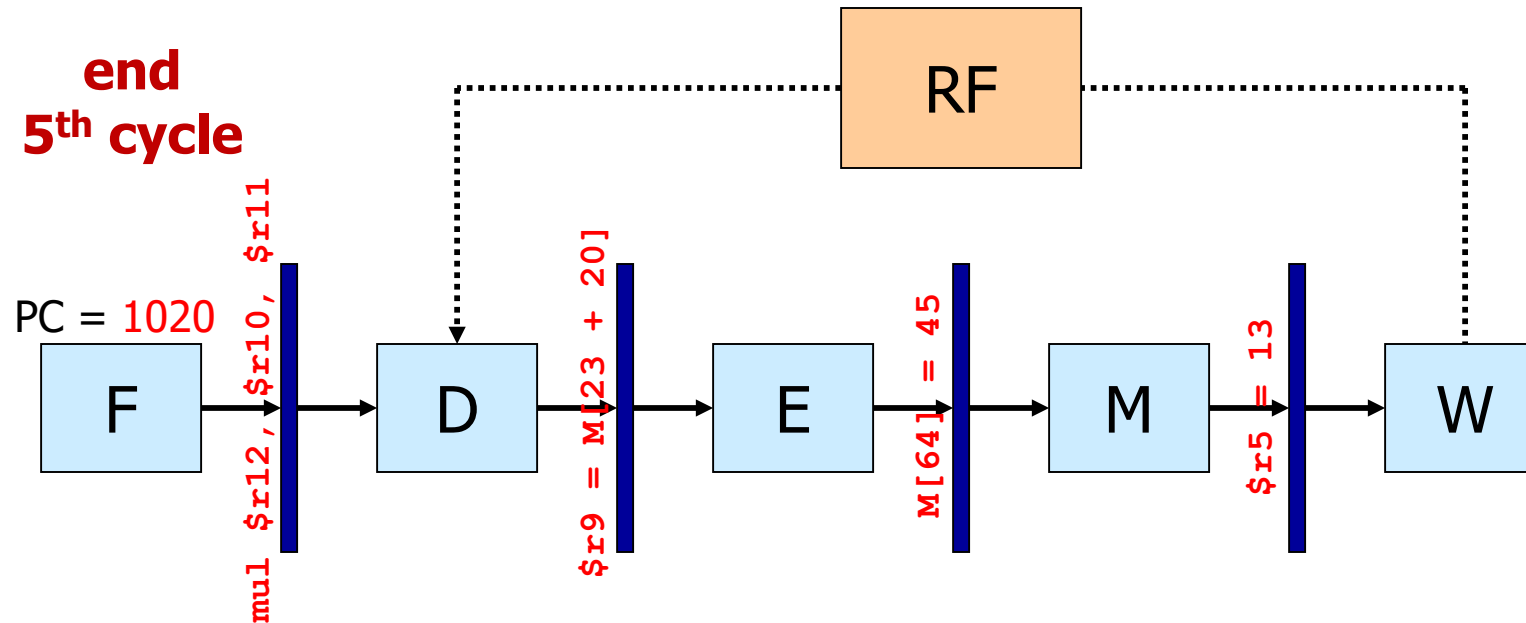
Example of Pipelined Execution



```

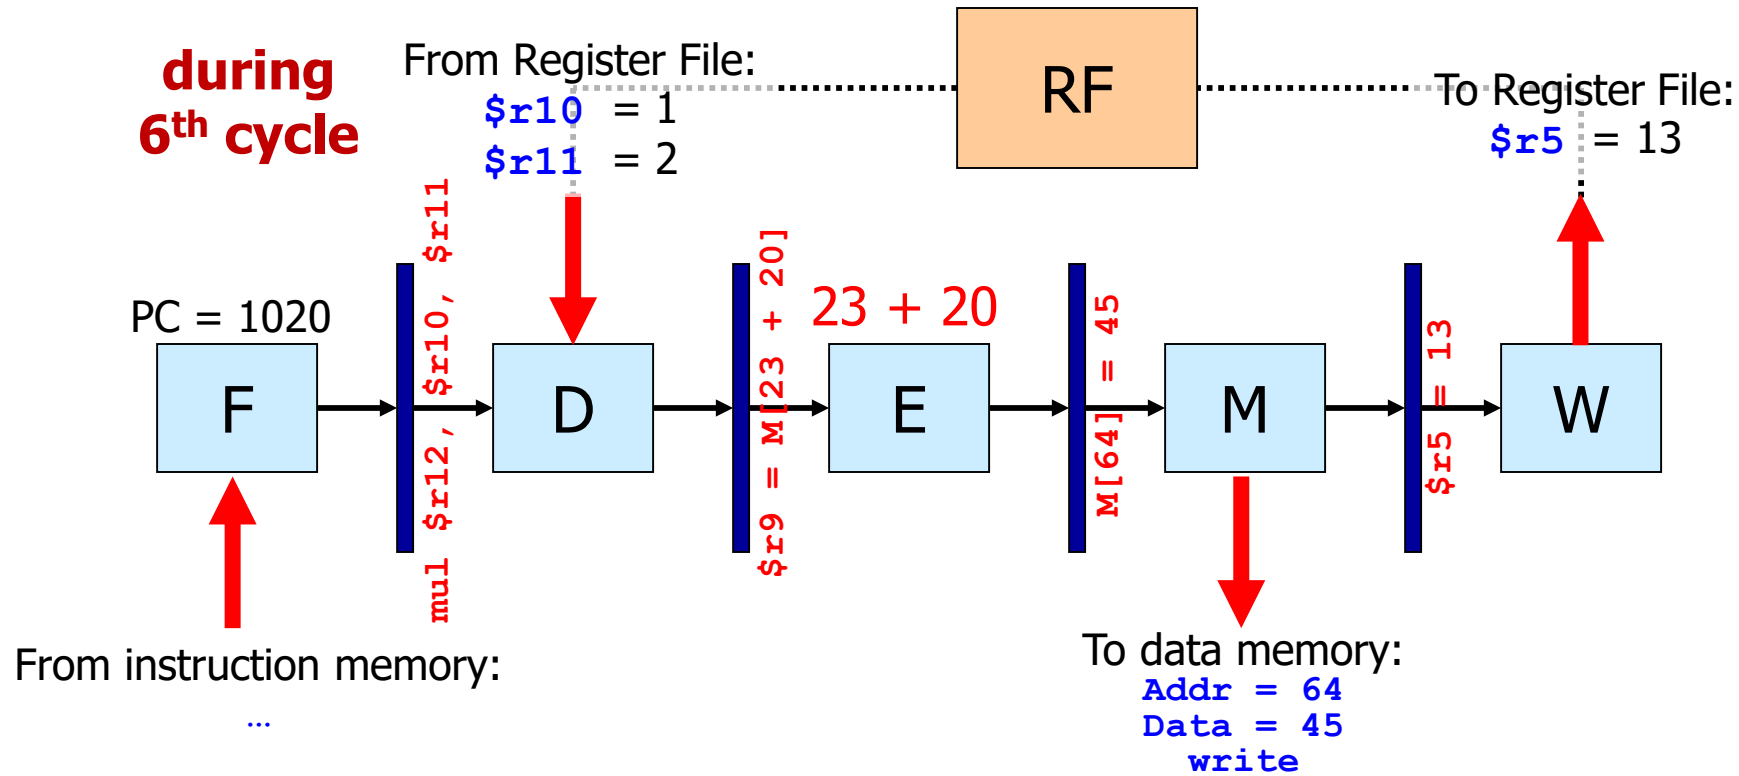
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
    
```


Example of Pipelined Execution



```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

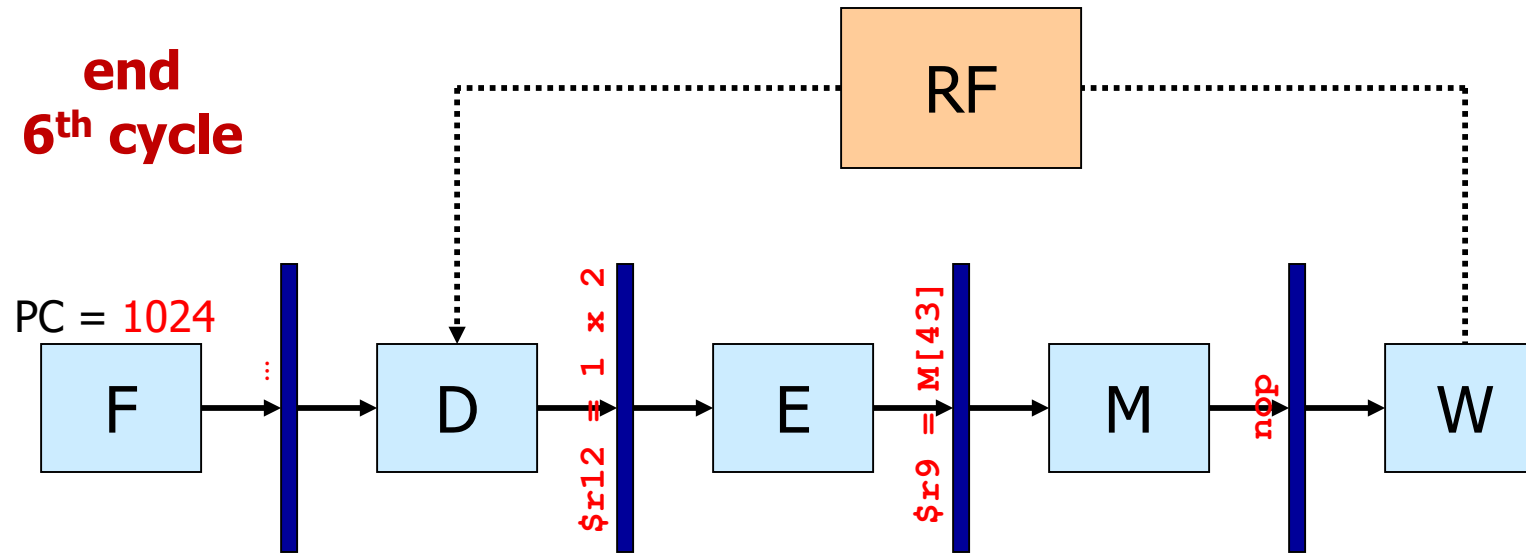
Example of Pipelined Execution



```

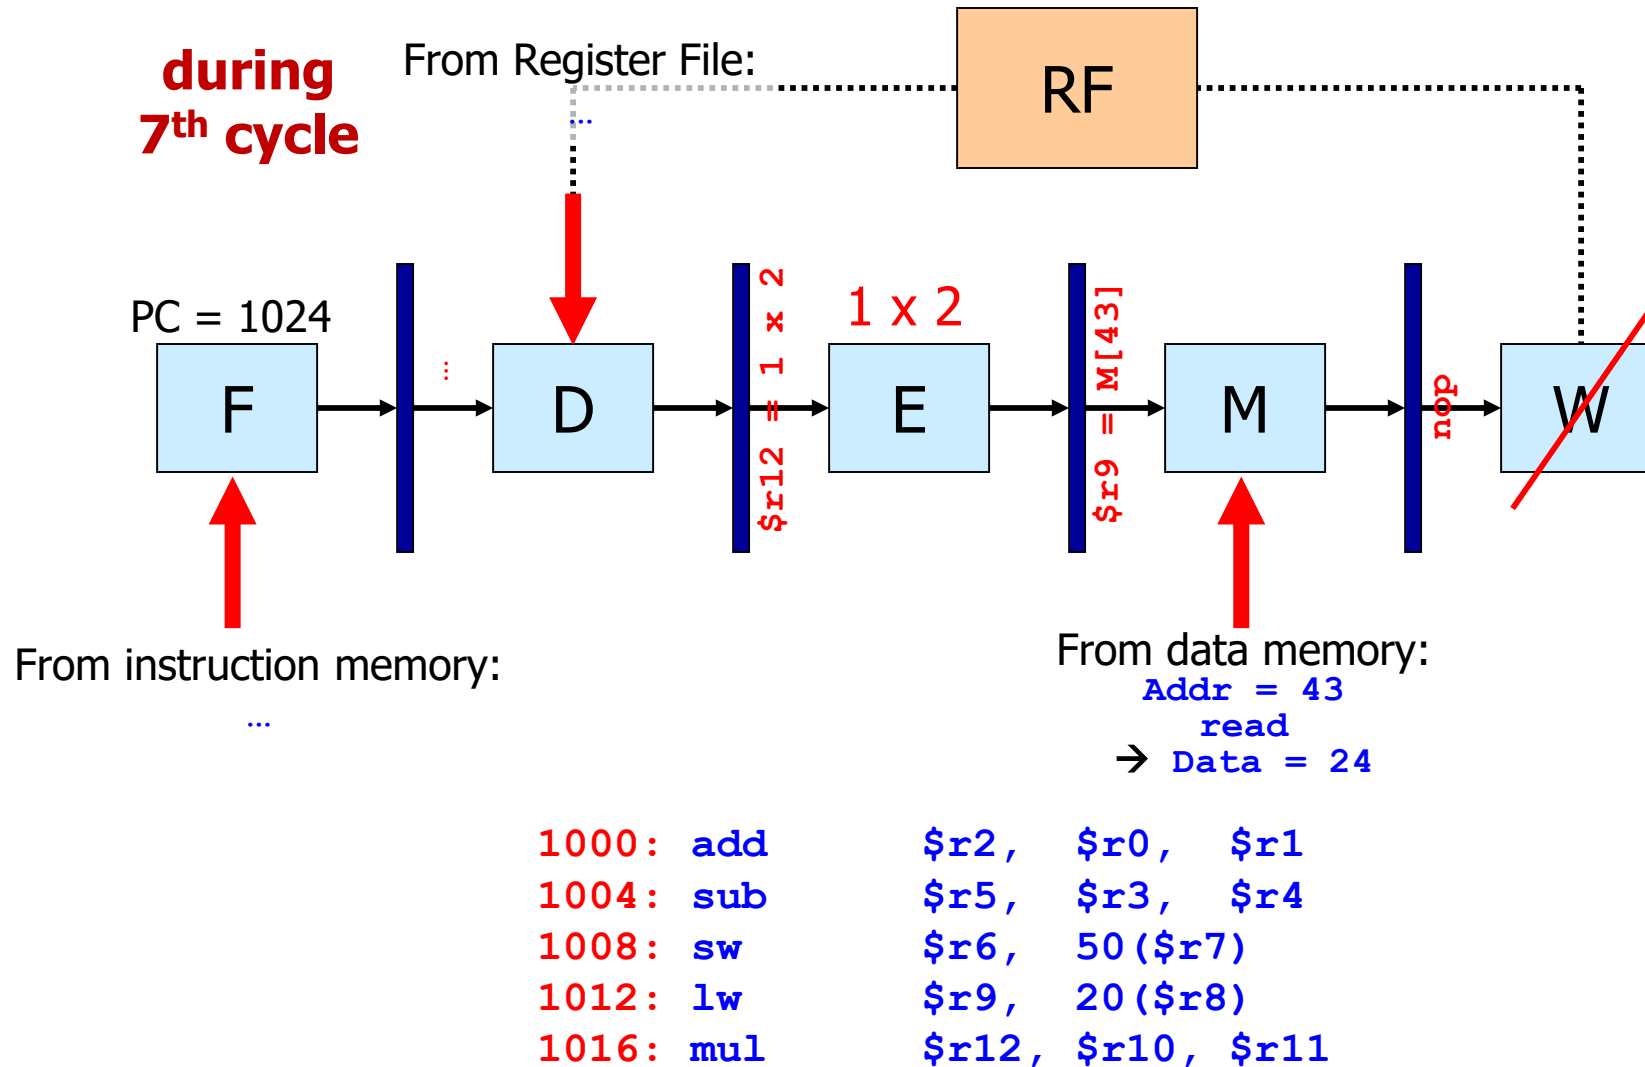
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
    
```

Example of Pipelined Execution

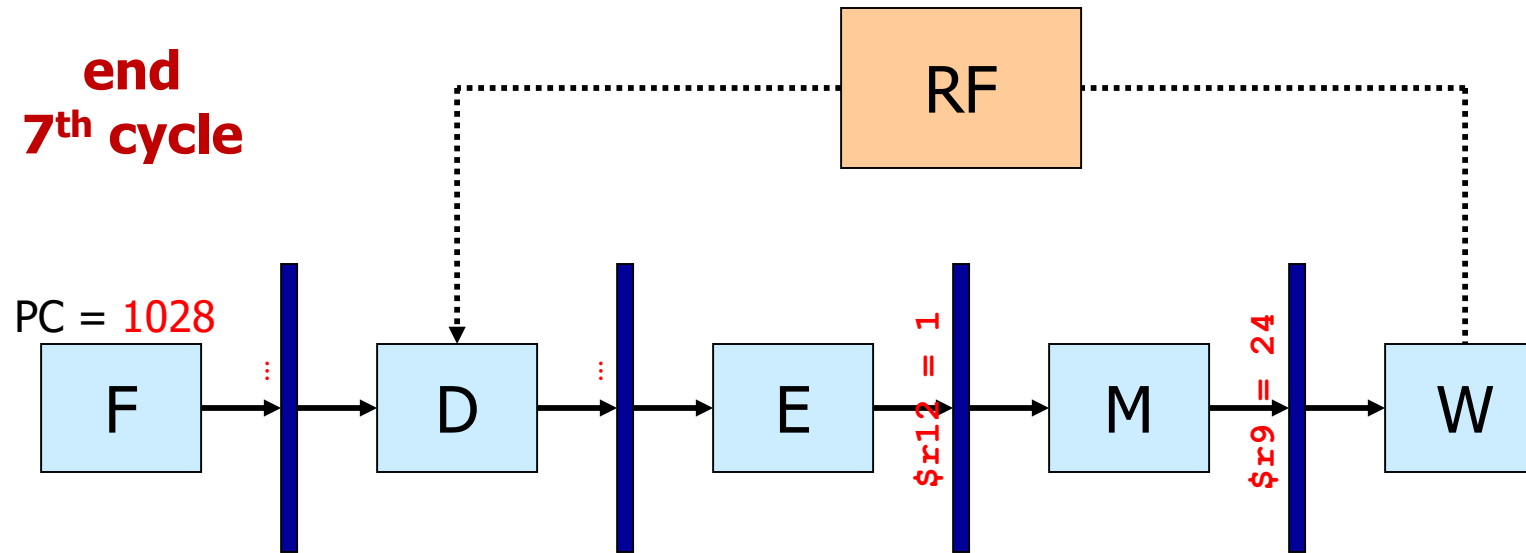


```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

Example of Pipelined Execution

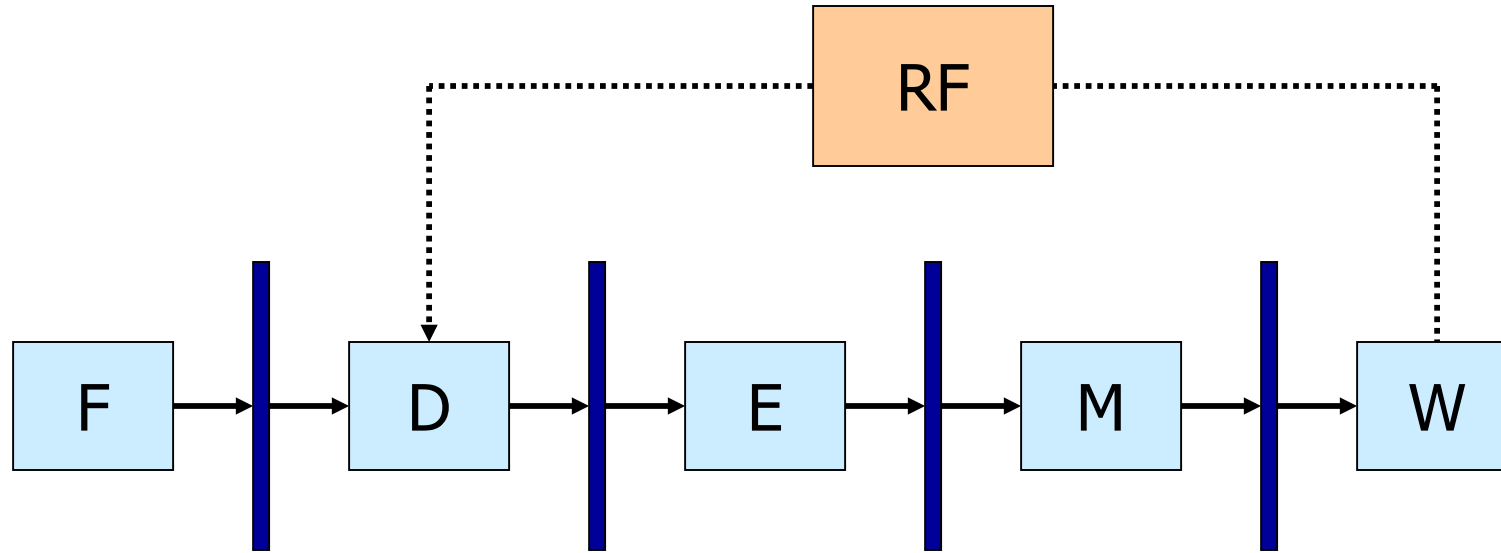


Example of Pipelined Execution



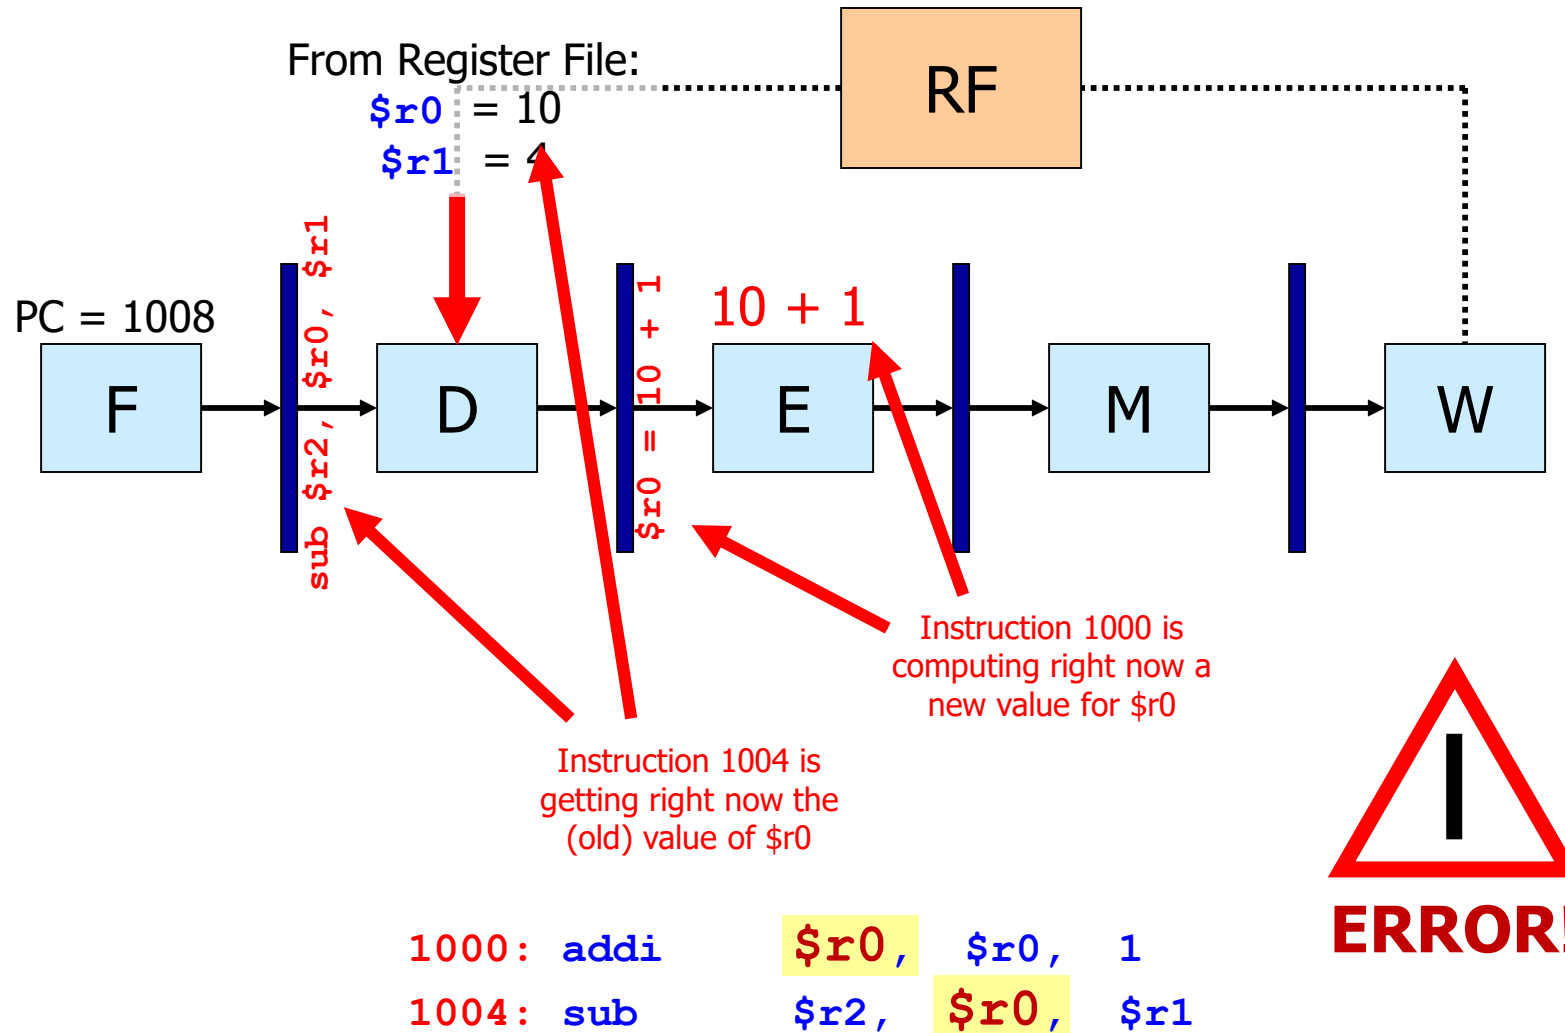
```
1000: add    $r2, $r0, $r1
1004: sub    $r5, $r3, $r4
1008: sw     $r6, 50($r7)
1012: lw     $r9, 20($r8)
1016: mul    $r12, $r10, $r11
```

Pipeline Schedule



1000:	F	D	E	M	W				
1004:		F	D	E	M	W			
1008:			F	D	E	M	W		
1012:				F	D	E	M	W	
1016:					F	D	E	M	W

Problem!



Two Main Problems

1. CISC vs. **RISC**

- Can we **build equally well a pipeline** for a Complex Instruction Set Computer as for a Reduced Instruction Set Computer?
- What does that mean?!

2. Instructions are **not** independent

- Can we execute code **correctly**?

RAW, WAR and WAW Dependences

Write After Read


```
divd    $f0, $f1, $f2
add     $f3, $f0, $f4
subd    $f4, $f5, $f6
adddi   $f0, $f5, 10
```

- **add** has a **RAW** dependence on **divd**
- **subd** has a **WAR** dependence on **add**
- **adddi** has a **WAW** dependence on **divd**

True or **data**
dependencies

Name
dependencies


Data Hazards



Causality violation!

We try to use a result before it is produced!

Data Hazards Solved by Stalling the Pipeline



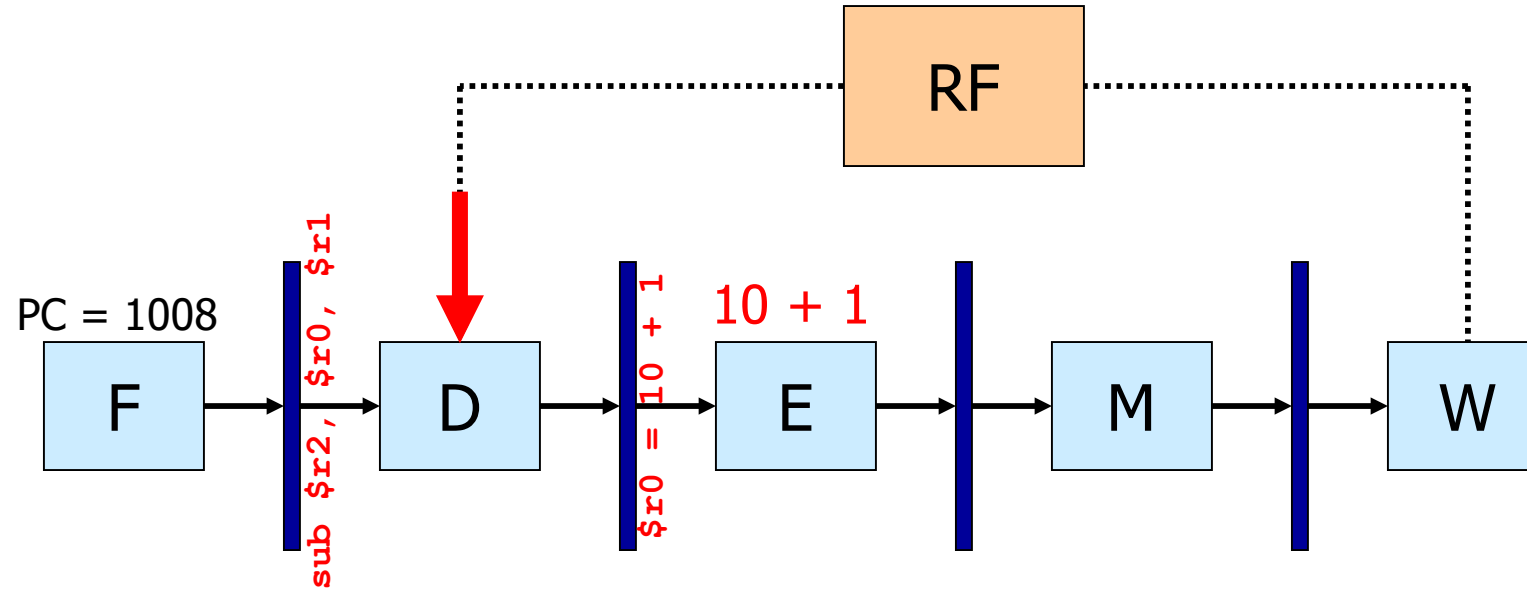
- The natural solution to **Data Hazards** caused by **RAW** dependences is to implement some logic in the processor to stop/repeat the decoding until the required value is available
- “**Stalling**” roughly means introducing **nop's** in the pipeline
- Due to the rigidity of the pipeline, if one stage is stalled (**D** in the example), all the preceding ones must be stalled too (e.g., **F**)

1000: addi

\$r0, \$r0, 1

1004: sub

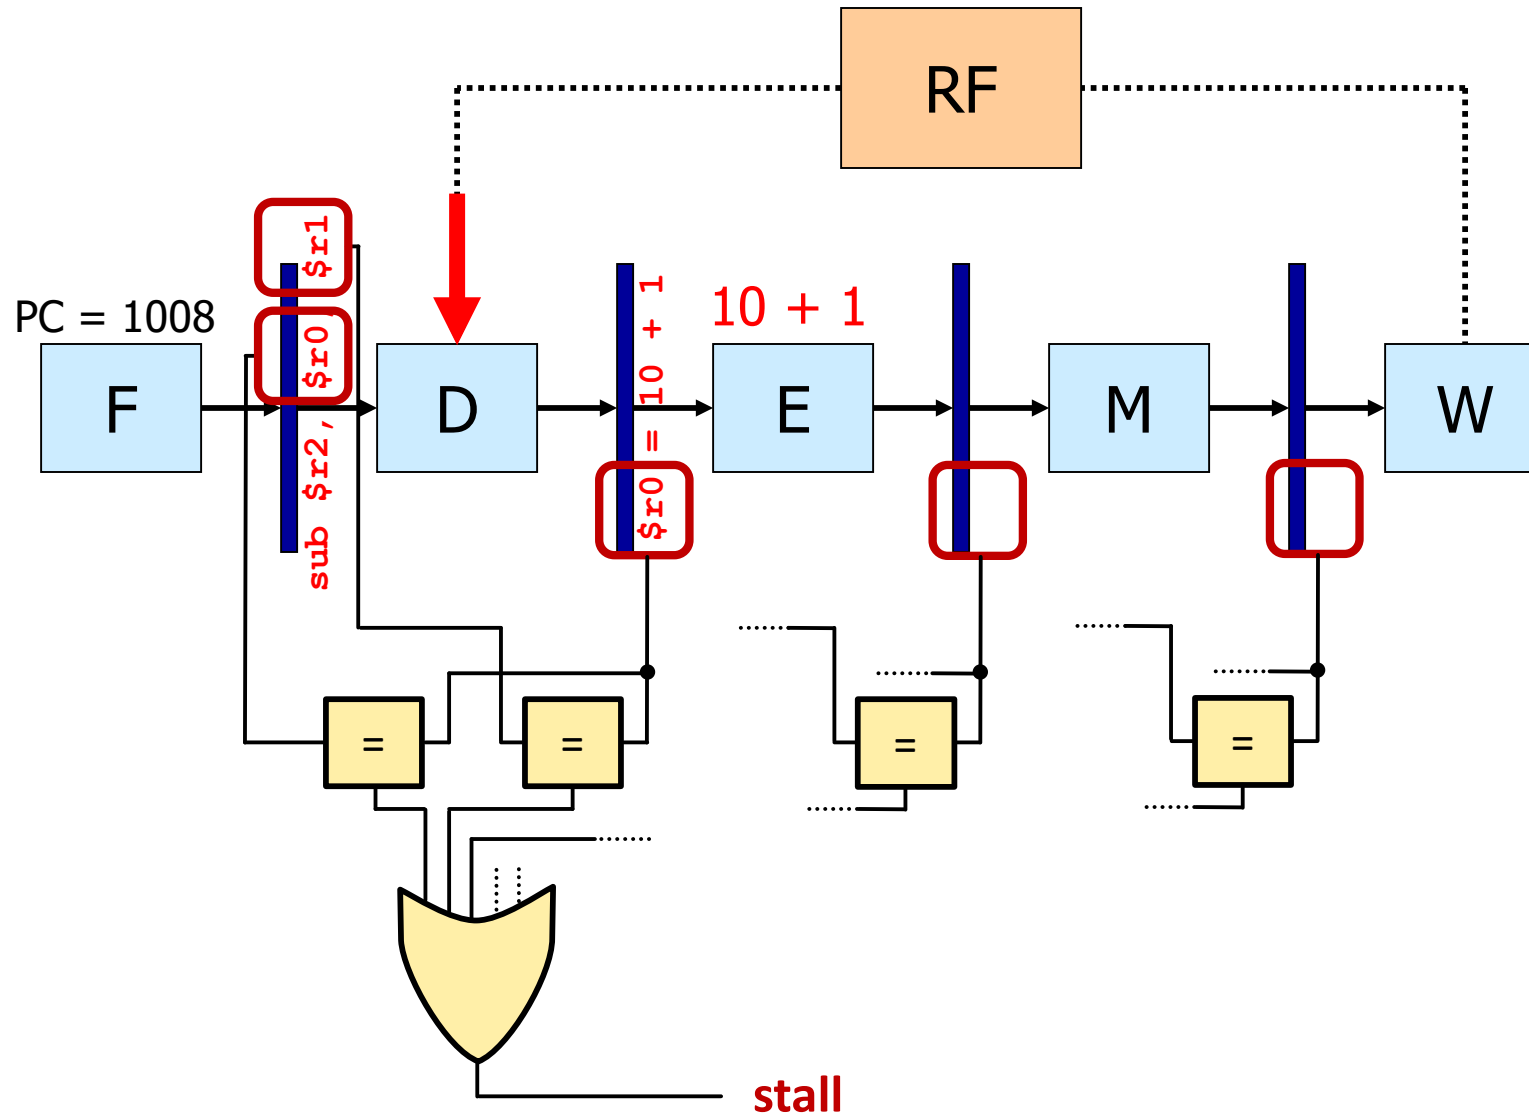
\$r2, \$r0, \$r1



\$r0, \$r0, 1

\$r2, \$r0, \$r1

Detecting

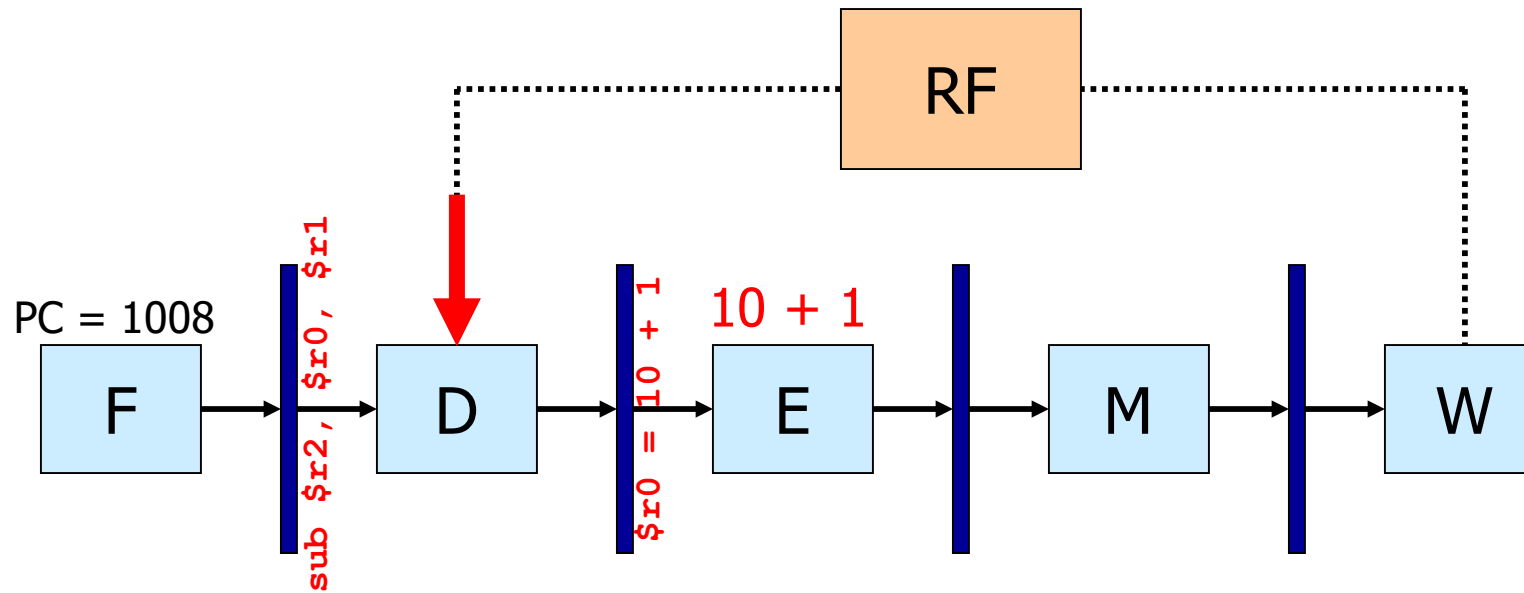


1000: addi

\$r0, \$r0, 1

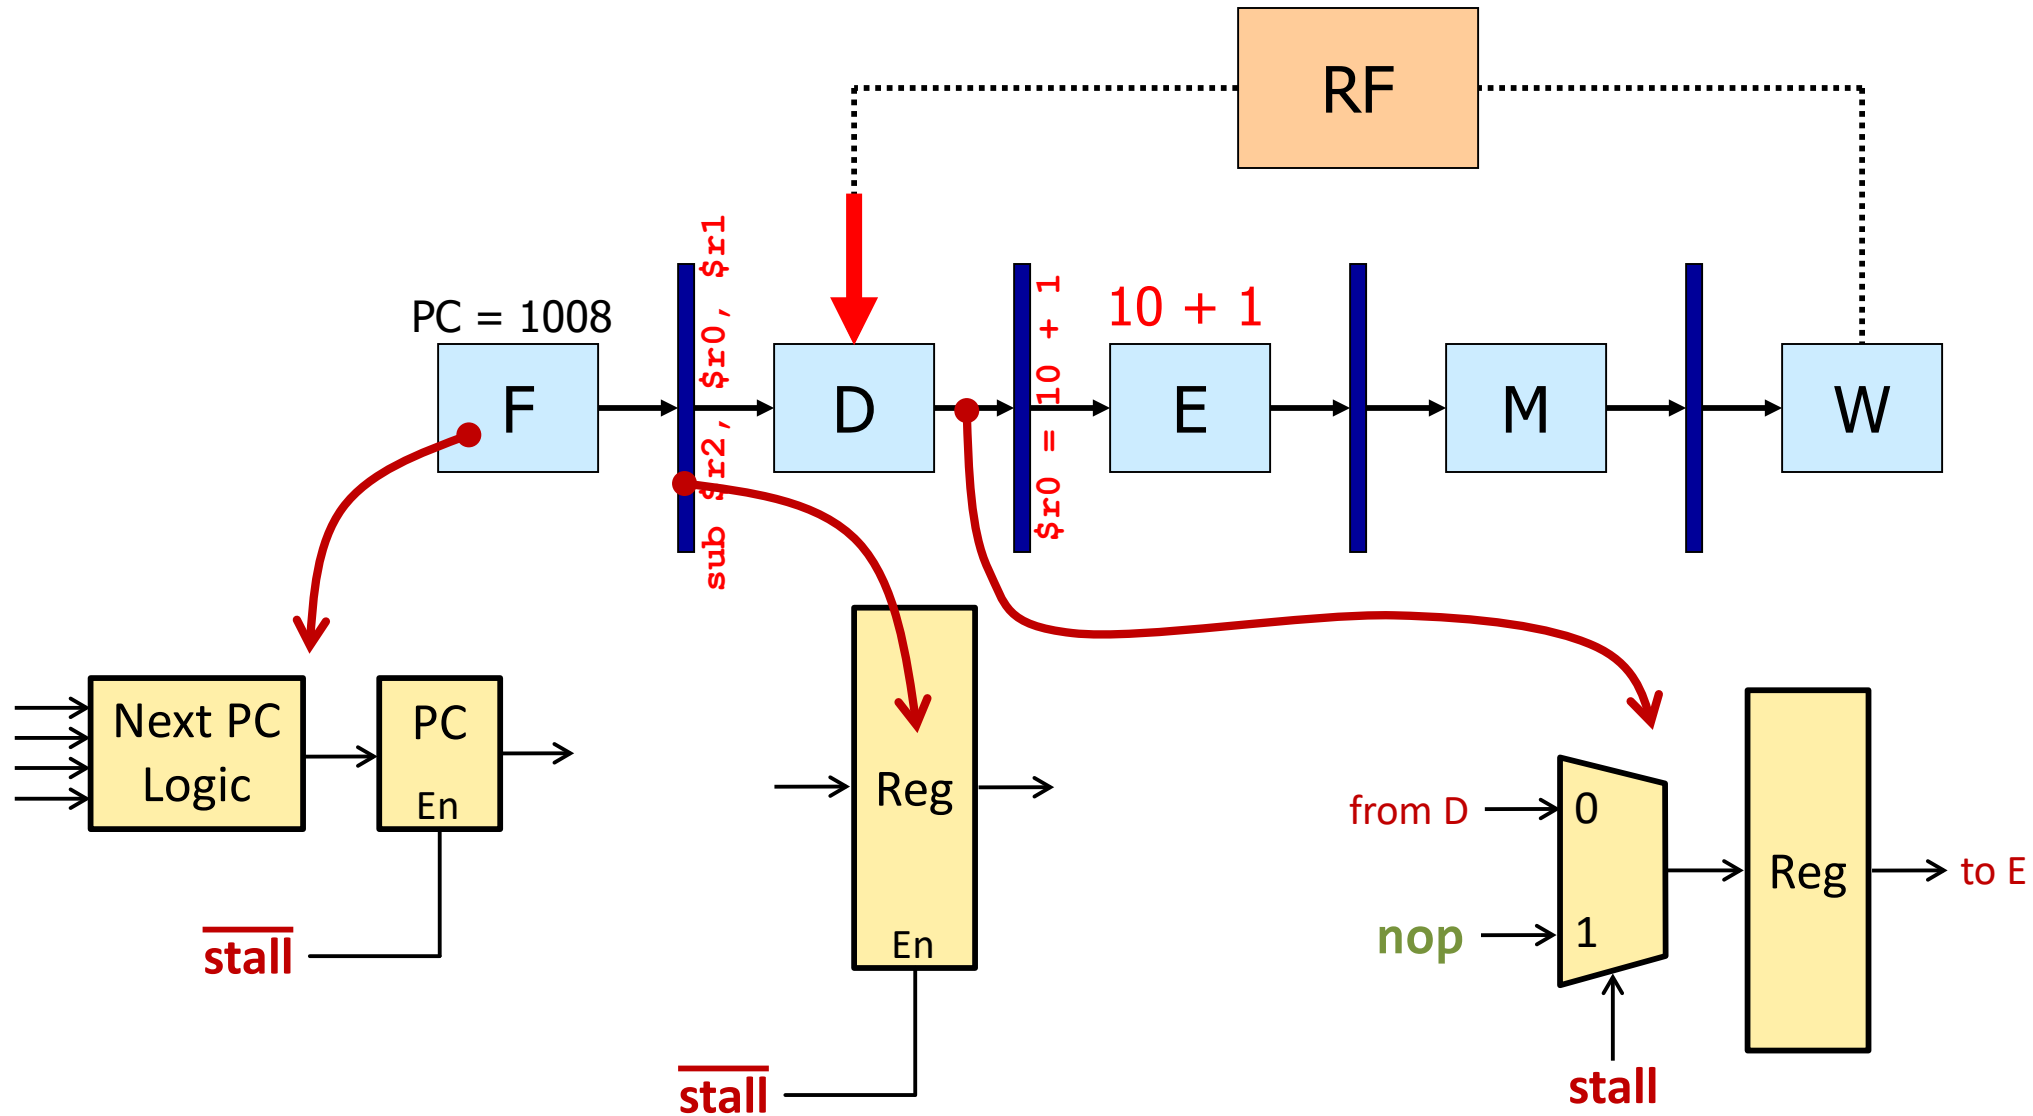
1004: sub

\$r2, **\$r0,** \$r1

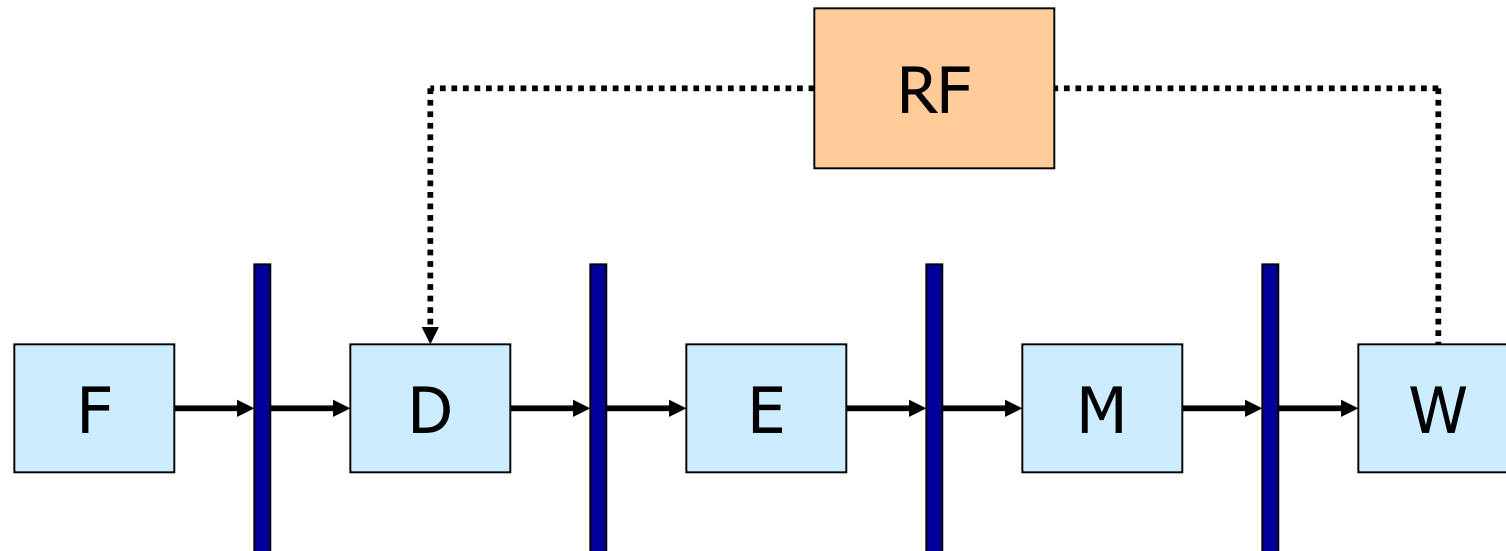


1000: addi \$**r0**, \$r0, 1
1004: sub \$r2, **\$r0**, \$r1

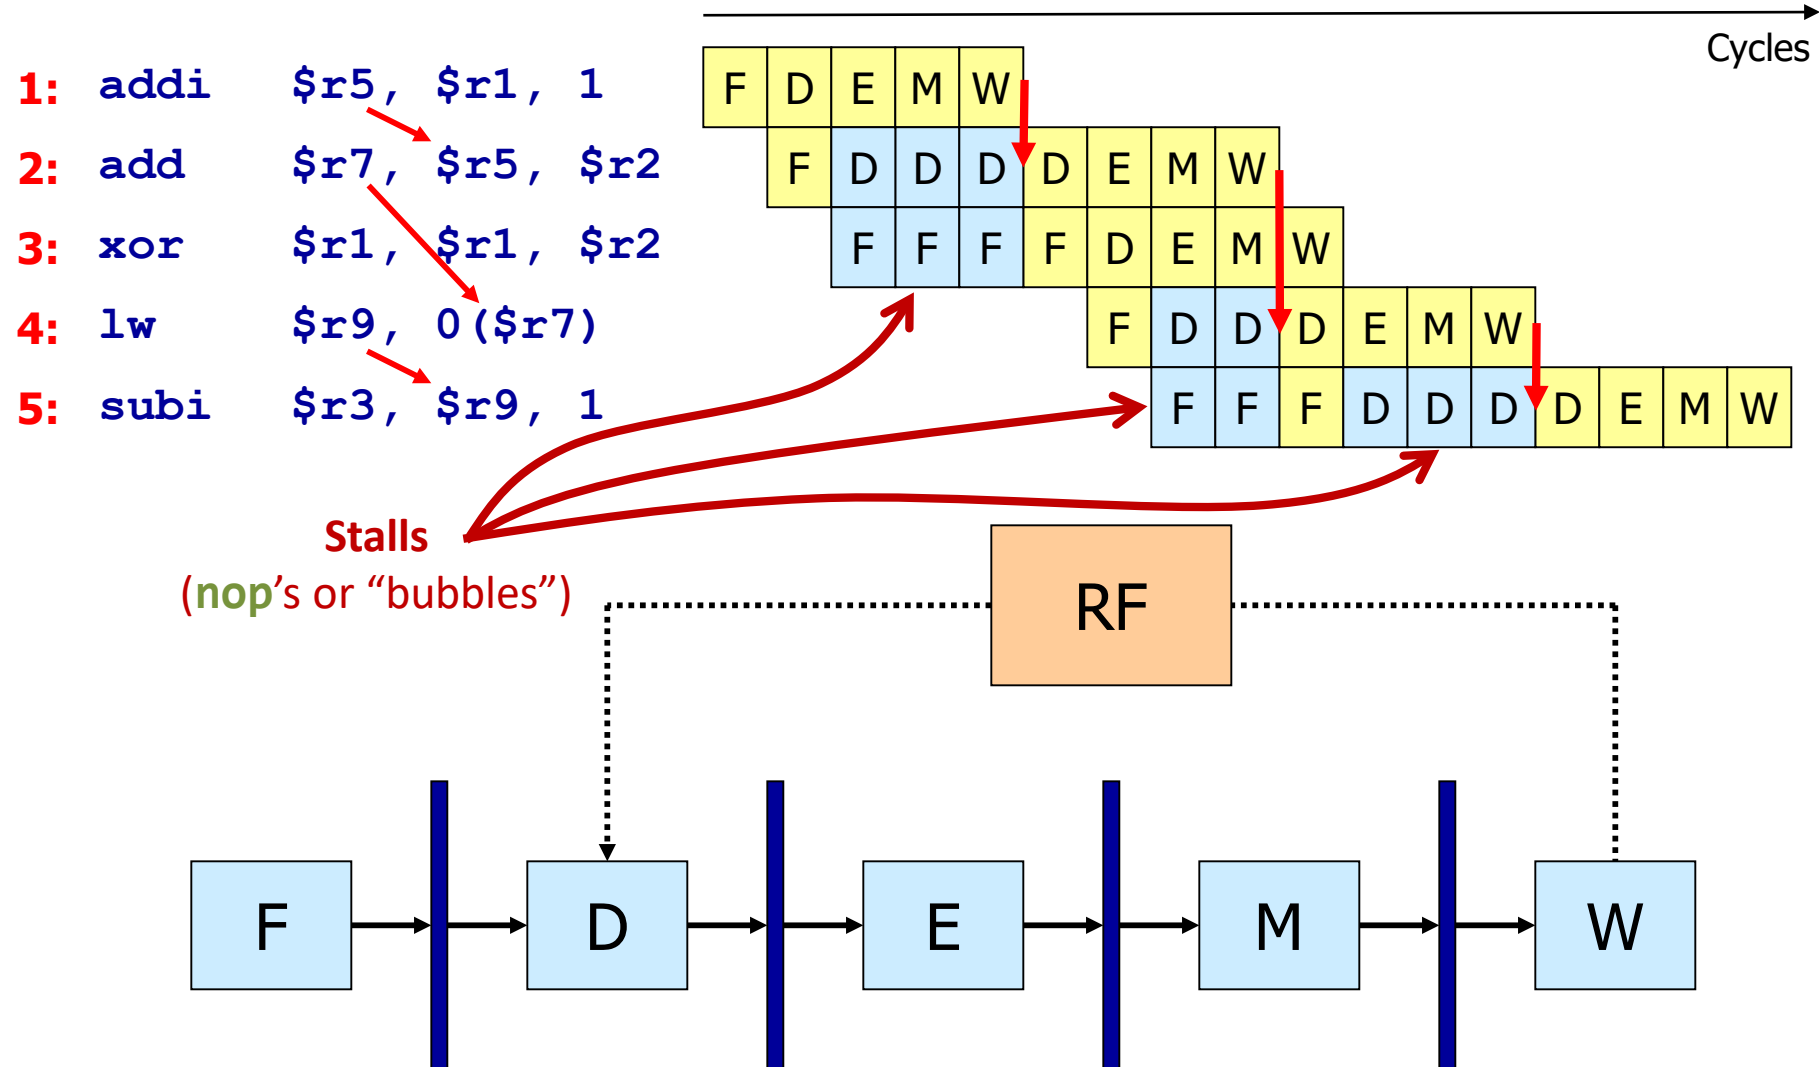
Stalling



1: addi \$r5, \$r1, 1
2: add \$r7, \$r5, \$r2
3: xor \$r1, \$r1, \$r2
4: lw \$r9, 0(\$r7)
5: subi \$r3, \$r9, 1



Data Hazards




Another Solution?

```

1: addi $r5, $r1, 1
2: add  $r7, $r5, $r2
3: xor  $r1, $r1, $r2
4: lw   $r9, 0($r7)
5: subi $r3, $r9, 1
    
```

This solution
needs some
hardware


This solution
needs a good
compiler



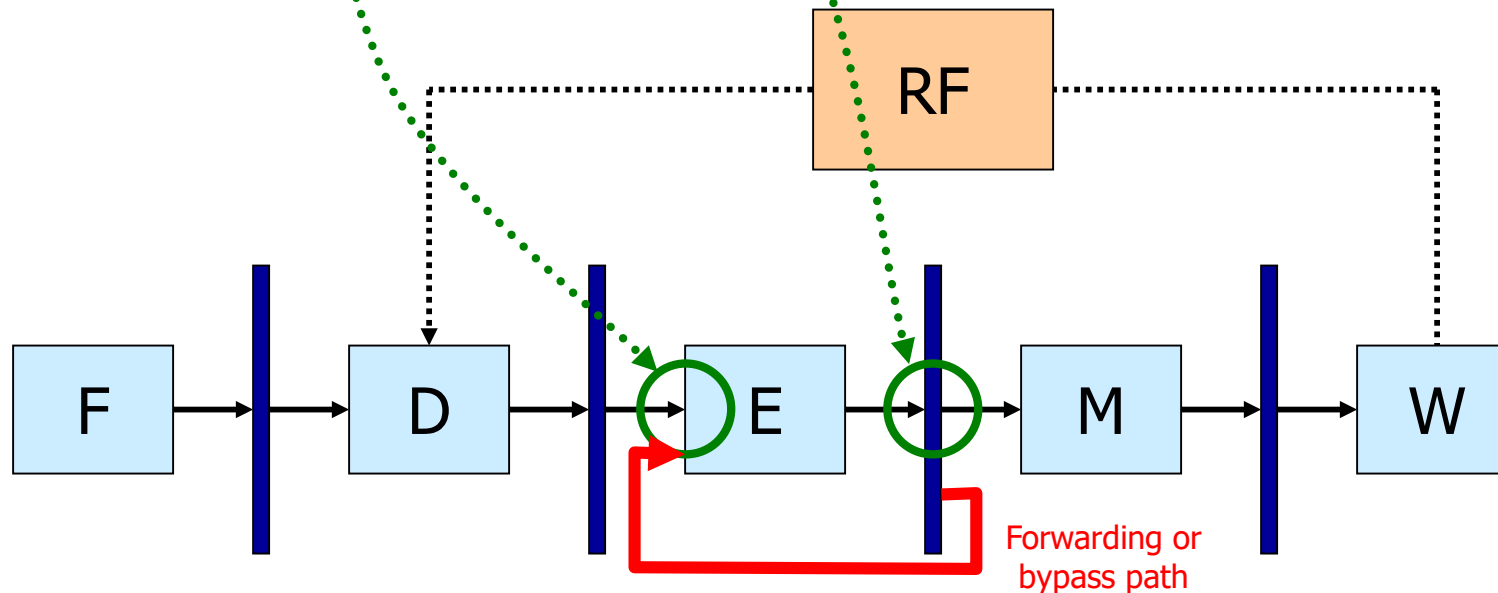
```

1: addi $r5, $r1, 1
2: nop
3: nop
4: nop
5: add  $r7, $r5, $r2
6: xor  $r1, $r1, $r2
7: nop
8: nop
9: lw   $r9, 0($r7)
10: nop
11: nop
12: nop
13: subi $r3, $r9, 1
    
```

Data Hazards Solved by Forwarding Values

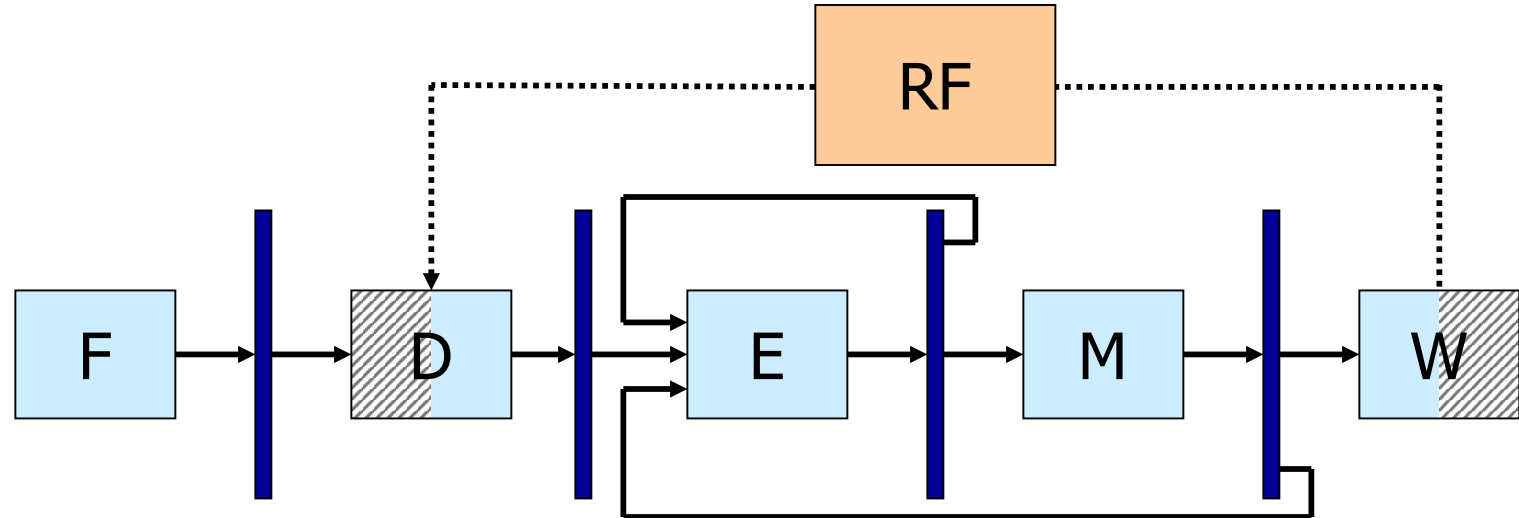


- Data are available at the **right time** but in the **wrong place**
- Add a **forwarding path** to bring it where it needs
- One also needs logic to **select (MUX) the right input** to the E stage

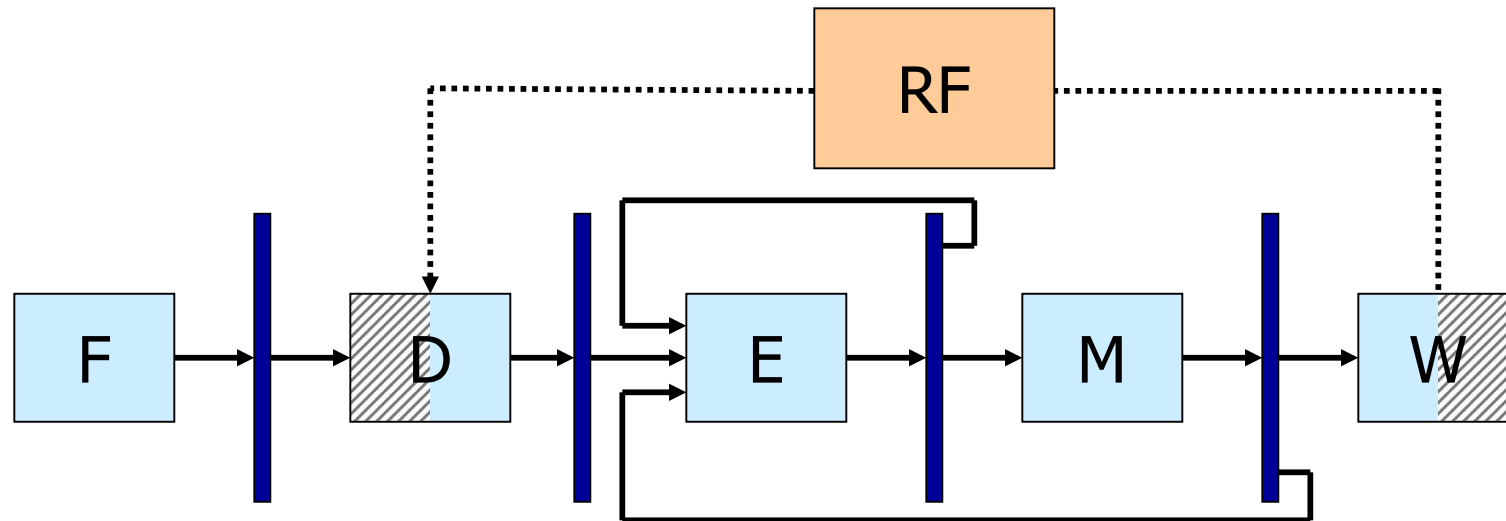


Classic MIPS Pipeline

- 5-stage pipeline with **all forwarding paths**:
 - $E \rightarrow E, M \rightarrow E$
 - $W \rightarrow D$
- The **register-file forwarding** ($W \rightarrow D$) is a special case:
 - During **W**, registers are written in the **first half** of the cycle
 - During **D**, registers are read in the **second half** of the cycle
 - a register can be correctly written and read in the same cycle
- Not always all forwarding path exist...




1: addi \$r5, \$r1, 1
2: add \$r7, \$r5, \$r2
3: xor \$r1, \$r1, \$r2
4: lw \$r9, 0(\$r7)
5: subi \$r3, \$r9, 1



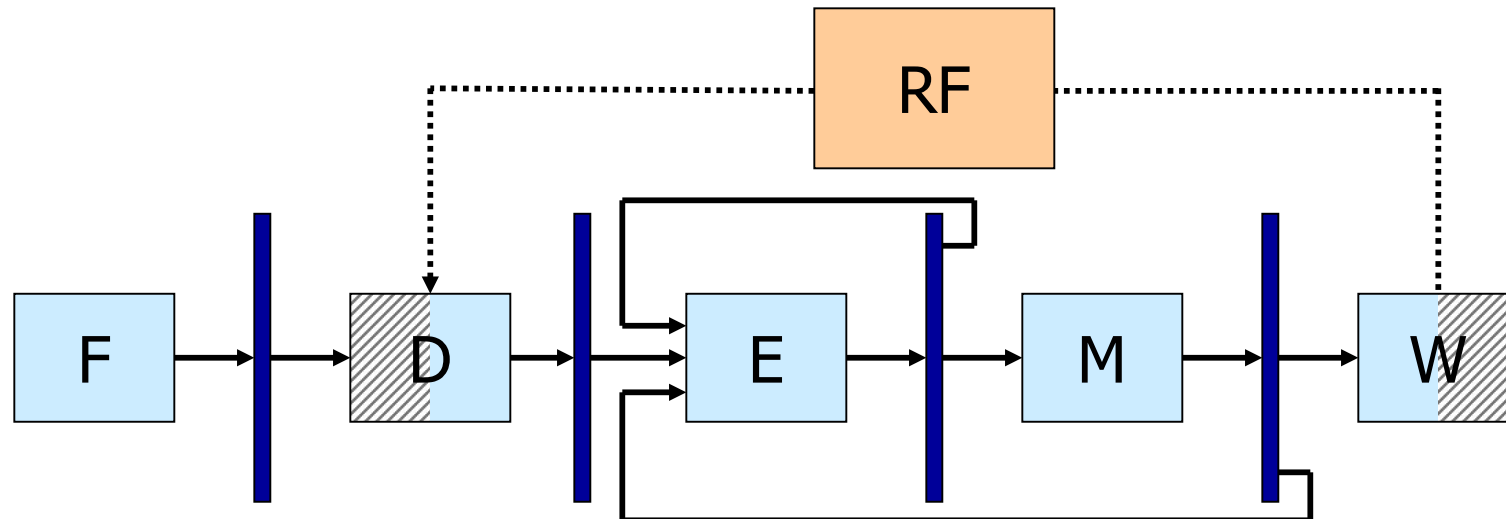
Reduced Data Hazards

```

1: addi $r5, $r1, 1
2: add  $r7, $r5, $r2
3: xor  $r1, $r1, $r2
4: lw   $r9, 0($r7)
5: sub  $r3, $r9, $r1
    
```




–
E→E
 –
M→E
M→E (\$r9)
W→D (\$r1)




Used forwarding paths
 (if any were missing,
 stalls would have been
 inserted)

Structural Hazards

- A **structural hazard** happens when different instructions compete for the same resource (e.g., pipeline stage)
- Structural hazards **cannot happen** in this pipeline
- If we did not stall also instructions following one missing an operand, we could have structural hazards




Control Hazards



Causality violation!

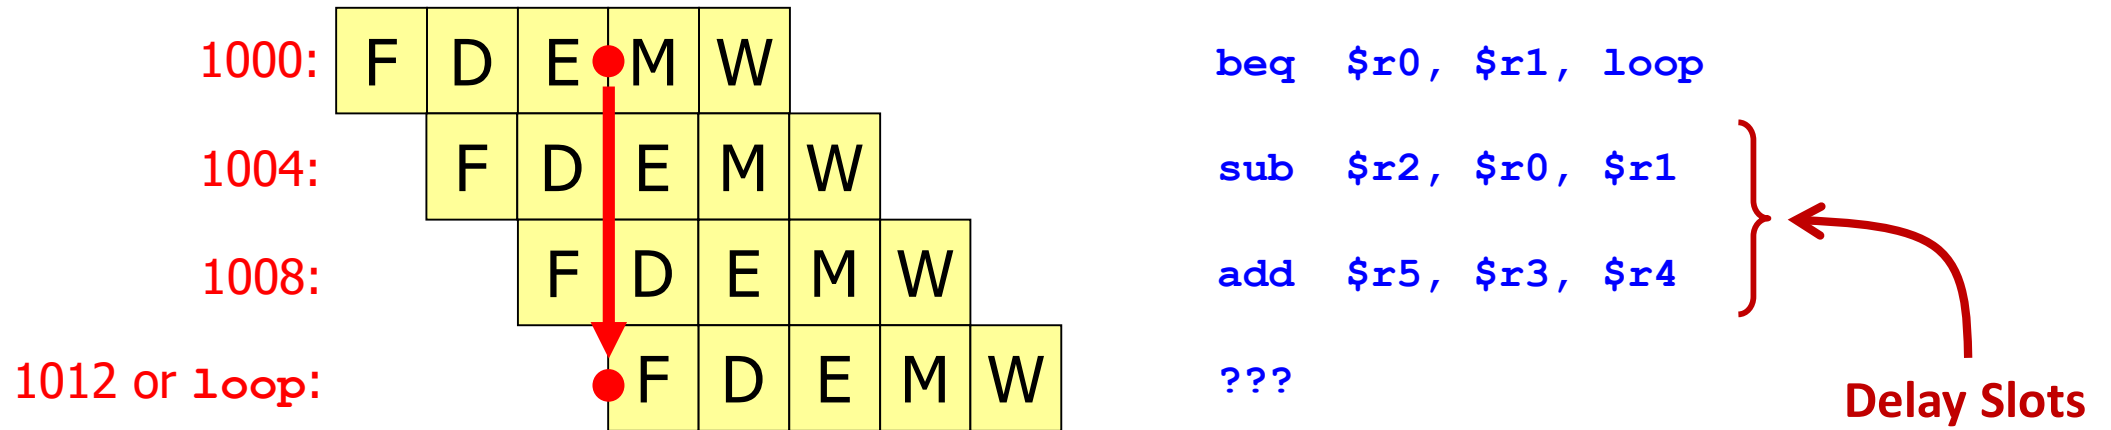
We fetch an instruction before we know which one!

Control Hazards Solved by Stalling the Pipeline



- Similarly to the way we solve data hazards, we can stall the pipeline once it is discovered, after **D**, that an instruction was a branch
- If, for instance, the correct address of the next instruction is known at the end of the **E** stage, 2 cycles are lost **every branch**

Control Hazards Solved by Defining Delay Slots




- Alternatively, we can **modify the definition of the architecture** and decide that the two instructions following a branch are **executed in any case** (branch taken or not)
- These instructions after the branches are called **delay slots**
- Note that code becomes **counterintuitive!!!**
- Quite rare in current architectures

Use of Delay Slots

- A simple way of using delay slots is to use them for **nop**'s—but then it is not better than stalling the pipeline
- A better idea is to put there instructions which precede the branch and on which the branch has **no dependence**
- Suppose an architecture with two delay slots:

```
1000: sub  $r2, $r0, $r7
1004: mul  $r1, $r6, $r7
1008: add  $r5, $r3, $r4
1012: beq  $r0, $r1, loop
1016: nop
1020: nop
1024: lw   $r8, 12($r9)
```



```
1000: mul  $r1, $r6, $r7
1004: beq  $r0, $r1, loop
1008: sub  $r2, $r0, $r7
1012: add  $r5, $r3, $r4
1016: lw   $r8, 12($r9)
```

RAW, cannot be moved

Delay slots

Branch Prediction

- A better strategy is to **guess the branch outcome** and fetch the corresponding instruction (either the next instruction or the branch destination, but not necessarily the former)
 - If the guess is correct, **no cycle is lost**
 - If the guess is wrong, what has been fetched and decoded is thrown away (**“squashed”**)
- Branch predictors of modern processors are extremely sophisticated: dynamic predictors **learn from previous executions** of a branch...
- Complex predictors can be **correct up to 95-99%** of the time
- The quality of branch predictors has made architectures with delay slots extremely rare


Architecture and Microarchitecture

- **Architecture**: what is in the ISA contract
 - Instructions, registers, etc.
- **Microarchitecture**: what is specific of an implementation
 - Multicycle vs. pipelined, FSM or pipeline structure, etc.


Some of the solutions evoked (reschedule instructions, add **nop**'s, delay slots) **expose typical microarchitectural aspects** (pipeline structure) **in the architecture** → the same binary **does not run** on different processors!

Starting Point (Programmer Model)

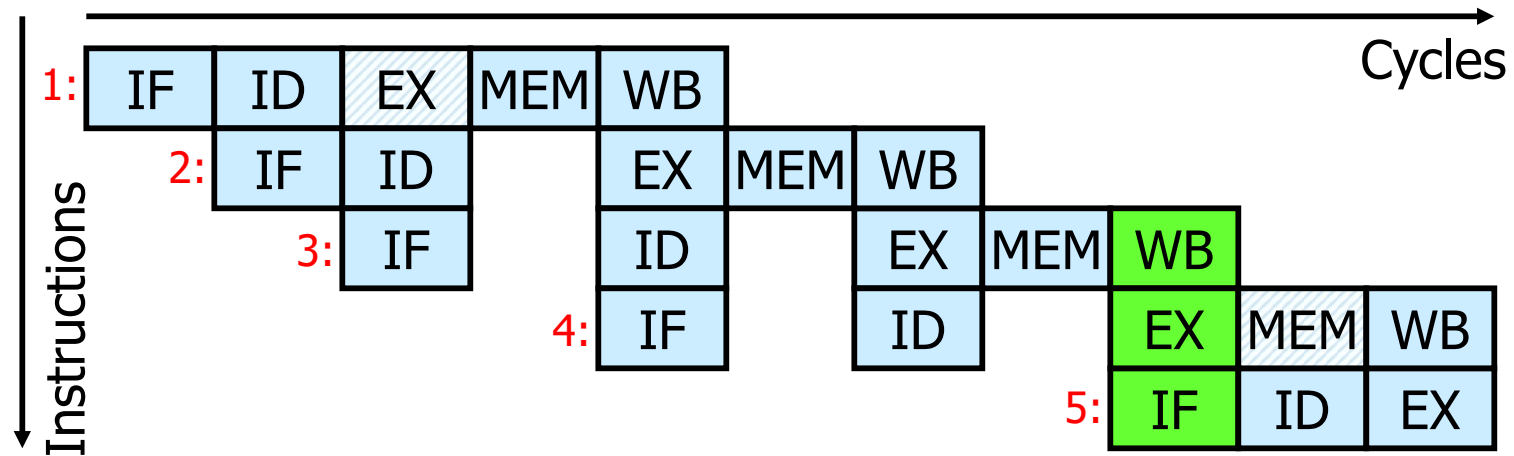
- Sequential multicycle processor



Instruction Level Parallelism?



First Step: Pipelining



- Simplest form of **Instruction Level Parallelism** (ILP): several instructions are now executed at once

Three Types of Hazards Hinder Pipelining

- **Data Hazards** (= data dependences)
 - *Solutions:*
 - Forwarding paths, wherever possible
 - Stalls, in all other cases
- **Control Hazards** (= jumps and branches)
 - *Solutions:*
 - Delay slots, if the architecture allows it
 - Branch prediction, to try to do the right thing
 - Stalls, if not
- **Structural Hazards** (= conflicting need for a resource)
 - *Solutions:*
 - Rigid pipelines which cannot have structural hazards by construction
 - Stalls, otherwise

References

- Patterson & Hennessy, COD – RISC-V Edition
– **Sections 4.6-4.9**